

# Functional Pearl: Holey Generators!

JOSEPH W. CUTLER, University of Pennsylvania, USA

HARRISON GOLDSTEIN, University of Pennsylvania, USA

KOEN CLAESSEN, Chalmers University of Technology, Sweden

JOHN HUGHES, Chalmers University of Technology, Sweden

BENJAMIN C. PIERCE, University of Pennsylvania, USA

An attractive feature of testing frameworks like QuickCheck is their domain-specific languages for custom generators of random test data. Programmers value such handcrafted generators for two reasons: they can guarantee *invariants*, such as ordering constraints on binary search trees, by construction, and they can control the *distributions* of generated values.

How easy is it to tune distributions to achieve desired properties? Surprisingly hard, as it turns out! We investigate the distributions produced by some familiar generation strategies for unlabeled binary trees and observe that it is quite challenging to achieve both a good distribution of tree sizes and a good distribution of tree shapes. The fundamental issue is *locality of control*—the way QuickCheck generators for recursive data structures are naturally written as recursive functions makes it *unnatural* to express generators that exert global control over distributional properties.

We propose instead a novel abstraction, *holey generators*, that makes a sequence of *global* random choices about where to incrementally extend a tree. This abstraction supports much more direct control over distributions; we show how changing a single parameter yields bushy, stringy, left-leaning, right-leaning, and (most challenging) uniformly distributed trees. Moreover, the core applicative combinators for building holey generators can be extended with a monadic interface, supporting generators for data with invariants, like ordered trees and heaps. Finally, we evaluate holey generators in practice, showing that they easily achieve testing performance comparable to expert-tuned classic generators. We concentrate on the case of binary tree structures, with and without invariants; at the end, we discuss prospects for generalizing to other data types.

## 1 INTRODUCTION

*This book fills a much needed gap.*

— Saul Gorn

Property-based testing (PBT) is a popular bug-finding technique, particularly in the Haskell community where QuickCheck [3] is the *de facto* testing tool of choice. QuickCheck allows users to define *properties*—Boolean-valued functions that validate a system’s behavior on a single given input—and test that they return `True` on many randomly generated arguments. The process of generating these random arguments can be automated much of the time, but to efficiently generate data structures with invariants, like binary search trees (BSTs), the programmer must write *generators*: functions expressed using QuickCheck’s domain-specific language which return randomly-chosen structures that are valid with respect to the invariants.

A simple QuickCheck generator for BSTs might be written like this:

```
data Tree = Leaf | Node Tree Int Tree

genBST :: (Int, Int) -> Gen Tree
genBST (lo, hi) | lo >= hi = return Leaf
genBST (lo, hi) =
  oneof [return Leaf,
        do
          x <- choose (lo, hi)
          l <- genBST (lo, x - 1)
```




Fig. 1. Left: The size distribution of one million generated BSTs. Right: the shape distribution of BSTs of size 8, ordered shortest to tallest by depth (note that the smallest possible depth is 4).

```
r <- genBST (x + 1, hi)
return (Node l x r)]
```

This produces valid BSTs by generating a value in a range and recursively generating children whose keys fall in appropriate smaller ranges. It uses QuickCheck’s monadic `Gen` abstraction and combinators such as `choose` (which samples from a discrete range) and `oneof` (which invokes a generator chosen at random from a list) to build up larger generators from smaller ones.

This rather naïve generator is not very useful for finding bugs, as many have observed. To see why, let’s look at the *distribution* of values that it produces (Figure 1). Each of these graphs points to a serious problem. The first shows that most of the values produced by this generator have size 1 or 2; these are far too small to catch many bugs, and even if they could, there is no point in trying Leaf 500,000 times! The second plot, showing the different shapes of size-8 trees that were generated, highlights that some shapes of trees are more popular than others. The generator seems to prefer the short, bushy trees on the left of the graph over the tall, stringy trees on the right.

But wait—doesn’t QuickCheck provide tools for addressing exactly these kinds of issues? Yes, it does! The “real-world” version of the above generator would be closer to this:

```
genBST :: (Int, Int) -> Gen Tree
genBST bnd = sized (aux bnd)
  where
    aux (lo, hi) n | lo >= hi || n <= 1 = return Leaf
    aux (lo, hi) n =
      frequency [(1, return Leaf),
                (5, do
                  x <- choose (lo, hi)
                  l <- aux (lo, x - 1) (n `div` 2)
                  r <- aux (x + 1, hi) (n `div` 2)
                  return (Node l x r)))]
```

This version uses two different techniques to get better distributional control. First, it uses `sized` to extract QuickCheck’s hidden *size* parameter and pass it around as the `n` parameter to the recursive `aux` function. By cutting off generation when `n` reaches 1 and halving it as the generator recurses, we ensure that the tree does not get too large. If the *size* parameter is set to 30 for example, this generator will not produce trees that are more than  $\log_2(30) \approx 5$  nodes deep. (During testing, the *size* parameter automatically ranges from 0 to 99; even at its maximum value, generated trees cannot be more than 6 nodes deep.) Additionally, this version replaces the `oneof` combinator (which

makes uniform choices) with frequency (which annotates choices with weights), preferring Nodes to Leafs at a 5-to-1 ratio.

Does this fix our problems? Sadly, not really. Take a look at the graphs now:




Fig. 2. Left: The size distribution of one million generated trees using the tuned generator. Right: the shape distribution of the tuned generator for trees of size 8, with x axis ordered by depth.

The sizes are more reasonable: using `sized` prevents *very* large trees and puts a bit more weight on mid-sized trees, between 6 and 12 nodes. But the shape distribution is *much worse*: every size-8 tree generated has depth exactly 4, leading to only 95 of the possible 1430 shapes occurring at all, with two of the shapes accounting for nearly a quarter of the total draws! This doesn't work either.

One could go down a rabbit hole here, attempting to solve this problem with increasingly complex generators. Some experienced QuickCheck users like to dynamically compute frequency weights using the size parameter, randomly split the size in different parts of the tree, and employ other sophisticated tuning strategies. These techniques can get testers closer and closer to their desired behavior, but they require more and more effort and expertise to implement. Ultimately we are left with a question: why is it so difficult to achieve declarative control over tree sizes and shapes?

The problem with the methods we've discussed so far is that control over distributions is applied *locally*, without regard for broader context. When we recursively build a binary search tree by building left and right subtrees, the recursive calls don't know how their results will be used in assembling a larger tree. In BSTs, this lack of communication results in generated trees with a short and bushy structure; long, stringy trees are very unlikely to occur.

What we need is a generic way for different parts of the generation process to depend on one another: decisions made to generate one subtree should be able to influence the decisions made to generate the next ones. In other words, instead of local control, we want *global* control over the distribution. Global control is difficult to achieve because it requires more complex program structure than straightforward recursive functions. We will need a better abstraction than classic QuickCheck generators if we hope to have a usable interface that allows for global control.

But why (you might ask), if we want more control over the distribution of tree shapes, do we not use a system like FEAT [4] instead of classic QuickCheck? This is a reasonable question! FEAT can generate values of any algebraic datatype at any specified size, with a controllable distribution, so in particular it can certainly generate binary trees that are better distributed than what we've seen. The problem with FEAT is that it has trouble with generators for data satisfying additional invariants (like *binary search trees*).

To see why, notice that the classic QuickCheck generator above relies on Haskell's `do`-notation to sequence generators that depend on one another:

```

do
  x <- choose (lo, hi)
  l <- aux (lo, x - 1) (n `div` 2)
  r <- aux (x + 1, hi) (n `div` 2)
  return (Node l x r)

```

Here, the calls to `aux` take `x` as an argument, which is only available after the call to `choose`. FEAT does not support this kind of monadic dependency: it only provides an `Applicative` abstraction that runs two component generators independently and combines their results.<sup>[15]</sup><sup>1</sup> It appears that FEAT and similar interfaces do not provide enough power for our use cases.

So... what to do? Classic QuickCheck generators provide a powerful monadic interface but give only local control over the distributions of shapes and sizes of the values generated. We do not seem to have a way to build generators that (a) allows users to write in an easy-to-use expressive monadic style, while also (b) providing *global* control over the distributional choices that influence the structures they generate.

The first step in addressing a problem is to understand that you have it. Accordingly, the first contribution of this paper is to bring this problem to the attention of the larger community. The deficiencies of local control appear in lots of abstractions for generating random data. All recursive QuickCheck generators must deal with the same issues that motivate this exploration, since they use the same inherently problematic local distribution control methods. This is clearly not a satisfactory state of affairs, and we hope this paper serves as a call to arms for PBT researchers to study the problem of global distributional control more closely.

Solving this problem in general appears quite challenging (see Section 7). As a first step, we attack the case of binary trees, with and without invariants, developing an approach to the problem of generation with global control and monadic invariants that solves this case. This *holey generator* technique is a new method for writing generators that combines global control with a monadic interface for invariant maintenance.

The presentation proceeds in two stages. In the first (Section 2), we introduce the basic idea of our technique and demonstrate how to use it to generate unlabeled binary trees using a simple applicative interface; Section 3 then shows how to instantiate this interface to obtain a uniform distribution over tree shapes of a given size. In the second stage (Section 4), we introduce the rest of the holey abstraction and explain how it gives global control over the distributions of labeled tree types with constraints on labels, like BSTs. Section 5 demonstrates that our newfound distributional control makes it painless to write a holey generator that performs as well as a painstakingly-tuned classic one on a slate of tests drawn from *How To Specify It!* [11]. Finally, Section 7 discusses future directions, including ways of generalizing the holey approach beyond binary trees, as well as other methods for distributional control.

## 2 HOLE-FILLING GENERATORS

As we learned through our experiments in Section 1, the fundamental problem behind the poor distributions generated by classic QuickCheck generators is that they make random decisions about the shape of their outputs *locally*, without taking into account the context that the generated value will be placed into. Our goal, then, is to design an abstraction that empowers users to make random choices that depend on the global state of the structure that's been generated so far, while still allowing the programmer to write generators in a familiar recursive style. It is not easy to do this

<sup>1</sup>Technically, FEAT *can* support a monadic interface, but random generation using it is intractable. The `bind (>>= :: Enumerate a -> (a -> Enumerate b) -> Enumerate b)` would need to eagerly enumerate all of the values of type `a` to generate a single value of type `b`.




Fig. 3. Generating a tree with hole filling

with the classic generator abstraction, since a generator (like any function) cannot observe the context in which it has been invoked without the programmer explicitly threading around some kind of state, which would interrupt the “recursive style” of generator writing.

In this section we describe our abstraction, explain why it works, and discuss how to use it to obtain specific distributions.

**Filling Holes.** To untie the knot created by recursive generators, we subtly shift how we think about generating values. Instead of generating entire structures recursively in one go, we take a step-by-step approach where structures are generated by repeatedly replacing one leaf, somewhere in the structure, with a node.

We can operationalize this new perspective by defining a recursive generator type `Holey a` as a sort of state machine whose states are binary tree structures of type `a` with “holes at the leaves.” Each transition in this state machine represents the “filling” of a hole by replacing it with either a new node constructor (with two new holes as children), or a leaf constructor. The state machine presentation of `Holey a` is suggestive of how it can be run: to generate a value of type `a`, we repeatedly transition the system some number of times and take the final state element as the resulting value. Figure 3 shows the process of building an unlabeled binary tree through repeated hole-filling.

Concretely, this machine state is represented as (1) a value of type `a` recording the tree that has been constructed so far, plus (2) a “skeleton” of a binary tree that we call an `HTree`, whose structure mirrors that of the value `a` exactly.

```
data Hole = Here | L Hole | R Hole
```

```
data HTree = HoleLeaf | DoneLeaf | HNode HTree HTree deriving (Eq,Ord,Show)
```

The `HTree` has two kinds of leaves: `HoleLeaves`, which mark the location of a hole, and `DoneLeaves`, which mark leaves in the current tree which cannot be further extended with a new node.<sup>2</sup>

Formally, a `Holey a` is a record with three fields.

```
data Holey a = Holey
  { done :: a,
    treeOfHoles :: HTree,
    fill :: Hole -> Holey a
  }
```

The first is `done :: a`, the current state of the partially-completed tree of type `a`. The second is an `HTree` whose structure mirrors that of the `done` value, but whose leaves mark either completed sub-trees, or holes which can yet be filled. The final field is the transition function `fill :: Hole -> Holey a`, which, given the location of a hole in the `HTree` as a path from its root to the leaf, returns a new state where the hole in question has been filled with a new structure, possibly containing more holes. In the case that `fill` is called with a path to a hole that is not actually present in the `HTree`,

<sup>2</sup>In a `Holey` gen for *unlabeled* trees, all of the leaves in the `HTree` are `HoleLeaves`. Because of this, the content of this section works just as well with `Holey a` defined without the `HTree` field, using the only `done` field as the state. While we will not see trees with labels until Section 4, it seems best to introduce the full `Holey` abstraction here.

the generator will fail. In practice, this will never happen, since `fill` is only used by the function that *runs* the holey generator.

Given a `Holey a`, we can repeatedly fill randomly chosen holes from the hole tree to build up a final value of type `a`. Crucially, since we have access to the `HTree` when we pick which hole to fill, the choice of hole can depend on the structure of the hole tree, and by proxy, the entire structure that has been generated so far. Moreover, since hole-filling adds one node at a time, we can control *exactly* how large our structures will be! A `Holey a` generator separately gives control over the shapes of your trees using global hole choices and control over the sizes by choosing the number of holes to be filled.

**Generators 'n Combinators.** To help users effectively build holey generators, we provide an instance for the Haskell `Applicative` typeclass, which defines a way of combining two independent holey generators. A precondition for this is that we must also provide an instance for the `Functor` typeclass: a function `fmap` which lifts a function between binary tree types `a -> b` to `Holey a -> Holey b`. Given `f :: a -> b` and `r :: Holey a`, we update `r` by applying `f` to the partial tree `done r` and post-composing the transition function `fill r` with a recursive call `fmap f` to transform the next states.

```
instance Functor Holey where
  fmap f r = r {fill = fmap f . fill r, done = f (done r)}
```

The applicative interface for `Holey` is more interesting, as it substantiates parts of the discussion that, so far, have been mostly hand-waving. The applicative “combination function” `<*> :: Holey (a -> b) -> Holey a -> Holey b` is where all the action happens.

```
instance Applicative Holey where
  pure x =
    Holey
      { treeOfHoles = DoneLeaf,
        fill = (error "No holes left to fill!"),
        done = x
      }
  rf <*> rx | isDone rf = done rf <$> rx
  rf <*> rx | isDone rx = ($ done rx) <$> rf
  rf <*> rx =
    Holey
      { treeOfHoles = HNode (treeOfHoles rf) (treeOfHoles rx),
        fill = \case
          L h -> fill rf h <*> rx
          R h -> rf <*> fill rx h
          Here -> error "No holes left to fill!",
        done = done rf (done rx)
      }
```

When we combine two holey generators into one with `<*>`, the combined generator’s remaining holes are those of the two arguments—that is, the tree of holes of the combined generator has as subtrees the two trees of the argument generators. Since the `treeOfHoles` of the combined generator has two subtrees, a call to `fill` will provide a `Hole` which is either `L h` or `R h`. In the former case, the hole `h` on the left argument `rf` is filled; in the latter case, the hole `h` in `rx`.

Critically, the `HTree` manipulations in `<*>` ensure that the shape of the current `HTree` tracks the shape of the current partial value being generated. When two nontrivial (i.e., “both sides not pure”) generators are combined with `<*>` (as in `UNode <$> holeyUTree <*> holeyUTree`) the internal `HTree` gains an `HNode`. But, if either of the two argument generators have no holes, they can be combined with the other *without* adding a node to the `HTree` using the `fmap` function from the functor instance

(written infix with `<$>`). This means that the internal state of the generator can stay in sync with the value it is generating.

The applicative interface also requires a function `pure :: a -> Holey a`, which constructs a “trivially holey” generator from a value. Given `x :: a`, it returns the generator with no holes to fill.

We next define a combinator we call `orFill`, which emulates a common use of the `oneOf` function in QuickCheck: to provide a base case to a recursively-defined generator. For a base-case value (usually a leaf) `x` and a holey generator `r`, we define `x `orFill` r` to be the generator whose current tree is a leaf node `x`, with a single hole that, when filled, results in the generator `r`. This combinator commonly provides the outermost structure of a generator, choosing between stopping the recursion at a leaf or continuing with a recursive call.

```
orFill :: a -> Holey a -> Holey a
orFill x r = Holey {treeOfHoles = HoleLeaf, fill = \Here -> r, done = x}
```

With these combinators, we can write a simple holey generator for unlabeled trees.

```
data UTree = ULeaf | UNode UTree UTree deriving (Eq, Ord, Show)
```

```
holeyUTree :: Holey UTree
holeyUTree = ULeaf `orFill` (UNode <$> holeyUTree <*> holeyUTree)
```

### Sampling from Holey Generators.

When deciding what holes to fill next, the generator can (1) inspect its internal state to find out the shape of the value that it has built so far, (2) pick one of the holes it sees, and (3) update its internal state in tandem with updates to the value being generated. This leaves the question of choosing a hole in step (2). How one makes these random choices determines the distribution over shapes that the generator will denote, and a good distribution makes all the difference in finding bugs quickly. To this end, we define a `HoleWeighting` to be a function `HTree -> [(Int,Hole)]` mapping states of the hole tree to a list of weighted holes in that tree to choose from. Holes with higher weight are chosen with higher probability, and lower weights are chosen with lower probability. Formally, given the weighted list `[(n1,h1), (n2,h2), ... (nk,hk)]`, we sample which hole to fill next from the categorical distribution over the holes `h1 ... hk` with probabilities equal to the associated weight divided by the sum of all the weights.

```
type HoleWeighting = HTree -> [(Int,Hole)]
```

Given a `HoleWeighting`, we can begin to sample from our random generators. In practice, we will accomplish this by *interpreting* holey generators into standard QuickCheck generators in QuickCheck’s `Gen` monad and then sample from those. The interpretation function recursively from `Holey a` into `Gen a` is a straightforward translation of the intuitive semantics of holey generators: it iteratively fills holes until done.

```
recursively :: HoleWeighting -> Holey a -> Gen a
recursively f p = sized $ go p 0
  where
    go r _ _ | isDone r = return (done r)
    go r n target | n == target = return (done r)
    go r n target = do
      i <- frequency (second pure <$> f (treeOfHoles r))
      go (fill r i) (n + 1) target
```

The auxiliary function `go :: Holey a -> Int -> Int -> Gen a` defines a single iteration of a loop that will run until either all of the holes have been filled (the first guard), or the structure has reached






Fig. 4. Frequency of each of the size-4 trees, ordered by increasing depth, in a draw of 10,000 trees from a depth-weighted distribution (left) and an inverse-depth-weighted distribution (right). X-axis labels are trees “serialized” by a preorder traversal (N = Node, L = Leaf).

the desired size (the second). The “body” of the loop passes the tree of holes to the user-specified hole-weighting function, then samples a hole from it. This hole is filled, and the loop proceeds.

**Controlling the Distribution.** Of course, we are still waiting to answer the most important question: How does one choose the `HoleWeighting` function? Ideally, one chooses it to ensure that the important parts of the input space are adequately explored. If, for example, the programmer suspects that bugs may be revealed by input trees that are long and stringy, they could assign weights to the holes that are exponentially growing with the depth of each hole. This way, in a partially completed tree, the deeper-down holes in the tree are exponentially more likely to be filled, leading to a “chain reaction,” where deep trees beget deeper trees.

```
depthWeighted :: HoleWeighting
depthWeighted t = (\h -> (f h, h)) <$> holes t
  where
    f h = 4 ^ holeDepth h
```

In the left part of Figure 4, we present a graph showing all of the trees of size 4 and their relative frequencies in a draw of 10,000 trees from a holey generator using `depthWeighted`. From left to right, the trees are ordered by increasing depth.

Conversely, suppose the programmer believes that a bug will make itself known if tested against trees that are short and squat? They could consider weighting deeper holes lower, as in `inverseDepthWeighted` below.

```
inverseDepthWeighted :: HoleWeighting
inverseDepthWeighted t = (\h -> (f h, h)) <$> hs
  where
    hs = holes t
    maxDepth = maximum (holeDepth <$> hs)
    f h = 4 ^ (maxDepth - holeDepth h)
```

This hole weighting gives holes exponentially more weight the closer they are to the root (compared to the current deepest hole). The graph in the right part of Figure 4 shows the opposite story from the graph in Figure 4, with shorter, squatter trees now much more likely.

What if the programmer believes that the bug will rear its head on inputs that are severely left or right skewed? In this case, they could heavily weight holes which are left-leaning; the weighting function below operationalizes this by weighting holes exponentially based on how many “left



393

394

395

396  
397  
398  
399  
400  
401

402  
403  
404  
405  
406  
407

408  
409  
410  
411  
412  
413  
414  
415  
416  
417  
418

419

420

421

423

425  
426  
427  
428  
429  
430

## 120

433  
434  
435  
436  
437  
438  
439  
440

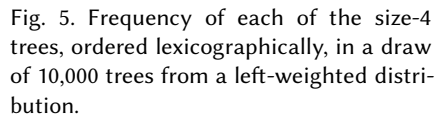


Fig. 5. Frequency of each of the size-4 trees, ordered lexicographically, in a draw of 10,000 trees from a left-weighted distribution.




Fig. 6. Frequency Charts,  $k = 4$ ,  $n = 10000$ . Left: unweighted holey generator. Right: classic QuickCheck Generator

For the case of binary trees, this problem is tractable. The key insight is to think of choosing the next hole as a *random walk* down the hole tree: a process that starts at the root of the tree and makes independently random choices to “go left” or “go right” until it reaches a leaf. The weight (or probability) of each hole in  $\text{uniform}_h$  will be the probability of that random walk ends up at that hole: the product of the probabilities of the choices it made along the way. The challenge of constructing  $\text{uniform}$  then reduces to the challenge of setting the random walk probabilities in such a way that the resulting hole weighting induces the uniform distribution on trees of each size.

Somewhat surprisingly, we can derive an efficiently computable solution for these random walk probabilities, where the probability of taking a left or right turn at a specific node during the walk depends only on the *sizes* of the left and right subtrees rooted at that node.

We also impose the helpful invariant that the generation process will be uniform “at every step.” This need not be the case—if you know that you intend to produce a uniformly random tree of size  $n$  by repeated node insertion, there is no reason a priori to insist that the terminating this process after  $k$  steps yield a uniformly random tree of size  $k$ . However, we will adopt this invariant as it greatly reduces the difficulty of computing the probabilities.

**Calculating the Weights.** To find the right random walk probabilities, let’s examine what happens when we add a node to a partially built tree by filling a hole chosen by the random walk. By determining what must be true of the walk probabilities in order for the outcome of this addition to be a uniformly chosen tree, we will derive constraints on the probabilities that can be turned into a method for computing them.

Suppose we have so far generated a tree  $t$  of size  $n$ , and that we’ve chosen the next hole to fill by taking a random walk down the tree. Let  $t'$  be the new tree (of size  $n + 1$ ) after filling this hole with a node. What is the probability that this hole is on the left side of the tree (i.e., that the first step in the random walk was to the left)? Formally, we would like to find  $P_n(d \mid l, r)$ : the probability that, when our walk encounters a tree of size  $n$  with left subtree of size  $l$  and right subtree  $r$ , it turns in the direction  $d \in \{L, R\}$ . These probabilities are sufficient to define the random walk, and hence the probabilities of filling each hole in every possible tree. To fix some more notation, let  $P_n(l, r)$  be the overall probability of generating a tree of size  $n$  with left and right subtrees of size  $l$  and  $r$ , and let  $P_n(l, r, d)$  be the probability of generating a tree of size  $n$  with a left subtree of size  $l$  and a right subtree of size  $r$  and then taking the first step in the walk down the tree to a hole in direction  $d$ .

For the moment, let’s suppose that both the left and right subtrees of  $t'$  are nonempty—we will return to the edge cases later. Let  $1 \leq k \leq n - 1$  be the size of the left subtree of  $t'$  (and  $n - k$  the

size of the right subtree). Given this setup, there are only two possibilities for the sizes of the two immediate subtrees of  $t$ : either the hole was chosen on the left, and it has a left subtree of size  $k - 1$  and right subtree of size  $n - k$ , or the hole was chosen on the right and it has a left subtree of size  $k$  and a right subtree of size  $n - k - 1$ . These two options are shown in diagram form in Figure 7.

This realization about the specific trees  $t$  and  $t'$  gives an important insight about the random walk probabilities. Since there are only two possible ways that we could have arrived at a tree with a left/right split like  $t'$ , namely  $(k, n - k)$ , the probability of generating a tree with that split has to be equal to the sum of the probabilities of (1) generating a tree of size  $n$  with left/right split  $(k - 1, n - k)$  and then the random walk going left, and (2) generating a tree of size  $n$  with left/right split  $(k, n - k - 1)$  and then the random walk going right. Symbolically, we have:




Fig. 7. The ways of getting to a binary tree of size  $n + 1$  with  $k$  nodes on the left, and  $n - k$  nodes on the right.

$$P_n(k - 1, n - k, L) + P_n(k, n - k - 1, R) = P_{n+1}(k, n - k) \quad (1)$$

But of course, the choices made during shape generation are all independent, and so the probability of encountering a tree with a particular left/right balance and then going left is the product of the probabilities of seeing such a tree, times the probability of going left at such a tree. So, we can re-write equation (1) as:

$$P_n(k - 1, n - k)P_n(L \mid k - 1, n - k) + P_n(k, n - k - 1)P_n(R \mid k, n - k - 1) = P_{n+1}(k, n - k) \quad (2)$$

But some of these probabilities are knowable. By the uniform-at-every-step assumption, the probability of encountering an  $n$ -node tree with left/right subtrees of size  $l$  and  $r$  is precisely the fraction of  $n$  node trees with subtrees of those sizes:

$$P_n(l, r) = \frac{C_l C_r}{C_n}$$

Moreover, the probabilities of going left and right are complements, since the walk always goes somewhere with probability 1:

$$P_n(L \mid l, r) + P_n(R \mid l, r) = 1$$

Using these facts we can simplify equation (2) to:

$$\frac{C_{k-1}C_{n-k}}{C_n}P_n(L \mid k - 1, n - k) + \frac{C_k C_{n-k-1}}{C_n}(1 - P_n(L \mid k, n - k - 1)) = \frac{C_k C_{n-k}}{C_{n+1}} \quad (3)$$

To simplify notation somewhat, we define  $P_n(k)$  to be  $P_n(L \mid k, n - k - 1)$ . Rewriting (3), we can plainly see that this equation defines a recurrence relation on  $P_n(k)$ , for  $1 \leq k < n - 1$

$$\frac{C_{k-1}C_{n-k}}{C_n}P_n(k - 1) + \frac{C_k C_{n-k-1}}{C_n}(1 - P_n(k)) = \frac{C_k C_{n-k}}{C_{n+1}} \quad (\star)$$

Given a base case for this recurrence, we can solve it and find the values of  $P_n(k)$ , as desired. The base case for this recurrence is derived from the two cases we ignored in our original analysis: when the left or right subtree of the tree  $t'$  that resulted after the node addition were empty, or in other words, in the cases  $k = 0$  and  $k = n$ . By an analysis similar to the one that brought us to the above solution, we find that the following two equations must hold on the boundary:

$$\frac{C_0 C_{n-1}}{C_n}(1 - P_n(0)) = \frac{C_0 C_n}{C_{n+1}} \quad \frac{C_{n-1} C_0}{C_n}P_n(n - 1) = \frac{C_n C_0}{C_{n+1}}$$

Intuitively, these equations hold because there is exactly one way to reach a tree with left/right split  $(0, n)$  or  $(n, 0)$ : by having a tree with split  $(0, n - 1)$ , and  $(n - 1, 0)$ , and going to the right and left, respectively. Solving, this yields the base case:

$$P_n(0) = 1 - \frac{C_n^2}{C_{n-1}C_{n+1}} \quad (\star\star)$$

Using  $(\star)$  and  $(\star\star)$ , we can easily compute  $P_n(k)$  for  $0 \leq k \leq n - 1$ . Through some serious algebraic simplification using the combinatorial fact  $\frac{C_n}{C_{n+1}} = \frac{n+2}{2(2n+1)}$ , we arrive at the equations:

$$P_n(0) = \frac{3}{(n+1)(2n+1)} \quad P_n(k) = \frac{2n-2k-1}{n-k+1} \left( \frac{n+2}{2n+1} - P_n(k-1) \frac{k+1}{2k-1} \right)$$

**Correctness.** While the above discussion hopefully conveys intuition, it does not constitute a proof. To derive the equations above, we chose a generation scheme—pick holes to fill via a random walk down the graph—and inferred constraints based on what must be true for iterating that process give a uniform distribution over trees of every size. But this by no means proves that using a hole-weighting function that picks holes by a random walk using the probabilities  $P_n(k)$  *must* induce the uniform distribution! To prove this, we need to be a bit more formal about our calculations with probabilities.

We begin by defining a random function called `add`. When given a tree, `add` takes a random walk down it by making independent left/right choices with our probabilities  $P_n(k)$ , and inserts a `Node` with `Children` children in place of the leaf at the bottom of its path. Formally, we define

$$\text{add}(\text{Leaf}) = \text{Node } \text{Leaf } \text{Leaf}$$

$$\text{add}(\text{Node } l \ r) = \begin{cases} \text{Node } \text{add}(l) \ r & \text{with probability } P_n(|l|) \\ \text{Node } l \ \text{add}(r) & \text{with probability } 1 - P_n(|l|) \end{cases}$$

We then define the formal analogue of our Holey uniform generator as a sequence of random variables  $T_n$  defined<sup>3</sup> by iterating the `add` function.

$$T_0 = \delta_{\text{Leaf}} \quad T_{n+1} = \text{add}(T_n)$$

By a routine induction we can see that the `add` function always sends trees of size  $n$  to trees of size  $n + 1$ , and so the trees  $T_n$  have support in the set of trees of size  $n$ , which we denote  $\text{Tree}_n$ .

Before prove uniformity, we need a lemma about the way the `add` function “balances” probabilities between different trees. The proof can be found in the Appendix.

LEMMA 3.1. *For all  $n \geq 0$  and all  $t$  of size  $n + 1$ , we have  $\sum_{t' \in \text{Tree}_n} \mathbb{P}(\text{add}(t') = t) = \frac{C_n}{C_{n+1}}$*

The lemma must hold in order to show that, in aggregate, the `add` function behaves  $\mathbb{P}(\text{add}(t') = t) = \frac{1}{C_{n+1}}$ : that every tree  $t'$  of size  $n$  has an equal chance of getting to any tree of size  $n + 1$ .

Now we can prove that the random variables  $T_n$  are in fact uniformly distributed over the trees of size  $n$ :

THEOREM 3.2. *For all  $n \geq 0$  and all  $t$  of size  $n$ ,  $\mathbb{P}(T_n=t) = \frac{1}{C_n}$*

PROOF. The proof proceeds by induction on  $n$ . The base case is trivial, since there is only one tree of size 0, namely `Leaf`. For the inductive step, let  $t$  be a tree of size  $n + 1$ . Then, unrolling definitions,

$$\mathbb{P}(T_{n+1}=t) = \mathbb{P}(\text{add}(T_n)=t)$$

<sup>3</sup> $\delta_{\text{Leaf}}$  is the “delta” random variable with law  $P(\delta_{\text{Leaf}} = \text{Leaf}) = 1$

Because the events  $\{T_n = t'\}_{t' \in \text{Tree}_n}$  are disjoint, we have that

$$\mathbb{P}(\text{add}(T_n)=t) = \sum_{t' \in \text{Tree}_n} \mathbb{P}(T_n=t', \text{add}(t')=t)$$

Because the randomness in  $T_n$  and the add function is independent, we have

$$\sum_{t' \in \text{Tree}_n} \mathbb{P}(T_n=t', \text{add}(t')=t) = \sum_{t' \in \text{Tree}_n} \mathbb{P}(T_n=t') \cdot \mathbb{P}(\text{add}(t')=t)$$

By the induction hypothesis,  $\mathbb{P}(T_n=t') = \frac{1}{C_n}$ , and so

$$\begin{aligned} \sum_{t' \in \text{Tree}_n} \mathbb{P}(T_n=t') \cdot \mathbb{P}(\text{add}(t')=t) &= \sum_{t' \in \text{Tree}_n} \frac{1}{C_n} \cdot \mathbb{P}(\text{add}(t')=t) \\ &= \frac{1}{C_n} \sum_{t' \in \text{Tree}_n} \mathbb{P}(\text{add}(t')=t). \end{aligned}$$

But by Lemma 3.1,  $\sum_{t' \in \text{Tree}_n} \mathbb{P}(\text{add}(t')=t) = \frac{C_n}{C_{n+1}}$ , and so we arrive at:

$$\frac{1}{C_n} \sum_{t' \in \text{Tree}_n} \mathbb{P}(\text{add}(t')=t) = \frac{1}{C_n} \frac{C_n}{C_{n+1}} = \frac{1}{C_{n+1}}$$

Stringing these equalities all together, we obtain  $\mathbb{P}(T_{n+1}=t) = \frac{1}{C_{n+1}}$ , as desired.  $\square$

**Implementation.** Now that we know it's correct, there are a few things to note about our solution to the uniform generation problem for binary trees. The first is that, in principle, these numbers need only be computed once! While there are  $O(n^2)$  probabilities required to generate a tree of size  $n$ , the same random walk probabilities can be used to generate any kind of binary tree of any size less than or equal to  $n$ , forever. This is incredibly convenient for using this distribution in practice: the combinatorics only need to be done once, and then can be re-used for any number of generation runs, in any number of tests, for any binary tree data type.

The second thing to remark is that the assumptions that we made in deriving these probabilities—uniform at every step, and that holes are chosen by a random walk from root to leaf—are not only simplifying assumptions that helped us derive the solution, but they also lend themselves nicely to a simple implementation of the `HoleWeighting` corresponding to the uniform distribution. Given a program to compute the probabilities  $P_n(k)$  (Figure 8), we use the fact that each step in the random walks are independent to compute the probability of a given hole being filled by multiplying out the probabilities of walking down the path to that particular hole. We can then use these hole probabilities to compute the `HoleWeighting` function for the uniform distribution by simply enumerating the holes in a tree and computing their probabilities.

```
calcProbs :: HTree -> Hole -> Rational
calcProbs HoleLeaf Here = 1
calcProbs t@(HNode l _) (L h) = (leftProbs (size t) !! (size l)) * (calcProbs l h)
calcProbs t@(HNode l r) (R h) = (1 - ((leftProbs (size t) !! (size l)))) * (calcProbs r h)
calcProbs DoneLeaf _ = error "Impossible: can't walk down to a done."
calcProbs _ _ = undefined
```

```
uniform :: HoleWeighting
uniform t = let hs = holes t in zip (weightify $ map (calcProbs t) hs) hs
```

With the `HoleWeighting` in hand, we can plug it into our `Holey` generator for the `UTree` type, and read off uniformly-at-random trees to our heart's content.

```

638 leftProbs :: Int -> [Ratio Integer]
639 leftProbs n = take n $ snd <$> iterate go (1,p0)
640 where
641   n' = toInteger n
642   p0 = 3 % ((n' + 1) * (2 * n' + 1))
643   go (k,pk_pred) =
644     let k' = toInteger k in
645     let a = (2 * n' - 2 * k' - 1) % (n' - k' + 1) in
646     let b = (n' + 2) % (2 * n' + 1) in
647     let c = (k' + 1) % (2 * k' - 1) in
648     let pk = 1 - a * (b - c * pk_pred) in
649     (k+1,pk)

```

Fig. 8. Code to compute the random walk probabilities.




Fig. 9. Frequencies of size-4 trees in a draw of 10,000 from the uniform distribution. Compare to Figure 6.

Figure 9 shows the frequencies of all size-4 trees over a draw of 10,000 from the generator recursively `genUTree uniform`. In expectation, one should find  $y = \frac{10,000}{C_4} \approx 714$  of each of the  $C_4 = 14$  trees. Of course, estimated probabilities will vary between draws, but we can see that all of the empirical frequencies deviate only slightly from this line, which is drawn in red through the top of the chart.

It's worth stepping back to consider what we've demonstrated here. Global control of the choices made during generation has given us complete access to shape the underlying distribution. This mechanism is powerful enough to encode combinatorially complex distributions like the uniform distribution over fixed sizes, simply by weighting which holes to fill.

#### 4 GENERATING HOLEY GENERATORS

We've established that `Holey` generators give powerful control over distributions, but they can't yet generate interesting structures—so far, they have really only been able to produce unlabeled trees. In this section, we show how to *generate holey generators*, reintroducing the monadic power of classic QuickCheck generators without giving up the control we fought so hard to obtain.

**Decorating the Trees.** The `holeyUTree` generator from Section 2 is rather bare. Indeed, it is boringly self-similar, defining an infinite tree of generators that all generate simple `UNodes`. What we need is

a way to *decorate* that tree with data. We can attempt to do this manually in a couple of different ways:

```

data Tree a = Leaf | Node (Tree a) a (Tree a) deriving (Eq, Ord, Show)

badGenHoleyTree1 :: Holey (Tree Int)
badGenHoleyTree1 = Leaf `orFill` (Node <$> badGenHoleyTree1 <*> pure 0 <*> badGenHoleyTree1)

badGenHoleyTree2 :: Holey (Tree Int)
badGenHoleyTree2 =
  Leaf `orFill`
    (Node <$> (Node <$> leaf <*> pure 1 <*> leaf)
      <*> pure 3
      <*> (Node <$> leaf <*> pure 5 <*> leaf))
  where
    leaf = pure Leaf

```

The first generator technically produces a labeled tree, but all of the labels are 0, so it's not good for much. The second is a bit more interesting, but it bottoms out after a few levels because every label needs to be written in manually. However, there's a good idea hiding here: if we had a way of building holey generators with labels already embedded inside, we could sample shapes from those trees without an issue!

To further illustrate this point, consider this holey generator that produces trees whose labels increase with depth:

```

incHoleyTree :: Int -> Holey (Tree Int)
incHoleyTree k =
  Leaf `orFill` (Node <$> incHoleyTree (k + 1) <*> pure k <*> incHoleyTree (k + 1))

```

Sampling from `incHoleyTree 1 >>=` recursively uniform with size 3 would produce trees like

```

Node Leaf 1 (Node (Node Leaf 3 Leaf) 2 Leaf)    and
Node (Node Leaf 2 Leaf) 1 (Node Leaf 2 Leaf)

```

with a nice uniform distribution over their shapes! The key here is that `incHoleyTree 1` represents an infinite tree of *hypothetical* nodes in a `Holey` structure that will eventually produce a tree via the hole-filling procedure from Section 2. Each node is hypothetical because recursively might choose to not expand that hole, leaving only a `Leaf`. But if the procedure *does* expand the node, the label that will appear in the node has already been chosen. This process is illustrated in Figure 10. All we need now is to replace these silly deterministic algorithms with random ones.




Fig. 10. A single binary tree, chosen from an infinite Holey structure.

Of course, we know how to do this—we can just use classic QuickCheck! This program in QuickCheck's `Gen` monad randomly generates a Holey generator:

```

genHoleyTree :: Gen (Holey (Tree Int))
genHoleyTree = do
  x <- arbitrary
  l <- genHoleyTree
  r <- genHoleyTree
  return (Leaf `orFill` (Node <$> l <*> pure x <*> r))

```

This function works just like the deterministic algorithm above, but it uses monadic sequencing to thread randomness through the program. This generator works just like `incHoleyTree`, but instead of incrementing the decorated labels down the tree it samples a new one each time using `arbitrary`. Every random seed gives a different holey generator with a different set of labels, and shapes can be sampled from that tree as desired.



Now we can use `genHoleyTree` to get a distribution over labeled trees that is uniform over their shapes, we first run `genHoleyTree` and then we fill holes with `recursively`:

```
genTree :: Gen (Tree Int)
genTree = genHoleyTree >=> recursively uniform
```

**What have we gained?** Before continuing, let's reflect on the many benefits of layering `Gen` and `Holey`.

*A Cleanly Staged Interface.* The type `Gen (Holey a)` is really a staged approach to generation. The first stage of generation controls the labels that might be in the tree, and the second controls the shapes of the actual trees generated. This is quite a natural interface to work with. Since these generators ultimately live in the usual `Gen` monad, they are entirely compatible with existing QuickCheck generators. Furthermore, by separating the random generation of labels from the recursive expansion procedure, it is relatively difficult to “get it wrong.” Structure is handled by the `Holey` abstraction, labels are handled by `Gen`, and the two are only interleaved at the last moment. Additionally, staging generation like this can lead to efficiency improvements. If a variety of shapes is deemed more important than a variety of labels, the tester can sample a single `Holey a` and use it to sample many different shapes of trees.

*Distributional Control.* This was illustrated above—it is easy to see that the distribution of decorated trees is exactly the same as the one that we got from `genHoleyUTree`. We'll need to relax this control a bit when preconditions come into play, but in many cases generating a `holey` generator gives the same level of control that was provided by the original abstraction discussed in Section 2.

*Size Control.* The way that generators like this are staged also makes size control a joy. A `holey` generator like `genHoleyTree` really has two sizes that one might care about. The first is the size of the labels in the tree—the range that they are chosen from—which is controlled by `sized` on the first line of the generator. The second is the size of the tree itself, which as we know is controlled by `Gen`'s size parameter when `recursively` is called. At first, it would seem that controlling both of these sizes with the same size parameter is a poor choice. Should we make one an explicit argument to the generator?

Actually, there is no need! Since generation of labels happens before generation of trees, we can leverage the `holey` generator's staging to stage the sizes too. Take a look at a modified version of `genBST`:

```
genBSTResize :: Gen (Tree Int)
genBSTResize = do
  g <- resize 30 genHoleyBST
  resize 5 (recursively uniform g)
```

We call QuickCheck's `resize` function twice to resize the two different aspects of BST size!

**Enforcing Preconditions.** Finally, we arrive at generators that enforce preconditions. The ones we present here, for BSTs and for binary heaps, are scarcely more complicated than `genHoleyTree`, but they are quite a bit more interesting and useful.

Our generator for BSTs follows the same strategy that testers use in classic QuickCheck, tracking the minimum and maximum labels allowed in each subtree:

```
genHoleyBST :: Gen (Holey (Tree Int))
genHoleyBST = sized (\n -> aux (-n, n))
  where
    aux (lo, hi) | lo > hi = return (pure Leaf)
    aux (lo, hi) = do
```

```

785     x <- choose (lo, hi)
786     l <- aux (lo, x - 1)
787     r <- aux (x + 1, hi)
788     return (Leaf `orFill` (Node <$> l <*> pure x <*> r))

```

The labels in this tree are constrained by the bounds that are passed from one recursive call to the next, and the tree is forced to “bottom out” with a `Leaf` if the range for the given subtree is empty. Similarly, this generator for binary heaps ensures ordering from the top to the bottom of the tree:

```

792 genHoleyHeap :: Gen (Holey (Tree Int))
793 genHoleyHeap = sized aux
794   where
795     aux hi | hi <= 0 = return (pure Leaf)
796     aux hi = do
797       x <- choose (0, hi)
798       l <- aux x
799       r <- aux x
800       return (Leaf `orFill` (Node <$> l <*> pure x <*> r))

```

We can run these holey generator generators just like `genHoleyTree`, either with fancy staged sizing or by simply inheriting QuickCheck’s size control.

```

804 genBST :: Gen (Tree Int)
805 genBST = genHoleyBST >=> recursively uniform

```

The distribution of this generator, shown in Figure 11 has a great variety of shapes, but it isn’t *exactly* uniform. This is a side-effect of the BST invariant that the generator is required to enforce. You see, while `genHoleyTree` produces infinite trees of hypothetical labels, `genHoleyBST` gives finite trees. This is obvious from a simple counting argument—BSTs do not allow duplicate elements and the initial range is finite, therefore the hypothetical tree must also be finite—but the mechanics of exactly what happens is interesting. We’ll use Figure 12 as an example. If we assume that  $lo = 1$  (i.e., we are in a subtree where 1 is the smallest valid label), then there are no valid labels to the left of 1 nor are there any to the right of 3. Accordingly, the branches of the Holey generator are cut off at those points, meaning that those parts of the tree will never be expanded to a Node.

What does any of this have to do with uniformity? Well, in the pathological case where the hypothetical tree has  $n$  labels and we want a tree of size  $n$ , there is only one tree to choose! So clearly there is no way to get a uniform distribution there. As the number of nodes in the tree increases things get much better, but recall from Section 3 that our uniformity argument still assumes that any node can be expanded to yield a new node with two children (this is how we get to  $C_n$  as the number of total trees of size  $n$ ). If that is not true for some nodes, then the uniformity argument cannot continue to hold in general.

This sounds unfortunate, but it is not the end of the world. If the range of labels allowed in the tree is sufficiently large, then the




Fig. 11. Frequency of each of the size-4 shapes of BSTs in a draw of 10,000 using the uniform hole weighting.




Fig. 12. A Holey generator for BSTs, pruned to enforce the invariant (assumes  $lo = 1$ ).

tree will be as good as infinite from the perspective of hole-filling. After all, filling  $n$  holes can't possibly explore more than  $n$  nodes deep in the hypothetical tree.

Additionally, while properties like uniformity suffer a bit from the addition of precondition constraints, your size almost certainly does not. This is because of the way hole-filling works: recursively tries to fill the appropriate number of holes no matter what, even if those holes are in unexpected places. If one part of the tree gets stuck due to a constraint, the rest of the size is allocated to another part of the tree without issue. The only potential problem arises if there are simply not enough nodes available in the hypothetical `Holey` tree, but this is rare. In our experiments, we found that as long as the range of potential labels is at least  $\frac{3}{2}$  the total number of desired nodes, the tree will be large enough to avoid worries. Since such a tight bound would lead to distributional issues, there is no reason to cut it that close.

We've successfully extended the `Holey` framework to generate the kinds of binary trees with invariants that occur in actual testing scenarios. These generators give the user fine control over their size and shape distributions, but does any of this actually help with testing? Yes, in fact it does, as we shall see next.

## 5 PUTTING IT INTO PRACTICE

To find out how `holey` generators perform in a realistic testing scenario, we needed a case study with many properties to experiment with. We found one in *How to Specify It!* [11], a tutorial on property-based testing that uses binary search trees as its running example, with all the code available online [10]. In this case, the trees represent finite maps, so contain both keys and values; the implementation itself is small but not trivial, consisting of 62 non-blank non-comment lines of Haskell code, with the following API:

```
find  :: Ord k => k ->      BST k v -> Maybe v
nil   ::                      BST k v
insert :: Ord k => k -> v -> BST k v -> BST k v
delete :: Ord k => k      -> BST k v -> BST k v
union  :: Ord k => BST k v -> BST k v -> BST k v
```

**A Case Study with a Wealth of Properties.** The code is accompanied by a test-suite of 49 properties from the paper—testing that each operation preserves the invariant, that expected postconditions hold, that various algebraic laws are satisfied, that behavior is consistent with a 'reference implementation' using lists of key-value pairs, and so on. The repository also includes 31 properties generated by QuickSpec [21], of which 11 duplicate properties presented in the paper. Seven properties in the paper quantify over pairs of equivalent trees (containing the same keys and values), which require a special generator<sup>4</sup>; we discard these and all duplicate properties, resulting in 57 different properties that we can use to evaluate binary tree generators.

The repository also contains eight buggy versions of the implementation, with bugs ranging from blatant (`insert` discards the tree it is passed, always returning a tree with a single node) to more subtle (when taking a `union` of trees with some keys in common, the specification says that values from the left argument should be preferred, but the buggy implementation does not do so consistently). We evaluate generators by measuring the *average number of tests needed to provoke failure*, for every failing property. In each case we took this average over 1,000 failing runs of QuickCheck, which translates into up to 200,000 tests per property (to find 1,000 counterexamples).

<sup>4</sup>Since two independently generated trees are highly unlikely to be equivalent, then we must generate such pairs of trees together, for example by generating one of the pair, extracting its keys and values, and then generating the other tree by choosing one of those keys to be at the root, splitting the keys and values into those for each subtree, and then generating the subtrees by the same method.

We ran a very large number of tests of *non-failing* properties, so it is very unlikely that we missed a potential failure.

**Three Approaches to Generation.** In all our tests, we generated *keys* in the range 0 to *size*, where *size* is the QuickCheck size parameter. Thus there are *size* + 1 possible keys, and when *size* is small, two independently generated keys are quite likely to be equal, while, when *size* is large, this is unlikely.

*API-based Generation.* The *How to Specify It!* experiment originally used the following tree generator, which inserts a random list of keys and values into the empty tree:

```
instance Arbitrary Tree where
  arbitrary = do kvs <- arbitrary :: Gen [(Key,Val)]
               return $ foldr (uncurry insert) nil kvs
```

We call this an ‘API-based’ generator, because it uses the API under test to generate test cases; this kind of generator is simple to write, and has the merit that generated trees are guaranteed to satisfy the invariant—provided *insert* is correct. There is a risk that testing may be incomplete, though, if some trees cannot be built at all using *insert* alone.

Using this generator we found 122 failing bug/property combinations; every bug provoked many properties to fail, ranging from 9 to 23 depending on the bug. Every bug can be found with fewer than 10 random tests *using the right property*—but this is a bit like saying that a bug is easy to provoke once you know the right test to run; it’s true, but not useful. It is not obvious in advance which property will be most effective at finding each bug.

Not all failures are found fast, though. To give us an idea of how well the generator worked overall, we computed the *total* number of tests needed (on average) to falsify *all* failing bug/property combinations (see Figure 13). We also looked at the *hardest* properties to falsify: the three hardest all involved a buggy *delete*, and they were (in order):

```
prop_DeleteDelete k k' t = delete k (delete k' t) == delete k' (delete k t)
prop_DeletePost k t k'   = find k' (delete k t) == if k==k' then Nothing else find k' t
prop_qs_24 k t t'       = delete k (union (delete k t) t') == delete k (union t t')
```

The very hardest of these, *propDeleteDelete*, required over 190 random tests to find a counterexample, on average.

*Classic Generation.* Next we replaced the generator with one in the classic style, very like *genBST* on page 2, except that we generate keys in the range 0 . . . *size*, and of course we also generate a random value in each node. As on page 2, we control tree size by halving the size bound in each recursion, and we favor branches over leaves by a ratio of 5-to-1.

The first surprise was that we found four more failures! Recall that the most blatant bug causes *insert* to return a single-node tree, no matter its arguments. When this version of *insert* is used in the API-based generator, it causes every generated tree to consist of only zero or one nodes, so properties that only fail for larger trees cannot be falsified. This is a graphic illustration of the risks of API-based generators.

Method	Total cost	Worst property
Original API-based	1607	193
Classic	1345	135
Holey	1075	103
Tuned API-based	1260	155
Tuned Classic	1006	96

Fig. 13. The holey generator competes with the others in terms of average number of tests needed to provoke all failures, and the failure of the most difficult property. Smaller is better.

As well as finding more failures, the classic generator did so at lower total cost (Figure 13); the same three properties were hardest to falsify, but fewer tests were needed to do so.

**Holey Generation.** Finally, we implemented a holey version of the generator. First we wrote an auxiliary

```
holeyTree :: Int -> Int -> Gen (Holey Tree)
```

very like `aux` in the definition of `genHoleyBST` on page 17; `holeyTree lo hi` chooses keys randomly (consistently with the invariant), and returns a `Holey Tree` that can be filled to obtain a tree of any size from 0 to `hi-lo+1` (because there are `hi-lo+1` possible keys). It remains to decide what size of tree to generate.

It might be tempting to generate trees with exactly *size* nodes, but this would be a mistake, because then properties with *two* trees as arguments would only be tested with arguments of the same size. Instead it's customary to treat the size parameter as a *bound*, so that the values generated at *size* + 1 are a superset of those generated at *size*. We achieve this by choosing the tree size uniformly in the range 0 . . . *size* + 1:

```
instance Arbitrary Tree where
  arbitrary = sized $ \n -> do
    ht <- holeyTree 0 n
    treeSize <- choose (0,n+1) -- OBS! n+1 different keys
    resize treeSize (recursively uniform ht)
```

This results in a nice variety of test data, including small trees with keys drawn from a small set, small trees with keys drawn from a large set, and even trees containing every possible key.

Repeating our measurements, we found that the holey generator found all the failures at an even lower total cost, and falsified the hardest property (still the same one) in little more than half the number of tests we started with (Figure 13). These improvements are useful, if not dramatic. The second- and third-hardest properties to falsify are actually different in this case: they test the same bug in `delete` as before.<sup>5</sup>

```
prop_qs_27 t k t' = union t (delete k (union t t')) ==== union t (delete k t')
prop_qs_30 k t t' = union (delete k (union t t')) t ==== union t (delete k t')
```

**Tuning.** One may wonder *why* this bug in `delete` is harder to find, and inspecting a generated counterexample provides a clue:

```
*BSTSpec> quickCheck prop_DeleteDelete
*** Failed! Falsified (after 2 tests):
Key 0
Key 1
Branch (Branch Leaf (Key 0) 0 Leaf) (Key 1) (-1) Leaf
[] /= [(Key 0,0)]
```

Recall that `prop_DeleteDelete` tests that deleting keys in either order yields the same result; in this case, because a comparison is the wrong way around, the buggy `delete` only works at the root of the tree. As a result, if we try to delete `Key 0` first, then it is not deleted. To falsify the property, we need to choose two different keys, both present in the tree, with one of them located at the root.

What is the probability that a randomly chosen key is present in a generated tree? We can measure this and find that, for the holey generator, it is 50%. This makes sense: we chose the tree size uniformly in the range 0 to the maximum, so on average a generated tree contains half the possible keys. But for the API-based generator, the probability is 36%, and for the classic generator

<sup>5</sup>These properties are not equivalent since `union` is not commutative.




Fig. 14. The “tuned classic” generator: performance vs “holey” on the hardest properties, and distribution of generated tree sizes.

it is only 29%. Thus these generators may perform worse for tests involving `delete`, simply because they test deletion of a key that is not present in the tree more of the time.

With this in mind, we can *tune* the other two generators to generate larger trees, so that a random key is more likely to be present. Since tests in which a key *is* present in a tree, and *is not* present in a tree, are arguably equally important, we aim for a 50% presence-probability.<sup>6</sup> For the API-based generator, one might have expected it to generate trees containing half the keys already, since it inserts a random list of keys and values, and QuickCheck generates lists with a length chosen uniformly from  $0 \dots \text{size}$ . However, since the keys in the list are chosen independently, then some are duplicates, and do not contribute to the final tree size. We can tune this generator by increasing the lengths of the lists, and we discovered empirically that lists  $\frac{5}{3}$  longer resulted in trees containing half the keys on average.

The classic generator is a bit harder to tune. We can (a) reduce the size bound by a smaller factor in each recursion, instead of dividing by 2, and (b) weight branch nodes more highly. We discovered empirically that we needed to do both: we removed the size bound altogether, and weighted branch nodes  $7\times$  higher than leaves, and by so doing achieved a membership-probability of 48.4%, which was the best we could do.

Tuning did result in better performance (Figure 13), and indeed the tuned classic generator exhibits a slightly lower total cost, and a slightly better worst-property cost, than the holey generator. However, the picture is not clear-cut: if we look at the five hardest properties to falsify (now the same ones for both generators), then we see that while tuning has much improved the classic generator’s performance for the two previously-hardest properties, it has also *worsened* it significantly for the three next hardest. Moreover, while the *average* tree size is now about right, the *distribution* of tree sizes is still skewed towards either empty or larger trees (Figure 14); this may explain the poorer performance on three of the harder properties, since in a larger tree, a random key that *is* present is less likely to be at the root.

**Lessons.** What have we learned? Firstly, that the natural way to write a holey generator resulted in better performance than the natural way to write either an API-based or classic generator. Secondly,

<sup>6</sup>The reader may feel it is less important to test `delete` often with keys that are not present in the tree, but consider `insert`: both inserting a *new* key, and updating the value associated with an *existing* key, are important cases. If you are tempted to suggest using different distributions to test `insert` and `delete`, be aware that some properties call both functions *with the same key*. Likewise, the reader may wonder whether we should generate *keys from trees*, instead of independently, by choosing one of the keys in the tree. But consider that some properties are parameterised on a key and *two* trees—which tree should one choose a key from in such cases? Customizing generation *for each property* leads to ad hoc decisions and complexity, and consequently risks errors which *reduce* the effectiveness of testing, and are hard to discover—this way lies madness! We prefer to focus on making generation work well for *all* the properties we write.



while it is *possible* to tune a classic generator to get comparable fault-finding power, it is quite painful to do so because the effect of tweaking each parameter is unpredictable; one is reduced to performing trial-and-error experiments to see if the goal is achieved. This process is fraught, and may be difficult for non-experts. By contrast, it is much easier to specify the desired outcome with holey generators, making them more accessible for average users.

## 6 RELATED WORK

Generator tuning and distributional control have been topics of interest to the testing community for some time now. Here are some good ideas that are related to ours.

*FEAT and Friends.* As discussed in the introduction, FEAT [5] is *almost* a solution to the problem of local distributional control. A FEAT enumerator can generate interesting structures with whatever size and shape distribution the programmer desires, but it cannot do so efficiently under complex semantic constraints. In the next section, we discuss potential ways to bring FEAT in the fold, combining it with the ideas we present here, but for now it is not quite what we're after.

However, FEAT is not the only tool that gives attractive, global distributional control. Rather than rely on a correct-by-construction generators, [2] use a backtracking scheme that narrows a generator's distribution, avoiding values that fall outside of an executable constraint. This technique is made surprisingly efficient with the help of Haskell's laziness, but fundamentally this approach can still run into trouble when constraints are particularly *sparse* (i.e., when few values satisfy them relative to the total number of values). In those cases, a manual generator, tuned using our techniques, is preferable.

*Derived Generators.* Distributional properties like size and shape often require manual tuning to get just right, but some properties can be enforced automatically with the help of some pre-computation. The DRAGEN tool uses metaprogramming to automatically derive generators for algebraic data types, which ensure that the constructors in the data type being generated appear at a predetermined rate [18]. For example, a DRAGEN generator for a tree with multiple `Node` constructors can be derived to ensure a particular ratio between `Node1` and `Node2`. A related tool gives similar control over any algebraic data type defined in the “à la carte” tradition [17]. These automatically derived generators are quite impressive, but as with FEAT and its ilk, they cannot express generators that enforce complex preconditions.

*Automatic Distributional Control.* While we focus on manually tuned generators in this discussion, there is a myriad of automatic approaches that manipulating a generator's distribution. Tools exist that guide test distributions using code coverage [12, 23], a variety of optimization functions [13], common examples [22], and even machine learning [8, 19]. Many of these approaches still ultimately rely on local tuning, which makes certain distributions more difficult to obtain than others, but in specific use-cases these tools are sometimes the best option for low-effort high-reward test generation.

## 7 CONCLUSIONS AND FUTURE WORK

What have we learned? In Section 1, we learned that the current state of the art for QuickCheck generators gives inadequate control to the user, even in the well-studied case of binary-tree data structures. In Section 2, we learned how to achieve better control in the setting of unlabeled binary trees with the help of `Holey` trees. We took a detour though the combinatorics of uniform generation in Section 3 to demonstrate the flexibility our abstraction. Our technical contributions concluded in Section 4, where we found that a staged approach—generating `Holey` generators—allows us to capture binary-tree-based data structures with invariants. Finally, in Section 5 we presented a case



study based on *How To Specify It!* that shows that holey generators compare favorably even with carefully-tuned classic ones.

To wrap up, let's look at where we might go from here.

**Alternative Abstractions.** We considered some other choices when defining our generator abstraction; while we think that the current presentation is best for current purposes, we plan to explore other options moving forward.

One interesting alternative is a different definition of `Holey`:

```
data Holey a where
  Pure :: a -> Holey a
  (:*) :: Holey (a -> b) -> Holey a -> Holey b
  OrFill :: a -> Holey a -> Holey a
```

This is essentially the *free applicative* [1] with an extra operation for delimiting recursion, and it is similar to existing approaches that use a free applicative structure to connect generators and parsers [9]. Amazingly, an argument due to Elliott [6] shows that, for our purposes, these definitions are essentially equivalent! Filling a hole in this symbolic version of `Holey` corresponds to a *symbolic derivative* and filling a hole in the state-machine version corresponds to an *automatic derivative*; moreover, the two operations are implemented with almost exactly the same code. We opted for the state-machine version presented here because it fit better with our presentation, but we are excited to work more with the syntactic version and explore further connections between generators and derivatives.

Another intriguing option is to eschew `Holey` entirely and generate FEAT generators instead—i.e., produce something like `genFeatBST :: Gen (Enumerate a)`. Indeed, it seems this should work: FEAT and `Holey` both suffer from the same applicative limitations, and wrapping either in `Gen` gives considerably more power. However, there seem to be a few reasons to prefer holey generators: they are both more efficient (FEAT may incur exponential cost when computing the size of a space) and more consistent (given certain constraints, FEAT will generate a tree that is too small, whereas `Holey` will simply choose to fill other holes). Still, FEAT does provide uniform generation of much more complex types than binary trees, so we plan to explore this avenue as a potential way to generalize our approach.

**Generalizing to More Complex Types.** More broadly, we are considering a number of strategies that might generalize the approach we discuss in this paper to work beyond binary trees.

First, it's important to understand why our current technique doesn't generalize well. As discussed in Section 2, to achieve global control, a `Holey` generator carefully maintains a *binary* `HTree` which matches the structure of the tree it is generating. Attempting to generate ternary-or-larger trees using this API will necessarily break the synchronization between the internal `HTree` and the tree being generated. If the data type being generated has a ternary constructor, the recursive generator expression `g = pure Leaf `orFill` Node `fmap` g <*> g <*> g` will internally yield unbalanced `HTrees` like `HNode (HNode HoleLeaf HoleLeaf) HoleLeaf`, as the `Holey` API attempts to interpret `<*>`s as binary nodes.

One potential approach could be to change the `HTree` type to mirror the constructors of the tree type being generated. This seems technically feasible—Template Haskell or GHC Generics [14, 20] come to mind as a way of constructing a new `Holey` type from an algebraic data type—but unappealingly ugly. More elegant might be to replace the `HTree` with a rose tree [16], but then the generator would have less static information to use during generation.

In either case, a larger problem is that it becomes difficult to know how to set hole weights. To see why, consider what it would take to achieve a uniform distribution like the one from Section 3 for general trees. Even for fixed-arity trees, the combinatorial logic we used wouldn't yield an

efficiently-computable formula, and for mixed arity trees, the “uniform at each stage” assumption breaks down entirely.

Although these generalizations seem challenging, there are a number of paths forward that seem promising. We hope to follow them, and find a unifying abstraction that gives the user fine-grained distributional control while generating all of their favorite data types.

## APPENDIX

LEMMA 3.1. *For all  $n \geq 0$  and all  $t$  of size  $n + 1$ , we have  $\sum_{t' \in \text{Tree}_n} \mathbb{P}(\text{add}(t') = t) = \frac{C_n}{C_{n+1}}$*

PROOF. By induction on  $t$ . The  $t = \text{Leaf}$  case is vacuous. Suppose  $n \geq 1$ . Let  $t = N(l, r)$ , and suppose that  $|l| = k$  (and so  $|r| = n - k$ .)

First, consider the case where  $1 \leq k \leq n - 1$ . If  $l$  and  $r$  both have nonzero size, then there are only two ways to arrive at  $t$  by adding a node to a tree  $t'$  of size  $n$ . Either  $t'$  has a left subtree of size  $k - 1$  (and add took the left branch), or  $t'$  has a right subtree of size  $n - k - 1$  (and add took the right branch). For any tree of size  $n$  with any other split of nodes in its left/right subtrees,  $P(\text{add}(t') = t) = 0$ . Therefore, denoting by  $T(a, b)$  the set of  $a + b + 1$ -node trees with left subtrees of size  $a$  and right subtrees of size  $b$ , we have:

$$\sum_{t' \in \text{Tree}_n} P(\text{add}(t') = t) = \sum_{t' \in T(k-1, n-k)} P(\text{add}(t') = t) + \sum_{t' \in T(k, n-k-1)} P(\text{add}(t') = t)$$

Re-writing  $t'$  inside the sums as  $N(l', r')$ , and the  $t$  as  $N(l, r)$ , the right hand side of the above is equal to:

$$\sum_{N(l', r') \in T(k-1, n-k)} P(\text{add}(N(l', r')) = N(l, r)) + \sum_{N(l', r') \in T(k, n-k-1)} P(\text{add}(N(l', r')) = N(l, r))$$

Again, most of these terms drop away. On the left side,  $P(\text{add}(N(l', r')) = N(l, r)) = 0$  if either  $r \neq r'$  or the add goes right: the only way to get to a tree in  $T(k, n - k)$  by adding a node to a tree in  $T(k - 1, n - k)$  is if you fill on the left, and the right subtrees were the same in the first place. A similar argument goes for the right hand side. Thus, we have:

$$\sum_{N(l', r') \in T(k-1, n-k)} P_n(k-1)P(\text{add}(l') = l)P(r = r') + \sum_{N(l', r') \in T(k, n-k-1)} (1 - P_n(k))P(\text{add}(r') = r)P(l = l')$$

Again, yet more terms drop out: the  $P(r = r')$  and  $P(l = l')$  are zero for  $r \neq r'$  and  $l \neq l'$ , and so we are left with

$$\sum_{l' \in \text{Tree}_{k-1}} P_n(k-1)P(\text{add}(l') = l) + \sum_{r' \in \text{Tree}_{n-k-1}} (1 - P_n(k))P(\text{add}(r') = r)$$

Pulling out the constants from both sides, we have

$$P_n(k-1) \sum_{l' \in \text{Tree}_{k-1}} P(\text{add}(l') = l) + (1 - P_n(k)) \sum_{r' \in \text{Tree}_{n-k-1}} P(\text{add}(r') = r)$$

But by the induction hypothesis,

$$\sum_{l' \in \text{Tree}_{k-1}} P(\text{add}(l') = l) = \frac{C_{k-1}}{C_k}$$

and

$$\sum_{r' \in \text{Tree}_{n-k-1}} P(\text{add}(r') = r) = \frac{C_{n-k-1}}{C_{n-k}}.$$

Substituting in, we have

$$\begin{aligned} P_n(k-1) \sum_{l' \in \text{Tree}_{k-1}} P(\text{add}(l') = l) + (1 - P_n(k)) \sum_{r' \in \text{Tree}_{n-k-1}} P(\text{add}(r') = r) \\ = P_n(k-1) \frac{C_{k-1}}{C_k} + (1 - P_n(k)) \frac{C_{n-k-1}}{C_{n-k}} \end{aligned}$$

But, we have picked the  $P_n(k)$  to satisfy  $(\star)$ , the equation from Section 3! Multiplying  $(\star)$  through by  $\frac{C_n}{C_k C_{n-k}}$  yields

$$P_n(k-1) \frac{C_{k-1}}{C_k} + (1 - P_n(k)) \frac{C_{n-k-1}}{C_{n-k}} = \frac{C_n}{C_{n+1}}$$

as required.

Now consider the case where  $k = 0$  (the  $k = n$  case is symmetric). If  $t$  has an empty left subtree and a right subtree of size  $n$ , then the only trees  $t'$  for which  $P(\text{add}(t') = t)$  are trees of the form  $N(\text{leaf}, r')$ , where  $|r'| = n - 1$ . Thus,

$$\sum_{t' \in \text{Tree}_n} P(\text{add}(t') = t) = \sum_{r' \in \text{Tree}_{n-1}} P(\text{add}(N(\text{leaf}, r')) = N(\text{leaf}, r))$$

But just like the last case,  $\text{add}(N(\text{leaf}, r')) = N(\text{leaf}, r)$  exactly when  $\text{add}$  goes right, and  $\text{add}(r') = r$ . So,

$$\begin{aligned} \sum_{r' \in \text{Tree}_{n-1}} P(\text{add}(N(\text{leaf}, r')) = N(\text{leaf}, r)) &= \sum_{r' \in \text{Tree}_{n-1}} (1 - P_n(0)) P(\text{add}(r') = r) \\ &= (1 - P_n(0)) \sum_{r' \in \text{Tree}_{n-1}} P(\text{add}(r') = r) \end{aligned}$$

but by the IH,

$$(1 - P_n(0)) \sum_{r' \in \text{Tree}_{n-1}} P(\text{add}(r') = r) = (1 - P_n(0)) \frac{C_{n-1}}{C_n}.$$

By  $(\star\star)$ ,  $1 - P_n(0) = \frac{C_n^2}{C_{n-1} C_{n+1}}$ , and so

$$(1 - P_n(0)) \frac{C_{n-1}}{C_n} = \frac{C_n^2}{C_{n-1} C_{n+1}} \frac{C_{n-1}}{C_n} = \frac{C_n}{C_{n+1}}$$

as desired.  $\square$

## REFERENCES

- [1] Paolo Capriotti and Ambrus Kaposi. 2014. Free applicative functors. *arXiv preprint arXiv:1403.0749* (2014).
- [2] Koen Claessen, Jonas Duregård, and Michal H. Palka. 2015. Generating constrained random data with uniform distribution. *J. Funct. Program.* 25 (2015). <https://doi.org/10.1017/S0956796815000143>
- [3] Koen Claessen and John Hughes. 2000. QuickCheck: a lightweight tool for random testing of Haskell programs. In *Proceedings of the Fifth ACM SIGPLAN International Conference on Functional Programming (ICFP '00), Montreal, Canada, September 18-21, 2000*, Martin Odersky and Philip Wadler (Eds.). ACM, Montreal, Canada, 268–279. <https://doi.org/10.1145/351240.351266>
- [4] Jonas Duregård, Patrik Jansson, and Meng Wang. 2012. Feat: Functional Enumeration of Algebraic Types. In *Proceedings of the 2012 Haskell Symposium (Copenhagen, Denmark) (Haskell '12)*. Association for Computing Machinery, New York, NY, USA, 61–72. <https://doi.org/10.1145/2364506.2364515>
- [5] Jonas Duregård, Patrik Jansson, and Meng Wang. 2012. Feat: Functional Enumeration of Algebraic Types. *SIGPLAN Not.* 47, 12 (sep 2012), 61–72. <https://doi.org/10.1145/2430532.2364515>
- [6] Conal Elliott. 2021. Symbolic and automatic differentiation of languages. *Proceedings of the ACM on Programming Languages* 5, ICFP (2021), 1–18.

- [7] A. Cayley Esq. 1859. LVIII. On the analytical forms called trees.–Part II. *The London, Edinburgh, and Dublin Philosophical Magazine and Journal of Science* 18, 121 (1859), 374–378. <https://doi.org/10.1080/14786445908642782> arXiv:<https://doi.org/10.1080/14786445908642782>
- [8] Patrice Godefroid, Hila Peleg, and Rishabh Singh. 2017. Learn&fuzz: Machine learning for input fuzzing. In *2017 32nd IEEE/ACM International Conference on Automated Software Engineering (ASE)*. IEEE, 50–59.
- [9] Harrison Goldstein and Benjamin C. Pierce. 2022. Parsing Randomness: Unifying and Differentiating Parsers and Random Generators. arXiv:[2203.00652](https://arxiv.org/abs/2203.00652) [cs.PL]
- [10] John Hughes. 2019. The code used in How to Specify It! <https://github.com/rjmh/how-to-specify-it>. (2019).
- [11] John Hughes. 2019. How to Specify It! *20th International Symposium on Trends in Functional Programming* (2019).
- [12] Leonidas Lampropoulos, Michael Hicks, and Benjamin C. Pierce. 2019. Coverage guided, property based testing. *PACMPL* 3, OOPSLA (2019), 181:1–181:29. <https://doi.org/10.1145/3360607>
- [13] Andreas Löschner and Konstantinos Sagonas. 2017. Targeted Property-Based Testing. In *Proceedings of the 26th ACM SIGSOFT International Symposium on Software Testing and Analysis* (Santa Barbara, CA, USA) (*ISSTA 2017*). Association for Computing Machinery, New York, NY, USA, 46–56. <https://doi.org/10.1145/3092703.3092711>
- [14] José Pedro Magalhães, Atze Dijkstra, Johan Jeuring, and Andres Löb. 2010. A Generic Deriving Mechanism for Haskell. In *Proceedings of the Third ACM Haskell Symposium on Haskell* (Baltimore, Maryland, USA) (*Haskell '10*). Association for Computing Machinery, New York, NY, USA, 37–48. <https://doi.org/10.1145/1863523.1863529>
- [15] Conor McBride and Ross Paterson. 2008. Applicative programming with effects. *Journal of functional programming* 18, 1 (2008), 1–13.
- [16] Lambert Meertens. 1988. First steps toward the theory of rose trees. In *IFIP Working Group 2.1*.
- [17] Agustin Mista and Alejandro Russo. 2019. Deriving compositional random generators. In *Proceedings of the 31st Symposium on Implementation and Application of Functional Languages*. 1–12.
- [18] Agustin Mista, Alejandro Russo, and John Hughes. 2018. Branching processes for QuickCheck generators. In *Proceedings of the 11th ACM SIGPLAN International Symposium on Haskell, Haskell@ICFP 2018, St. Louis, MO, USA, September 27-17, 2018*, Nicolas Wu (Ed.). ACM, 1–13. <https://doi.org/10.1145/3242744.3242747>
- [19] Sameer Reddy, Caroline Lemieux, Rohan Padhye, and Koushik Sen. 2020. Quickly generating diverse valid test inputs with reinforcement learning. In *ICSE '20: 42nd International Conference on Software Engineering, Seoul, South Korea, 27 June - 19 July, 2020*, Gregg Rothermel and Doo-Hwan Bae (Eds.). ACM, 1410–1421. <https://doi.org/10.1145/3377811.3380399>
- [20] Tim Sheard and Simon Peyton Jones. 2002. Template Meta-Programming for Haskell. In *Proceedings of the 2002 ACM SIGPLAN Workshop on Haskell* (Pittsburgh, Pennsylvania) (*Haskell '02*). Association for Computing Machinery, New York, NY, USA, 1–16. <https://doi.org/10.1145/581690.581691>
- [21] Nicholas Smallbone, Moa Johansson, Koen Claessen, and Maximilian Alghed. 2017. Quick specifications for the busy programmer. *Journal of Functional Programming* 27 (2017).
- [22] Ezekiel Soremekun, Esteban Pavese, Nikolas Havrikov, Lars Grunske, and Andreas Zeller. 2020. Inputs from Hell Learning Input Distributions for Grammar-Based Test Generation. *IEEE Transactions on Software Engineering* (2020).
- [23] Michal Zalewski. 2018. AFL quick start guide. <http://lcamtuf.coredump.cx/afl/QuickStartGuide.txt>.