

The Search for Constrained Random Generators

ANONYMOUS AUTHOR(S)

One of the most significant challenges in property-based testing (PBT) is the *constrained random generation problem*: given a predicate on program values, randomly sample from the set of all values satisfying that predicate (and only those values). Efficient solutions to this problem are critical for effective testing, since the executable specifications used by PBT often have preconditions that input values must satisfy in order to be useful for testing, and satisfying values are often sparsely distributed.

We present a novel approach to the constrained random generation problem using ideas from deductive program synthesis. We propose a set of synthesis rules, based on a denotational semantics of generators, that give rise to an automatic procedure for synthesizing correct generators. To deal with recursive predicates, we rewrite the predicate as a catamorphism (i.e., a fold) and then match that with an appropriate anamorphism (i.e., unfold); this is theoretically simpler than other approaches to synthesis for recursive functions yet still extremely expressive. Our implementation, PALAMEDES, is an extensible library for the Lean theorem prover.

1 INTRODUCTION

Property-based testing (PBT) [11] is an approach to software testing that bridges the gap between traditional testing and heavier-weight formal methods [54] by allowing developers to test software systems against formal specifications. PBT is used to find bugs in a wide range of real-world software [1, 3, 4, 7], but significant work remains. A key challenge is the *constrained random generation problem*.

Consider a classic example of the kind of property one would test with PBT:

$$\forall x \ t, \text{isBST}(t) \implies \text{isBST}(\text{insert}(x, t))$$

This says that, assuming a tree t is a valid binary search tree (BST), inserting a new value x into that tree yields another valid BST. To test this property, a PBT framework uses a program called a *generator* to randomly sample hundreds or thousands of values for x and t and check that the property holds for each pair of values. But the property is vacuously true if t is not a valid BST to start with—if the random generation procedure is not carefully designed, most generated trees will not actually be used to test the insert function. This is the main motivation for the constrained random generation problem: to test a conditional property effectively, we need a way to randomly generate *all values and only values* that satisfy its precondition.

This problem has been studied extensively by programming languages researchers over the years [10, 17, 26, 50], with many proposed solutions that address it to some degree. But as of 2024, a study on PBT’s usability in practice [16] still cited the availability of generators as a key challenge for PBT adoption. The key issue is speed: this study showed that, in practice, developers run PBT often, so they need it to be fast. But most of the established approaches to constrained generation perform some kind of combinatorial search during generation (e.g., rejection sampling, calling an SMT solver, or using feedback from code coverage), which slows down generation significantly.

We propose a different approach. Rather than searching during generation, we search ahead of time, using a *deductive synthesis* algorithm to search for generators that are guaranteed to be fast and to produce only valid values, and any valid value, by construction. This synthesis algorithm can run ahead of time, allowing for much faster and more effective testing. The algorithm is based on a denotational semantics that captures the values that a given generator can generate. If we want a generator for a property whose precondition is φ , we compute a generator g satisfying

$$\forall a, a \in \llbracket g \rrbracket \iff \varphi(a),$$

where $\llbracket g \rrbracket$ is the set of values the generator can generate. In other words, a can be generated by g if and only if $\varphi(a)$ holds. Inspired by systems like Synquid and SuSLik [42, 43], our synthesis algorithm builds a *proof* that there exists an appropriate g for a given φ by applying a series of proof rules, then extracts a proof witness—a functional program—that can be used for testing.

For straight-line code, our proof rules are roughly one-to-one with common *generator combinators*—the basic functions used to build generators by hand. To dealing with recursive data structures, we use *recursion schemes*; in particular, we observe that many predicates that can be represented as a fold (or catamorphishm) have an associated generator represented as an unfold (or anamorphism). The upshot is that we can deal with recursive predicates without relying on cyclic proofs [22], instead using pre-derived induction principles.

We implement our synthesis algorithm as a library in the Lean theorem prover that builds Lean generators from predicates expressed as Lean functions. Working in a theorem prover affords us significant benefits. First, synthesized generators come with mechanized proofs of correctness, ensuring that any generator produced by our algorithm is appropriate for testing. Next, working in Lean means the implementation of the synthesis procedure itself can be relatively simple—in fact, we use a popular proof search tactic, Aesop [29], to do much of the heavy lifting. Finally, extending the algorithm with new primitives or search tactics is as simple as proving a lemma or writing a macro. The synthesizer is already quite powerful—it can, for example, synthesize generators for BSTs and well-typed simply-typed lambda calculus (STLC) programs—and the set of predicates it can handle will continue to grow modularly.

In Section 2 we provide more concrete motivation and background, along with a preview of our approach. The remaining sections offer the following contributions:

- We propose a system of deductive synthesis rules for correct generators (Section 3). These rules are proven correct relative to a denotational semantics of generators, meaning that the synthesized generators are correct by construction.
- We extend our synthesis algorithm to work for generators of recursive data types, leveraging intuitions from the literature on recursion schemes (Section 4). In particular, we establish a relationship between predicates written as folds and generators written as unfolds.
- We describe PALAMEDES, an implementation of our synthesis procedure that is embedded as a library in the Lean proof assistant (Section 5). In this context, we can achieve performant synthesis that maintains mechanized proofs of correctness, while borrowing key parts of the algorithm from Lean’s existing infrastructure.
- We evaluate our approach with a suite of case studies, demonstrating that our synthesis procedure produces generators comparable to handwritten generators for a wide range of predicates (Section 6). Our most complex case studies—generators for binary search trees, AVL trees, and well-typed STLC terms—have been PBT benchmarks for decades.

We conclude with a discussion of limitations (Section 7), related work (Section 8), and future work (Section 9).

2 MOTIVATION

We begin by explaining the constrained random generation problem (Section 2.1), reviewing current approaches to addressing the problem (Section 2.2), and motivating a refinement (Section 2.3)—the constrained generator synthesis problem—which we address in the rest of the paper.

2.1 Fast Generators for Property-Based Testing

Property-based testing uses random test data to validate executable program specifications. In this paper we focus on the *generators* that produce the random test inputs used to exercise the system

under test, which are critical to performant and effective testing. For example, when testing the property in [Section 1](#), a developer might write a generator like the one in [Figure 1](#).

```

def genBST lo hi :=
  if lo > hi then
    pure leaf
  else
    pick
      (pure leaf)
      (do
        let x <- choose lo hi
        let l <- genBST lo (x - 1)
        let r <- genBST (x + 1) hi
        pure (node l x r))

```

Fig. 1. A hand-written generator for binary search trees.

This generator produces random BSTs with values in a given range. In the case where the range is empty, pure creates a constant generator that always returns leaf. Otherwise, pick is used to make a choice: either generate a leaf or generate a node by selecting a value in the appropriate range and recursively generating subtrees with truncated ranges. The **do** notation sequences generators by sampling from a generator, binding the sampled value to a variable, and then continuing.

While generators like this are familiar to PBT experts, a novice might struggle to come up with this generator from scratch. Moreover, a recent study of PBT users [16] found that even experts, who can in principle write generators

like the one above, still see writing effective generators as a distraction from the other testing tasks.

What makes an effective generator? We can start with two key properties:

- (1) **Soundness** The generator should produce *only* values that satisfy a given validity predicate. This is important for properties that have preconditions (e.g., the BST one above).
- (2) **Completeness** The generator should be able to produce *any* value that satisfies the validity predicate. This ensures that the generator does not miss important parts of the input space.

We can now phrase the constrained random generation problem more formally:

Definition 2.1 (Constrained Random Generation Problem). Given a predicate φ , sample random values in a way that is sound and complete with respect to φ .

Solving this problem—and, in particular, finding fast solutions—is critical if we want to build PBT libraries that are usable for developers and powerful enough to find bugs in software.

2.1.1 Aside About Distributions. Readers familiar with PBT literature might have expected us to say something about the *distribution* of a generator, but (aside from this aside) we are not going to. We treat generators as nondeterministic programs, focusing on the set of values that they can produce and ignoring the probabilities with which they produce them. This is not because probabilities are unimportant—they have a huge impact on testing performance—but because they can be handled separately. Recent work [53] has shown that, once a generator is written, it can be tuned by external processes to achieve various desirable distributional qualities; we defer distributional concerns to such processes. See [Section 9](#) for more.

2.2 Search-Based Approaches to Generation

Many approaches to the constrained random generation problem have been proposed over the years, most based around some kind of search procedure. (There are some notable exceptions [25, 27], but we delay comparisons with these to [Section 8](#), where we can go into more technical detail.)

Most of the available search procedures start with a naïve, complete PBT generator (e.g., derived from type information [33, 56]) and prune its generation space “online,” during the generation process, to ensure soundness. The simplest way to prune the space is via *rejection sampling*—discarding any invalid values and retrying generation—but this results in poor testing performance

for even moderately sparse predicates. More advanced mechanisms for pruning include laziness [10], Brzozowski derivatives [17], reinforcement learning [46], and constraint solving [48, 50]. Others try to search for valid inputs “from scratch,” for example with the help of a large language model [55].

But there is something unsatisfying about all of these search-based approaches: genBST, above, does no searching at all! It produces valid inputs by construction, every time, essentially as quickly as possible. Could we automatically build generators like genBST that do no searching?

2.3 The Constrained Generator Synthesis Problem

We propose a refinement of the constrained random generation problem that focuses specifically on building a correct generator directly, rather than relying on search processes:

Definition 2.2 (Constrained Generator Synthesis Problem). Given a predicate φ , synthesize an efficient generator g that is sound and complete with respect to φ .

The difference between this and Definition 2.1 is subtle but important: we are specifically interested in solutions that synthesize an efficient generator first and then use it to sample valid values; solutions that actively search for valid values during generation are excluded. We will discuss what we mean by “efficient” in the next section. This version of the problem is harder, but solutions also have the potential to be much more effective for PBT users, since generators are typically written once and then run many times.

3 DEDUCTIVE SYNTHESIS FOR GENERATORS

We propose a novel approach to the constrained generator synthesis problem. We start by introducing our representation of generators and some key definitions (Section 3.1); next we present our core deductive synthesis rules for constructing generators (Section 3.2); finally, we discuss an optimization procedure that can be applied to generators after synthesis (Section 3.4).

The presentation in this section is phrased in terms of an unspecified ambient dependent type theory; in Section 5 we make the definitions and proofs concrete in the Lean proof assistant.

3.1 Generator Representation

While it is often standard to represent generators as sampling functions of type `Seed → a`, we opt for a representation with a bit more flexibility. Taking inspiration from work on *free generators* [17], we represent generators as data structures that can be interpreted in multiple ways, including as sampling functions. These generators are represented by the following inductive data type:

inductive Gen where
`pure` : $\alpha \rightarrow \text{Gen } \alpha$
`bind` : $\text{Gen } \beta \rightarrow (\beta \rightarrow \text{Gen } \alpha) \rightarrow \text{Gen } \alpha$
`pick` : $\text{Gen } \alpha \rightarrow \text{Gen } \alpha \rightarrow \text{Gen } \alpha$
`indexed` : $(\mathbb{N} \rightarrow \text{Gen } (\text{Option } \alpha)) \rightarrow \text{Gen } \alpha$
`assume` : $(b : \mathbb{B}) \rightarrow (b = \text{true} \rightarrow \text{Gen } \alpha) \rightarrow \text{Gen } \alpha$

We often write `bind` with the infix notation `>>=`.

While these constructors are simply data, they each have a standard interpretation as a procedure for sampling values (we discuss the interpretation in more detail in Section 5). The `pure` constructor represents a constant generator that always produces the same value. The `>>=` constructor sequences generators, sampling from one and passing the sampled value to a function producing another. The `pick` constructor represents a choice between generators. The `assume` constructor represents a partial generator. It checks a boolean condition; if true, it simply calls its argument, but if false it

fails, generating nothing. Finally, the **indexed** constructor represents an infinite family of generators, indexed by natural numbers.¹

Next, we define the support of a generator.

Definition 3.1. The *support* of a generator g is the set of values that g can produce. We denote it by $\llbracket g \rrbracket$. Support is defined as follows:

$$\begin{aligned} a \in \llbracket \text{pure } a' \rrbracket &\iff a = a' \\ a \in \llbracket x \ggg f \rrbracket &\iff \exists a', a' \in \llbracket x \rrbracket \wedge a \in \llbracket f a' \rrbracket \\ a \in \llbracket \text{pick } x y \rrbracket &\iff a \in \llbracket x \rrbracket \vee a \in \llbracket y \rrbracket \\ a \in \llbracket \text{assume } b \text{ in } x \rrbracket &\iff b = \text{true} \wedge a \in \llbracket x \rrbracket \\ a \in \llbracket \text{indexed } f \rrbracket &\iff \exists n, \text{ some } a \in \llbracket f n \rrbracket \end{aligned}$$

These definitions follow the intuition given above and agree with the denotational semantics of generators presented in prior work [39].

Example 3.1 (Generator for Natural Numbers). The following generator uses all but one of the above constructors to generate the set of all natural numbers:

```
def arbNat : Gen ℕ :=
  let rec go (fuel : ℕ) : Gen (Option ℕ) :=
    match fuel with
    | 0 => pure none
    | fuel' + 1 =>
      pick
        (pure (some 0))
        (go fuel' >> λ on' =>
          match on' with
          | none => pure none
          | some n' => pure (some (1 + n'))))
  indexed go
```

It defines a recursive function `go` that takes some fuel and produces a potentially failing generator of natural numbers. If the fuel has run out, the generator fails with `none`. Otherwise, the generator makes a random choice between returning 0 and returning $1+n'$, where n' is generated by recursively calling `go`. We use `indexed` to turn this indexed family of partial generators into a total generator.

The support of this generator is precisely \mathbb{N} (we prove this in our Lean development).

Example 3.2 (Partial Generator). While our ultimate goal is to avoid generators that search during generation, we need such generators to be representable in our language (we will see why later when synthesizing generators for BSTs). Our synthesis procedure works by first producing potentially partial generators and then attempting to optimize them into total generators.

Here is an example of a generator that is partial:

```
def backtracks := pick (pure 1) (assume false in pure 2)
```

The support of this generator is $\{1\}$ —that is to say, when it generates a value, it always generates 1—but it will sometimes choose the right side of the `pick` and fail.

We can use support to formally define what it means for a generator to be correct.

¹In lazy languages like Haskell this constructor is not necessary. But since we will be working in Lean (which is strict) we need this constructor to be able to represent generators of infinite sets (e.g., a generator of all natural numbers).

Definition 3.2 (Correctness). A generator g is *correct* with respect to a predicate φ if it is both *sound*, i.e., $\forall a, a \in \llbracket g \rrbracket \implies \varphi(a)$, and *complete*, i.e., $\forall a, \varphi(a) \implies a \in \llbracket g \rrbracket$.

Notation 3.1 (Correct Generator). The type of correct generators with respect to a predicate φ on α is denoted $\text{Gen}_\alpha \varphi$. The type of the generated value becomes a subscript; we may leave it off if it is clear from context.

Example 3.3. The `arbNat` generator above can be given the type $\text{Gen}_{\mathbb{N}} (\lambda n \Rightarrow \top)$; backtracks has type $\text{Gen}_{\mathbb{N}} (\lambda n \Rightarrow n = 1)$.

We now return to the notion of “efficiency” that we skipped in §2.3. Our gold standard for generators is demonstrated by `genBST`—generators that are like the ones written by expert users to produce valid inputs by construction. The synthesis procedure we describe over the next few sections meets this standard in most cases, which we demonstrate in Section 6 by comparing synthesized generators with user-written ones, but there are situations in which it is not completely successful. In particular, the `assume` constructor means that the synthesis procedure can occasionally produce partial generators that backtrack more than an expert-written generator might.

To classify these suboptimal generators, we define the following notion:

Definition 3.3 (Assume Freedom). A generator g is *assume-free* iff it does not use the `assume` constructor (including in sub-generators that it calls).

Each of our case studies in Section 6 is labeled to clarify whether or not the generator is assume-free; all are correct by construction.

3.2 Core Synthesis Algorithm

We will now outline an algorithm to solve the Correct Generator Synthesis Problem. The algorithm uses *deductive program synthesis*, constructing a generator by working backwards from the structure of the predicate. It creates a proof that witnesses the generator’s correctness and creates the generator itself *en passant*. This approach to synthesis can be found in systems like SuSLik [43] and Synquid [42]. Concretely, we start with the statement

$$\frac{}{\Gamma \vdash ? : \text{Gen}_\alpha \varphi} ?$$

and successively refine the generator by applying a series of *synthesis rules* to build a complete derivation.

Pure and Pick. Here are our first two basic synthesis rules:

$$\frac{\Gamma \vdash a' : \alpha}{\Gamma \vdash \text{pure } a' : \text{Gen}_\alpha (\lambda a \Rightarrow a = a')} \text{S-PURE}$$

$$\frac{\Gamma \vdash x : \text{Gen}_\alpha P \quad \Gamma \vdash y : \text{Gen}_\alpha Q}{\Gamma \vdash \text{pick } x \ y : \text{Gen}_\alpha (\lambda a \Rightarrow P a \vee Q a)} \text{S-PICK}$$

The former says that we can synthesize a value that is equal to a constant using `pure`, and the latter says that we can synthesize a disjunction by using `pick`. We can use these rules to synthesize a generator for the predicate $\lambda a \Rightarrow a = 1 \vee a = 2$:

$$\frac{\frac{\frac{}{\cdot \vdash 1 : \mathbb{N}} \Downarrow}{\cdot \vdash \text{pure } 1 : \text{Gen } (\lambda a \Rightarrow a = 1)} \text{S-PURE} \quad \frac{\frac{}{\cdot \vdash 2 : \mathbb{N}} \Downarrow}{\cdot \vdash \text{pure } 2 : \text{Gen } (\lambda a \Rightarrow a = 2)} \text{S-PURE}}{\cdot \vdash \text{pick } (\text{pure } 1) (\text{pure } 2) : \text{Gen } (\lambda a \Rightarrow a = 1 \vee a = 2)} \text{S-PICK}$$

While the core of the proof tree is focused on the synthesis rules, there are other things going on. In particular, when applying S-PURE on the left, we also need to prove that $\cdot \vdash 1 : \mathbb{N}$. In this case the proof is trivial, but in general our synthesizer may need to be able to discharge certain nontrivial theorems automatically. In §5 we discuss the specifics of this process, but for now we assert that any rule labeled “ \Downarrow ” can be discharged automatically.

Sometimes our synthesis rules do not apply directly to the goal as stated. For example consider the goal:

$$\frac{}{\cdot \vdash ? : \text{Gen} (\lambda a \Rightarrow 1 = a)} ?$$

This is obviously equivalent to a goal that we know how to deal with (i.e., with $a = 1$), but it is not syntactically the same. In these cases, we need to apply the *conversion rule*:

$$\frac{\frac{}{\Gamma \vdash \varphi = \psi} \Downarrow \quad \Gamma \vdash g : \text{Gen} \psi}{\Gamma \vdash g : \text{Gen} \varphi} \text{CONVERT}$$

The resulting derivation looks like this:

$$\frac{\frac{}{\cdot \vdash (\lambda a \Rightarrow a = 1) = (\lambda a \Rightarrow 1 = a)} \Downarrow \quad \frac{}{\cdot \vdash \text{pure } 1 : \text{Gen} (\lambda a \Rightarrow a = 1)} \text{S-PURE}}{\cdot \vdash \text{pure } 1 : \text{Gen} (\lambda a \Rightarrow 1 = a)} \text{CONVERT}$$

Assumptions and Functions. If the generator we need is already available in the typing context, Γ , we can just use it.

$$\frac{(x : \text{Gen}_\alpha \varphi) \in \Gamma}{\Gamma \vdash x : \text{Gen}_\alpha \varphi} \text{S-ASSUMPTION}$$

When the synthesis goal is a function returning a generator, we can apply an introduction rule.

$$\frac{b:\beta, \Gamma \vdash x : \text{Gen}_\alpha \varphi}{\Gamma \vdash \lambda b \Rightarrow x : (b : \beta) \rightarrow \text{Gen}_\alpha \varphi} \text{S-INTRO}$$

This rule says that if we can produce an appropriate generator given $b:\beta$ in the context, then we can produce a (dependent) function from β to that generator.

We also need a special case for dealing with functions that take tuples as arguments, which appear frequently as a result of our rules for recursive functions (see Section 4):

$$\frac{\Gamma \vdash f : (b : \beta) \rightarrow (c : \gamma) \rightarrow \text{Gen}_\alpha (\varphi b c)}{\Gamma \vdash \lambda (b, c) \Rightarrow f b c : (p : \beta \times \gamma) \rightarrow \text{Gen}_\alpha (\varphi (\text{fst } p) (\text{snd } p))} \text{S-UNCURRY}$$

In words, this rule says that we are free to synthesize a curried function and then uncurry it, if the goal is to produce a function with a tuple argument.

Bind. By analogy with the definition of support for `bind`, we can define a synthesis rule for composing generators:

$$\frac{\Gamma \vdash x : \text{Gen}_{\alpha'} P \quad \Gamma \vdash f : (a' : \alpha') \rightarrow \text{Gen}_\alpha (Q a')}{\Gamma \vdash x \gg f : \text{Gen}_\alpha (\lambda a \Rightarrow \exists (a' : \alpha'), P a' \wedge Q a' a)} \text{S-BIND}$$

The goal of this rule requires that the predicate we want to generate for is an existential statement with two conjuncts; P should be some statement constraining a value a' , and then Q should constrain the final value, a , given a particular a' . A generator for a predicate of this form looks like a generator x of a' 's satisfying P and a function f that takes an a' and produces a generator for a s satisfying the predicate $Q a'$.

We can use \gg to chain generators together:

$$\frac{\Gamma \vdash \dots : \text{Gen}(\lambda j \Rightarrow j = 1 \vee j = 2) \quad \frac{\Gamma, j:\mathbb{N} \vdash \text{pure}(j+3) : \text{Gen}(\lambda i \Rightarrow i = j+3)}{\Gamma \vdash \lambda j \Rightarrow \text{pure}(j+3) : j:\mathbb{N} \rightarrow \text{Gen}(\lambda i \Rightarrow i = j+3)}}{\Gamma \vdash \text{pick}(\text{pure } 1)(\text{pure } 2) \gg (\lambda j \Rightarrow \text{pure}(j+3)) : \text{Gen}(\lambda i \Rightarrow \exists j, (j = 1 \vee j = 2) \wedge i = j+3)}$$

The goal here is to generate an i that is equal to $j+3$, where j is further constrained to be either 1 or 2. To synthesize an appropriate generator, we use S-BIND, which gives two sub-goals. On the left, the goal is now to synthesize a generator for j , which we complete by the process above. On the right, we apply S-INTRO followed by S-PURE to produce the continuation of the bind. The final generator generates 1 or 2 and then adds 1.

The conclusion S-BIND may look strange, considering that our ultimate goal is to synthesize generators for the kinds of predicates that real PBT users care about. Indeed, existential quantification cannot be expressed directly in most programming languages. However, many executable predicates are logically equivalent to ones with existentials. For example, the predicate `isSome x` is equivalent to $\exists a, x = \text{some } a$. This means that the CONVERT rule is critical for S-BIND—it allows the synthesizer to logically manipulate the predicate before applying the rule to uncover the existential quantification implicit in a given predicate. We discuss those logical manipulations in Section 5.2.

Case Splitting. If a generator needs to do different things based on the value of a variable in the context, it can use one of the following rules to split that variable into cases.

$$\frac{\Gamma \vdash x : \text{Gen}_\alpha(\varphi \text{ true}) \quad \Gamma \vdash y : \text{Gen}_\alpha(\varphi \text{ false})}{b:\mathbb{B}, \Gamma \vdash \text{match } b \text{ with } | \text{true} \Rightarrow x | \text{false} \Rightarrow y : \text{Gen}_\alpha(\varphi b)} \text{S-SPLITBOOL}$$

$$\frac{\Gamma \vdash x : \text{Gen}_\alpha(\varphi 0) \quad n':\mathbb{N}, \Gamma \vdash y n' : \text{Gen}_\alpha(\varphi (n' + 1))}{n:\mathbb{N}, \Gamma \vdash \text{match } n \text{ with } | 0 \Rightarrow x | n' + 1 \Rightarrow y n' : \text{Gen}_\alpha(\varphi n)} \text{S-SPLITNAT}$$

Here we give rules for booleans and natural numbers; rules for other inductive types can be derived from their definitions (see Section 5). Each of these rules corresponds to a match in the final generator; they pick out a variable in the context, test it, and then choose a different generator depending on the outcome of that test. Note that these splits are *not* recursive; we deal with recursive generators in the next section.

A key benefit of this approach, and deductive synthesis in general, is that the generators we arrive at by applying these rules are correct by construction. The process that produces the generator simultaneously builds a proof that the generator's support is equivalent to the desired support. If this process finds a generator, a developer can be confident that the generator is appropriate for their testing needs.

3.3 Standard Library Functions

The rules from the previous section are the core of our synthesis algorithm, and they can be used to build surprisingly complex generators, but the real power of this approach comes from its extensibility. Rather than ask the synthesis process to synthesize “all the way down,” we can provide it with a library of building blocks that it can use to build more complex generators. For the examples in the rest of this paper, we will need a few generators that are standard in PBT libraries.

Choose. The choose generator picks a natural number in a defined range. We define it by recursion:

```
def choose (lo hi := if lo = hi then pure lo else pick (pure lo) (choose (lo + 1) hi))
```

We can characterize its support as follows:

Lemma 3.1 (Choose Support). If $lo \leq hi$, then $a \in \llbracket \text{choose } lo \ hi \rrbracket \iff lo \leq a \leq hi$.

And we can give the corresponding synthesis rule:

$$\frac{\Gamma \vdash lo \leq hi}{\Gamma \vdash \text{choose } lo \ hi : \text{Gen}_{\mathbb{N}} (\lambda a \Rightarrow lo \leq a \leq hi)} \text{S-CHOOSE}$$

In a synthesis context, it may not always be easy to show that $lo \leq hi$ (indeed, it might not even be true). This motivates a second way to synthesize choose that checks its precondition explicitly.

$$\frac{}{\Gamma \vdash \text{assume } lo \leq hi \text{ in choose } lo \ hi : \text{Gen}_{\mathbb{N}} (\lambda v \Rightarrow lo \leq v \leq hi)} \text{S-CHOOSEPARTIAL}$$

This rule is valid, but it introduces the potential for the generator to fail: if $lo > hi$ when this generator is executed, it will not be able to produce a value. Critically, this generator is still correct—it is complete, and it is sound in the sense that if it produces a value then that value satisfies the given condition—but it is not ideal for testing. Luckily, we can usually optimize this issue away.

Greater Than and Less Than. The greaterThan and lessThan generators take a single natural number and can generate any number greater or less than that number. Their implementation and associated lemmas are similar the ones for choose, so we delay them to [Appendix A](#).

Elements. Finally, the elements generator picks a random value from a list:

def elements (xs : List α) := match xs with | [x] \Rightarrow pure x | x :: xs' \Rightarrow pick (pure x) (elements xs')

It has the following support and synthesis rules:

Lemma 3.2 (Elements Support). If $xs \neq []$, then $a \in \llbracket \text{elements } xs \rrbracket \iff a \in xs$.

$$\frac{\Gamma \vdash xs \neq []}{\Gamma \vdash \text{elements } xs : \text{Gen}_{\alpha} (\lambda a \Rightarrow a \in xs)} \text{S-ELEMENTS}$$

$$\frac{}{\Gamma \vdash \text{assume } xs \neq [] \text{ in elements } xs : \text{Gen}_{\alpha} (\lambda a \Rightarrow a \in xs)} \text{S-ELEMENTSPARTIAL}$$

3.4 Optimizing Generators to Avoid Assumes

In most cases where the synthesizer inserts assumptions, they can be optimized away. For example, consider the following synthesized generator:

$$\frac{\dots \quad lo:\mathbb{N}, hi:\mathbb{N} \vdash \text{assume } lo \leq hi \text{ in choose } lo \ hi : \text{Gen} (\lambda a \Rightarrow lo \leq a \leq hi)}{lo:\mathbb{N}, hi:\mathbb{N} \vdash \text{pick (pure 0) (assume } lo \leq hi \text{ in choose } lo \ hi) : \text{Gen} (\lambda a \Rightarrow a = 1 \vee lo \leq a \leq hi)}$$

Logically, this `assume` is not necessary. As written, the generator makes a choice and then fails if it happened to choose the right side and $lo > hi$. But it could just as well check $lo \leq hi$ first and only choose the right branch if the check succeeds. We leverage this observation by designing *optimization rules* that rewrite generators to avoid failures. The rule we need for the above case is

$$\text{pick } x \text{ (assume } b \text{ in } y) \rightsquigarrow \text{if } b \text{ then pick } x \ y \text{ else } x,$$

which rewrites a `pick` containing an `assume` to an if statement that checks the assumption before the choice. Concretely, we have:

$$\begin{aligned} & \text{pick (pure 0) (assume } lo \leq hi \text{ in choose } lo \ hi) \rightsquigarrow \\ & \text{if } lo \leq hi \text{ then pick (pure 0) (choose } lo \ hi) \text{ else pure 0} \end{aligned}$$

In total, we use six optimization rules to facilitate uncovering `assumes` and lifting them out of choices:

$$\text{pure } v \gg\!\!\gg f \rightsquigarrow f v \quad (1)$$

$$(x \gg\!\!\gg g) \gg\!\!\gg f \rightsquigarrow x \gg\!\!\gg (\lambda a \Rightarrow g a \gg\!\!\gg f) \quad (2)$$

$$(\text{assume } b \text{ in } x) \gg\!\!\gg f \rightsquigarrow \text{assume } b \text{ in } (x \gg\!\!\gg f) \quad (3)$$

$$x \gg\!\!\gg (\lambda a \Rightarrow \text{assume } b \text{ in } (f a)) \rightsquigarrow \text{assume } b \text{ in } (x \gg\!\!\gg f) \quad \text{if } a \notin \text{fv}(b) \quad (4)$$

$$\text{pick } (\text{assume } b \text{ in } x) y \rightsquigarrow \text{if } b \text{ then pick } x y \text{ else } y \quad (5)$$

$$\text{pick } x (\text{assume } b \text{ in } y) \rightsquigarrow \text{if } b \text{ then pick } x y \text{ else } x \quad (6)$$

Rules (1) and (2) are standard monad equivalences, rules (3) and (4) describe how `assumes` interact with binds, and (5) and (6) actually lift `assumes` out of choices.

Lemma 3.3 (Optimizations Correct). Rules (1)–(6) do not change the support of the generator.

These rules are not complete; they may still leave `assumes` in the generator, e.g., when there are incompatible assumptions on the two sides of a `pick`. But, for most of the examples in Section 6, they are enough to create a generator that is `assumeFree`.

4 SYNTHESIZING GENERATORS FOR RECURSIVE STRUCTURES

We next describe how our synthesis procedure handles predicates over inductive data structures like lists and trees. Since our approach is based on *recursion schemes*, we start with some background on those (Section 4.1). Then we outline our synthesis procedure at a high level (Section 4.2), building up to a fully detailed picture of generators for recursive data (Section 4.3). Finally, we discuss the end-to-end process in a bit more detail (Section 4.5).

At a high level, the approach in this section takes some recursive predicate, transforms it into a normal form, and then applies a synthesis rule like the ones from the previous section. The pipeline that we implement is shown in Figure 2. While the details in this section are important, there is one key takeaway: rather than give the synthesizer direct access to `indexed` and recursion, we synthesize all recursive generators through higher-level rules. This approach does have limitations—we cannot directly synthesize generators like elements that iterate over one structure and produce another—but it is highly effective for predicates that directly constrain data structures.




Fig. 2. How a list predicate is normalized for synthesis.

4.1 Background: Recursion Schemes

Recursive functions are a common challenge for program analysis and synthesis tasks, even in strongly normalizing languages where they are guaranteed to terminate. While there are techniques available for synthesizing recursive programs directly from recursive specifications [42, 43], we take a different approach that is theoretically simpler and easier to embed in Lean.

In the functional programming community, there is a rich literature on *recursion schemes*. Rather than express recursive functions directly, via unstructured general recursion, recursion schemes abstract recursion into structured forms that are easier to reason about.

4.1.1 *Folds*. The simplest recursion scheme is a *fold* or *catamorphism*. Here is an implementation of fold for the List datatype:

```
def List.fold (f :  $\alpha \rightarrow \beta \rightarrow \beta$ ) (z :  $\beta$ ) (xs : List  $\alpha$ ) :  $\beta$  :=
  match xs with
  | []  $\Rightarrow$  z
  | x :: xs'  $\Rightarrow$  f x (List.fold f z xs')
```

This function takes as arguments a “base-case” z and a “step function” f . We call the type β the “collector” for the fold.² When the list is empty, we return z . When the list is a cons-cell, we recursively call List.fold $f z$ on the tail of the list and then use f to combine the resulting value with the value at the head.

Note that information here flows backward, from the tail of the list to the head.³ We first compute something about the tail of the list, without considering the value at the head, and only at the end do we actually consider the information at the head. This will be useful to remember later.

We can use List.fold to implement common functions over lists, for example length:

```
def length xs := List.fold ( $\lambda x b \Rightarrow 1 + b$ ) 0 xs
```

4.1.2 *More Advanced Recursion Schemes*. While List.fold can represent all terminating functions over lists [21], it is not always ergonomic to do so. For this reason, researchers have identified dozens of specialized recursion schemes, each capturing some common pattern of recursion that arises in functional programs [59]. Of note for this paper is the *accumulation* [40], which we present with some minor simplifications:

```
def List.accu (st :  $\alpha \rightarrow \sigma \rightarrow \sigma$ ) (s :  $\sigma$ ) (f :  $\alpha \rightarrow \beta \rightarrow \sigma \rightarrow \beta$ ) (z :  $\sigma \rightarrow \beta$ ) (xs : List  $\alpha$ ) :  $\beta$  :=
  match xs with
  | []  $\Rightarrow$  z s
  | x :: xs'  $\Rightarrow$  f x (List.accu st (st x s) f z xs') s
```

Accumulations are similar to folds, but they pass information in both directions. The collector value β is still passed backwards through the list, and we add a new type parameter σ representing an “accumulation state” that flows forward. The function takes an initial state s as input, along with a “state update function” st that says how the state changes based on the value at the head of the list.

We can use List.accu to implement some functions that would be awkward to implement with List.fold. For example, this function checks if a list of natural numbers is sorted:

```
def sorted xs := List.accu ( $\lambda x \_ \Rightarrow x$ ) 0 ( $\lambda x b lo \Rightarrow lo \leq x \wedge b$ ) ( $\lambda \_ \Rightarrow \text{true}$ ) xs
```

The accumulation state is the minimum value allowable in the remaining segment of the list; it is initialized to 0 and updated to be the most recently seen value at each step. The step function then checks that $lo \leq x$ and conjoins that with the value computed from the tail of the list (which ensures that the tail of the list is also sorted). We could implement this same function with List.fold, but this would require the collector to be a higher-order function. This higher-order function can be difficult to work with, so accumulations provide a convenient alternative representation.

²Other texts call this value the “accumulator,” but we use “accumulation” to refer to a type of fold [40] and want to avoid confusion.

³This is a “right fold” over the list (i.e., List.foldr). In this paper we drop the “r” for consistency across data structures; left folds are natural for lists, but they do not have an analog for algebraic data types with branching recursion like trees.

4.1.3 *Optional Folds*. We need versions of `List.fold` and `List.accu` that capture *optional computations*:

```

def List.foldOpt
  (f :  $\alpha \rightarrow \beta \rightarrow \text{Option } \beta$ )
  (z :  $\text{Option } \beta$ )
  (xs :  $\text{List } \alpha$ ) :  $\text{Option } \beta :=$ 
  match xs with
  | []  $\Rightarrow$  z
  | x :: xs'  $\Rightarrow$ 
    match List.foldOpt f z xs' with
    | none  $\Rightarrow$  none
    | some b'  $\Rightarrow$  f x b'

def List.accuOpt
  (st :  $\alpha \rightarrow \sigma \rightarrow \sigma$ ) (s :  $\sigma$ )
  (f :  $\alpha \rightarrow \beta \rightarrow \sigma \rightarrow \text{Option } \beta$ )
  (z :  $\sigma \rightarrow \text{Option } \beta$ )
  (xs :  $\text{List } \alpha$ ) :  $\text{Option } \beta :=$ 
  match xs with
  | []  $\Rightarrow$  z s
  | x :: xs'  $\Rightarrow$ 
    match List.accuOpt st (st x s) f z xs' with
    | none  $\Rightarrow$  none
    | some b'  $\Rightarrow$  f x b' s

```

We discuss use cases for these in the next section; for now, notice that they behave the same way as `List.fold` and `List.accu`, but if any step of the iteration evaluates to none, the whole function does.

4.1.4 *Unfolds*. The final recursion scheme we will examine in detail is an *unfold* or *anamorphism*. Unfolds are the inverse of fold: whereas folds collapse data structures into compact values, unfolds expand values into data structures.

The following function nondeterministically “unfolds” a value into a whole data structure:

```

def List.unfoldGen (g :  $\beta \rightarrow \text{Gen } (\text{Option } (\alpha \times \beta))$ ) (b :  $\beta$ ) :  $\text{Gen } (\text{List } \alpha) :=$ 
  let rec go b fuel :=
    match fuel with
    | 0  $\Rightarrow$  pure (none)
    | 1 + fuel'  $\Rightarrow$ 
      g b  $\gg=$   $\lambda$  step  $\Rightarrow$ 
      match step with
      | none  $\Rightarrow$  pure (some [])
      | some (x, b')  $\Rightarrow$ 
        go b' fuel'  $\gg=$   $\lambda$  mxs  $\Rightarrow$ 
        match mxs with | none  $\Rightarrow$  pure none | some xs  $\Rightarrow$  pure (some (x :: xs))
  indexed (go b)

```

The internal function `go` takes a seed value `b` and some fuel. If the fuel is gone, the function returns none. Otherwise, it samples `g b` to obtain a “step”—if the step is none then generation terminates with an empty list, and if the step is some (x, b') then generation continues with a node containing the value `x` and a new seed value `b'`. We use the `indexed` constructor to unify this indexed family of generators into a single generator for lists.

A key benefit of `List.unfoldGen` is that it is guaranteed to make exactly one recursive call for each element of the list it produces. This is important for efficiency: it means that generators implemented with `List.unfoldGen` (as opposed to arbitrary general recursion) are guaranteed to be efficient as long as their step functions are efficient.

4.2 Generators for Inductive Data Types

Many real-world properties take inductive data as input, so it is important that our synthesis procedure be able to handle predicates over inductive data. The powerful tools for structuring recursion that we reviewed in the previous subsection will allow us to do just that.

For a first example, consider the following predicate, which checks that a list has a given length:

```
def hasLengthK (k : ℕ) (xs : List Nat) := List.fold (λ x b => 1 + b) 0 xs = k
```

This uses the definition of length that we saw earlier, which relies on `List.fold` for recursion. We can also use `List.unfoldGen` to write a generator for values satisfying this predicate:

```
def genLengthK (k : ℕ) : Gen (List ℕ) :=
  List.unfoldGen
    (λ n =>
      match k with
      | 0 => pure none
      | 1 + k' => arbNat >> λ x => pure (some (x, k'))))
    n
```

At each unfolding step, the generator checks the seed value n . If $n = 0$ then it generates `none`, indicating that the list should end (this makes sense, since n is the target length of the list). Otherwise, it generates an arbitrary natural number x and yields `some (x, n - 1)` to indicate that the list should continue with a cons-cell containing x , plus a new target length.

How might we derive `genLengthK` from `hasLengthK`? The key observation is that `hasLengthK` and `genLengthK` have an inverse relationship—whenever `genLengthK` takes a step, it is guaranteed that `hasLengthK` can undo that step. We can make that observation concrete with the following propositions:

$$\begin{array}{ll}
 \text{none} \in \llbracket (\lambda n \Rightarrow & \text{some } (x, b') \in \llbracket (\lambda n \Rightarrow \\
 \text{match } k \text{ with} & \text{match } k \text{ with} \\
 | 0 \Rightarrow \text{pure none} & | 0 \Rightarrow \text{pure none} \\
 | 1 + k' \Rightarrow & | 1 + k' \Rightarrow \\
 \text{arbNat} \gg \lambda x \Rightarrow & \text{arbNat} \gg \lambda x \Rightarrow \\
 \text{pure (some (x, k')) } b \rrbracket & \text{pure (some (x, k')) } b \rrbracket \\
 \iff & \iff \\
 b = 0 & b = (\lambda x b \Rightarrow 1 + b) x b'
 \end{array}$$

We can see that `genLengthK`'s step function returns `none` precisely when b is 0—the initial collector value for `hasLengthK`. Likewise, it returns `some (x, b')` precisely when b is $1 + b'$ —which is the result of applying `hasLengthK`'s step function to x and b' .

We can state a more general version of this relationship as a lemma.

Lemma 4.1 (Fold-UnfoldGen-Inverse for Lists). If, for all values b , the following relationship holds between an unfold's step function g and a fold's arguments f and z ,

$$\begin{array}{l}
 \text{none} \in \llbracket g b \rrbracket \iff b = z \\
 \forall x b', \text{ some } (x, b') \in \llbracket g b \rrbracket \iff b = f x b',
 \end{array}$$

then the following relationship holds of the unfold and fold

$$\forall xs, xs \in \llbracket \text{List.unfoldGen } g b \rrbracket \iff \text{List.fold } f z xs = b.$$

The informal argument for this lemma's correctness bears repeating: the fold and unfold are inverses because, for each step the unfold takes, the fold is guaranteed to be able to “fold that step back up” and get the seed. Another perspective comes from the observation that a fold passes information backwards in a list from the tail to the head; the unfold does the opposite, passing information forwards while ensuring that the fold would always compute the same information going the other way.

We can use Lemma 4.1 to prove the following synthesis rule correct:

$$\frac{\Gamma \vdash g : (b : \beta) \rightarrow \text{Gen}_{\text{Option } (\alpha \times \beta)} (P b)}{\Gamma \vdash \text{List.foldGen } g \ b : \text{Gen}_{\text{List } \alpha} (\lambda xs \Rightarrow \text{List.fold } f \ z \ xs = b)} \text{S-UNFOLDGEN'}$$

where $P b = \lambda step \Rightarrow (step = \text{none} \wedge z = b) \vee (\exists x \ b', \ step = \text{some } (x, b') \wedge f \ x \ b' = b)$

Indeed, our system can use this rule to synthesize `genLengthK` from only the definition of `hasLengthK`.

4.3 Handling More Complex Folds

The S-UNFOLDGEN' rule works for predicates whose folds have exact matches as unfolds, but other folds require pre-processing to ensure synthesis is effective. For example, consider a simple predicate that checks that all elements of a list are equal to 2:

```
def allTwo (xs : List Nat) := List.fold (λ x b ⇒ x = 2 ∧ b) true xs = true
```

Using S-UNFOLDGEN', we could turn this directly into a generator of the form:

```
List.foldGen
  (λ b ⇒
    match b with
    | true ⇒ pick (pure none) (pure (some (2, true)))
    | false ⇒ ...) true,
```

But note that the false branch will never be executed: `b` starts as true and remains true every step through the unfold. We would prefer to avoid synthesizing the false branch at all.

The key observation is that `allTwo` has a hidden invariant that the S-UNFOLDGEN' cannot make use of—if the step function ever returns false, the whole fold returns false. We can make this invariant available to the synthesizer by reinterpreting `allTwo` as an optional fold:

```
def allTwo (xs : List Nat) :=
  List.foldOpt (λ x () ⇒ if x = 2 then some () else none) (some ()) xs = some ()
```

Now β is Unit, and the step function simply checks if `x = 2` and, if not, fails. The invariant we wanted falls out of the definition of `List.foldOpt`—if the step function fails at any step, the whole computation fails.

All we need now is a synthesis rule for `List.foldOpt`:

$$\frac{\Gamma \vdash g : (b : \beta) \rightarrow \text{Gen}_{\text{Option } (\alpha \times \beta)} (P b)}{\Gamma \vdash \text{List.foldGen } g \ b : \text{Gen}_{\text{List } \alpha} (\lambda xs \Rightarrow \text{List.foldOpt } f \ z \ xs = \text{some } b)} \text{S-UNFOLDGEN''}$$

where $P b =$

$\lambda step \Rightarrow (step = \text{none} \wedge z = \text{some } b) \vee (\exists x \ b', \ step = \text{some } (x, b') \wedge f \ x \ b' = \text{some } b)$

This rule looks roughly the same as S-UNFOLDGEN', but it enforces the step-wise invariant we are interested in. Now, if we know the fold should return some, we can also assume each step should return some. We can use the new rule to obtain the following much simpler generator:

```
def genAllTwo : Gen (List Nat) :=
  List.foldGen (λ () ⇒ pick (pure none) (pure (some (2, ()))) ()
```


To reach our most general recursion synthesis rule, we target `List.accuOpt` rather than `List.foldOpt`. We define one more synthesis rule, `S-UNFOLDGEN`, that subsumes both `S-UNFOLDGEN'` and `S-UNFOLDGEN''`:

$$\frac{\Gamma \vdash g : (b : \beta) \rightarrow (s : \sigma) \rightarrow \text{GenOption } (\alpha \times (\sigma \times \beta)) \ (P \ b \ s)}{\Gamma \vdash \text{List.unfoldGen } g' \ (b, s) : \text{Gen}_{\text{List } \alpha} \ (\lambda \ xs \Rightarrow \text{List.accuOpt } st \ s \ f \ z \ xs = \text{some } b)} \quad \text{S-UNFOLDGEN}$$

where $P \ b \ s =$

$$\begin{aligned} & \lambda \ step \Rightarrow (step = \text{none} \wedge z \ s = \text{some } b) \vee (\exists \ x \ b', \ step = \text{some } (x, b') \wedge f \ x \ b' \ s = \text{some } b) \\ & g' \ b \ s = \\ & \quad g \ b \ s \ggg \lambda \ mstep \Rightarrow \\ & \quad \text{match } mstep \text{ with} \\ & \quad | \text{none} \Rightarrow \text{pure none} \\ & \quad | \text{some } (x, b') \Rightarrow \text{pure (some } (x, (b', st \ x \ s))) \end{aligned}$$

This allows for both a collector that passes information backward from the tail of the list and a state that passes data forward, and it allows for this short-circuiting behavior.

If we rewrite sorted from the beginning of this section as

$$\begin{aligned} \text{def sorted } xs &:= \\ \text{List.accuOpt } (\lambda \ x _ \Rightarrow x) \ 0 \ (\lambda \ x \ b \ lo \Rightarrow \text{if } lo \leq x \text{ then some } () \text{ else none}) \ (\text{some } ()) \ xs, \end{aligned}$$

we can use `S-UNFOLDGEN` to derive:

$$\begin{aligned} \text{def genSorted} &: \text{Gen (List Nat)} := \\ & \text{List.unfoldGen} \\ & \quad (\lambda \ (lo, ()) \Rightarrow \\ & \quad \quad \text{pick (pure none) (pick (greaterThan } lo) (\text{pure } lo) \ggg \lambda \ x \Rightarrow \text{pure (some } (x, (x, ()))))) \\ & \quad (0, ())) \end{aligned}$$

This generator behaves the same as one an expert user might write. At each step, the generator either ends the list or generates a new value x that is greater than or equal to lo , puts that value in the list, and continues with $lo = x$.

4.4 Tupling Predicates

The `S-UNFOLDGEN` rule works for predicates that are written as a single pass over a data structure, but sometimes predicates have multiple independent constraints. For example,

$$\text{def allTwoLengthK } (k : \mathbb{N}) \ (xs : \text{List Nat}) := \text{allTwo } xs \wedge \text{hasLengthK } k \ xs$$

combines two predicates that we have seen before into a single predicate.

We have two options for handling these situations. The first is to draw from the literature on program calculation and apply a *tupling* transformation [6, 41] to combine the conjuncts. These transformations were designed for functional program optimization, but they are a natural fit for this problem. Another is to try to adapt the “merging” procedure for inductive relations described by Prinz and Lampropoulos [44]. It turns out that, in the case of predicates written with `List.accuOpt`, these concepts coincide!

We introduce a transformation `tupleAccuOpt` which takes two predicates P and Q , each expressed with `accuOpt`, and combines them into a single predicate (also expressed with `accuOpt`) that computes $P \ xs \wedge Q \ xs$. The definition is surprisingly straightforward—the state and collector arguments are simply combined in a tuple and computed in parallel—so we do not replicate it here. This transformation means that our synthesis approach automatically benefits from the merging optimizations that users would need to apply manually in systems like QuickChick [44].

It is also worth pointing out that tupling is not the only useful program transformation that the early program calculation literature proposed for recursive programs. If we find that other transformations (e.g., fusion) are useful for the kinds of predicates that users care about, we could easily extend our pipeline with them.

4.5 Putting it All Together

The previous sections have introduced a lot of machinery; now we put that machinery together into a standardized workflow for handling predicates on inductive data.

The normalization workflow is presented at a high level in Figure 2.

- (1) Normalize the predicate to be of the form: $\lambda xs \Rightarrow F_1 xs = b_1 \wedge \dots \wedge F_n xs = b_n$.
- (2) Rewrite each F as a fold. Concretely, search for z and f satisfying the equations $F [] = z$ and $F (x :: xs) = f x (F xs)$. If these equations are satisfied, then F can be rewritten as $\text{List.fold } f z$; this is sometimes referred to as the “universal property of fold” [21, 31].
- (3) Rewrite each fold as an optional accumulation based on its return type. For example, `allTwo`, which collects an always-true boolean, would take a different path from `sorted` which collects a higher-order function.
- (4) Tuple the n different branches of the predicate together.

After normalization, we apply `S-UNFOLDGEN` to the final accumulation and obtain a generator by recursively synthesizing the step function. This workflow is automated through tactics in Lean, which we discuss in the next section.

4.6 Other Inductive Types

Everything in this section so far has revolved around the `List` data type, but there is actually nothing in the above workflow that is specific to lists. All recursive data types that are composed of products and sums admit operations that are analogous to `List.fold`, `List.accuOpt`, etc.⁴ This means that the pipeline from Figure 2 can be directly generalized to a wide range of recursive data structures.

The case studies in Section 6 required us to implement this pipeline (along with other utilities, e.g. for case splitting) for 5 data structures: lists, binary trees, STLC types, STLC terms, and stacks (from [20]). The process is mechanical, as the definitions follow the structure of the data type and its constructors. We believe it will be straightforward to automate via meta-programming or by leveraging *quotients of polynomial functors* (QPFs) [5]. We leave this engineering exercise as future work, as it does not change the feasibility of our approach.

5 PALAMEDES: SYNTHESIZING GENERATORS IN LEAN

In this section we describe PALAMEDES, our implementation of PBT generator synthesis as a library in the Lean theorem prover. We begin with an overview of the library (Section 5.1), then describe the synthesis algorithm in detail (Section 5.2).

5.1 Synthesizer as Library

PALAMEDES is implemented as a library in the Lean theorem prover. Because of Lean’s powerful meta-programming capabilities, this does not require compiler plugins or additional build steps as other synthesizers might: the synthesizer is implemented in standard Lean code.

Starting with a definition of the BST predicate

```
def isBST (t : Tree Nat) : Nat × Nat -> Bool := fun (lo, hi) =>
  match t with
```

⁴Technically speaking, these functions exist for all algebraic data types that arise as the least fixed-point of a traversable functor [32].

```

785 | leaf => true
786 | node l x r =>
787   (lo <= x && x <= hi) && isBST (lo, x - 1) l && isBST (x + 1, hi) r,

```

we can use PALAMEDES to synthesize a generator of valid BSTs:

```

789 def genBST (lo hi : Nat) : Gen (Tree Nat) := by
790   generator_search (fun t => isBST t (lo, hi) = true)

```

More generally, the user defines some predicate that they want their data to satisfy (such as `isBST` above), and they use the `generator_search` tactic to find an efficient generator that is sound and complete with respect to that predicate. The above call finds the following generator:⁵

```

795 def genBST (lo hi : Nat) : Gen (Tree Nat) :=
796   Tree.unfold
797     (fun ((), lo, hi) =>
798       if lo <= hi then
799         pick (pure none)
800           (choose lo hi (...) >>= fun x =>
801             pure (some ((((), lo, x - 1), x, (((), x + 1, hi))))))
802       else pure none) ((), lo, hi)

```

This code uses `Tree.unfold` to manage recursion and termination, but it operates exactly the same way as `genBST` from Section 2. As we discuss later in this section, the above code can be pasted directly into the user’s file; from there, they can modify it to tweak the distribution and make any other changes that they see fit.

While there are no proofs visible to the user in this workflow, they exist under the hood. The `generator_search` tactic proves that the generator it synthesizes is sound and complete with respect to the provided predicate and also assume-free. This is one of the major benefits of working inside a theorem prover like Lean—we do not need to rely on meta-arguments that our synthesis procedure is correct. Each synthesized generator comes with mechanized proofs that the generator is indeed appropriate for the user’s testing needs. (We provide alternative versions of `generator_search` that give users access to those proofs if needed.)

5.2 The Synthesis Algorithm

At a high level, “`generator_search P`” implements the following steps (expressed as a series of Lean tactic statements):

```

818 let g : CorrectGen P := by synthesize
819 let g' : CorrectGen P := by optimize g
820 let _ : AssumeFree g' := by prove_assume_free
821 exact g'

```

First, we synthesize an initial generator `g`. This generator has type `CorrectGen P` (`Genα P` from above). It is implemented in Lean as:

```

824 def CorrectGen {α : Type} (P : α -> Prop) := {g : Gen α // ∀ v, v ∈ [g] <-> P v}

```

In Lean this type is called a “subtype;” it is a dependent pair of a value and a proof that that value satisfies a given predicate. A value of type `CorrectGen P` is therefore a pair of a generator and a proof that that generator’s support is equivalent to `P`. Next, we optimize the generator with the rewrites described in Section 3.4; this procedure also produces a `CorrectGen`, ensuring that optimization has not changed the support of the generator. Finally, we prove that the generator is assume-free (Definition 3.3). We now describe these steps in more detail.

⁵For clarity, we manually added a pattern match on `(((), lo, hi))`—the actual synthesized code uses projections.

5.2.1 Step 1: Synthesize. The `synthesize` tactic solves a goal of type `CorrectGen P` by applying the synthesis rules described in [Section 3.2](#). The procedure uses `Aesop`, a tactic in Lean that performs best-first proof search [29]. `Aesop` takes a large list of Lean tactics and applies them in a loop; when all tactics fail to solve a particular sub-goal, the search backtracks to try a different route. Relying on `Aesop` in this way turned out to be extremely effective both in terms of results (in our case studies) and in terms of ease of implementation.

The rules we provide to `Aesop` mirror the ones described in [Section 3](#) and are given in [Table 1](#). The core synthesis rules each apply a function that builds a term of type `CorrectGen P` for some `P`; for example, `S-PICK` is defined as:

```
def s_pick (x : CorrectGen P) (y : CorrectGen Q) :
  CorrectGen (fun a => P a ∨ Q a) :=
  Subtype.mk (pick x y) (...)
```

This combines the rule for synthesizing `pick` with a proof (elided in the above definition) that the rule is valid with respect to the definition of the generator's support. Every synthesis rule is constructed this way, ensuring that the end-to-end synthesis process is correct by construction.

As discussed in [Section 3.2](#), the synthesis rules may not apply to a given goal directly. The `CONVERT` rule allows the synthesizer to change the predicate to an equivalent one, as long as we can prove the equivalence:

```
def convert (h : Q = P) (x : CorrectGen P) : CorrectGen Q := Subtype.mk x (...)
```

In practice, rather than allowing the synthesizer to apply `convert` arbitrarily, we instead combine it with the various generator builders. For example,

```
apply convert (by match_pick) (s_pick _ _)
```

uses `convert` with a pre-determined end goal (the conclusion of `s_pick`). This tactic tries to apply `s_pick`, using a custom tactic, `match_pick`, to prove that the rule actually applies.

To see this in action, consider the goal:

```
CorrectGen (fun a => (a = 3 ∨ a = 2) ∧ True)
```

Simply trying to apply `s_pick` here will fail, but applying the `convert` version allows `match_pick` to simplify away the no-op term and apply the rule.

As mentioned in [Section 3](#), the rule for `s_bind` is especially dependent on `convert`. At a given point in the synthesis process `s_bind` may apply in a number of different ways, corresponding to different orders that values may be generated in. Concretely, for the predicate $\exists x y, \dots$, it is not clear whether to apply the `bind` rule to `x` or `y` first. For this reason, we need multiple versions of the `s_bind` rule, each targeting a different generation order. For the examples in the following section, we only need two: the version using `match_bind 0` tries prioritizing the first existentially quantified variable and `match_bind 1` tries the second. More complex examples may require more of these rules. The `match_bind` tactic also rearranges conjuncts in the predicate to better match the pattern required by `s_bind`.

For inductive data types, we follow the same pattern as the core rules. Rules like

```
apply convert (by match_List_unfold) (List.s_unfold _)
```

implement the pipeline from [Figure 2](#), turning user-written recursive predicates into calls to `List.accuOpt`.

Finally, we implement the `S-SPLIT*` rules similarly to `s_bind`. The tactic

```
split_cases 1 Nat.split
```

for example, looks for the second `Nat` in the context and attempts to split it with `S-SPLITNAT`.

Table 1. The synthesis rules used to synthesize our core examples. Each rule has a precedence; rules with 100% precedence are tried first and never backtracked. Other rules are tried in order of precedence, and higher-precedence branches of the proof are explored first. Rules containing $\langle T \rangle$ are replicated once for each of the recursive data structures we consider.

Rule	Precedence
uncurry_intro	100%
assumption	100%
apply convert (by match_pure) s_pure	100%
apply convert (by match_pick) (s_pick _ _)	50%
apply convert (by match_bind 0) (s_bind _ _)	50%
apply convert (by match_bind 1) (s_bind _ _)	50%
apply convert (by match_greaterThan) s_greaterThan	50%
apply convert (by match_lessThan) (s_lessThan (by solve_lessThan))	50%
apply convert (by match_lessThan) s_lessThan_partial	50%
apply convert (by match_between) (s_between (by solve_between))	50%
apply convert (by match_between) s_between_partial	50%
apply convert (by match_elements) (s_elements (by solve_elements))	50%
apply convert (by match_elements) s_elements_partial	50%
apply convert (by match_<T>_unfold) (<T>.s_unfold _)	50%
split_cases 0 <T>.split	5%
split_cases 1 <T>.split	5%

5.2.2 Step 2: Optimize. The optimization tactic applies the optimization rules presented in Section 3.4. It is implemented as a meta-level function, operating directly on the AST of the generator. This makes it easy to write rules like rule (4), which matches on the body of a lambda abstraction. Being written at the meta level means that we cannot prove once and for all that optimization is correct in Lean (it is straightforward to prove on paper). Instead, we use proof automation to show that each optimized generator is equivalent to its unoptimized counterpart.

5.2.3 Step 3: (Optionally) Prove Assume-Free. While PALAMEDES does not categorically reject generators that contain assumes (they can still be useful in some cases), we would prefer to prove that generators are assume free. We use straightforward proof automation to attempt to prove that a generator does not make nontrivial use of the `assume` constructor; if we fail to prove this fact, we output a warning for the user.

5.2.4 Step 4: Render to the User. At this point, the user has a choice. If they think their predicate may change over the course of development, or if they simply want to simplify the codebase, they can choose to leave the call to `generator_search` as the definition of their generator. Lean will try to cache the generator when possible, and otherwise it will re-synthesize the generator when reloading the file. Users may also want to *render* the synthesized generator as a concrete program. Rendered generators are static, which means they are not re-synthesized when the file loads, and they can be further manipulated and tuned by the user. The user might, for example, want to add weights that bias the distribution or add size bounds to ensure generated values do not get too big.⁶ When the user wants to render a generator, they can use a variant of the search tactic: `generator_search?`. This version does the same synthesis procedure and then provides the user

⁶This may make the generator incomplete, so do not do it during synthesis, but an expert user might decide the incompleteness is worth it, on a case-by-case basis.

with a “try this” widget ⁷ in their editor. Clicking on the hint pastes the full text of the generator into their file, which they can then edit as normal.

We make a couple of choices during the synthesis process to make rendering possible, both of which address subtle technical details. First, we mark all `CorrectGen` constructors as `reducible`. This tells Lean’s evaluator that those definitions can (and should) be reduced during elaboration. Second, we ensure that all `CorrectGen` constructors are of the form `Subtype.mk g pf`, where `g` is a plain generator and `pf` is a proof about that generator. This means that, when we project out the generator `g`, it is guaranteed to be independent of the proof. In some other type theories [2], all terms of type `CorrectGen` could be guaranteed to reduce to a call to `Subtype.mk`; but in Lean, a term might reduce to a type cast (i.e., `h ► e`) from which we cannot readily extract the generator component. These choices do slightly complicate our synthesis procedure—it means we need to use Aesop somewhat cautiously—so we hope to be able to relax these requirements in the future.

5.2.5 Step 5: Interpret and Run the Generator. The final step of the process is to actually test code with the generator by *sampling* from it. As discussed in Section 3, the generators we synthesize are data structures, not programs, so they cannot be run directly. The sampling interpreter gives meaning to generators as maps from random seeds to values, mirroring Definition 3.1. In order to be consistent with a generator’s support, the sampling interpretation needs to handle backtracking and non-termination carefully. Specifically, it re-samples values the case of failure (e.g., from `assume`), and functions wrapped by `indexed` are given increasing fuel after running out.

6 EVALUATION

In this section, we evaluate our synthesis algorithm by using PALAMEDES to synthesize over 40 benchmark generators. Our synthesis procedure is not complete (see Section 7), but we demonstrate that it works for a wide range of interesting and useful predicates.

We answer two key research questions:

RQ1 Can PALAMEDES synthesize a variety of generators in an acceptable time-frame?

RQ2 Are the generators that PALAMEDES synthesizes comparable to ones that expert users write?

We answer the first of these questions in Section 6.1 and the second in Section 6.2.

6.1 Benchmark Overview and Timing

To answer **RQ1**, we benchmark PALAMEDES on a wide range of predicates. Some demonstrate specific aspects of our synthesis algorithm (including low-level examples that appear in the text above), while others are drawn from the literature on PBT [25, 27, 44].

A selection of the benchmark predicates are presented in Table 2; the rest appear in Appendix C. The synthesis times range from around 40 milliseconds for very basic examples to up to 34 seconds for a complex example (AVL trees) with multiple nested case splits. It takes under 2 seconds to synthesize `genBST` and under 4 seconds for the generator for well-typed STLC terms.

We argue that these synthesis times are well within the acceptable range for a procedure like this. Our assumption, based on prior work [16], is that developers tend to run their tests far more frequently than they change their definitions. That means that these modest synthesis costs will be amortized across many test runs. Critically, any costs incurred by search-based approaches to generation are paid every time the developer runs their tests.

Backtracking Generators. The vast majority of our synthesized generators are proved `assume-free`; only three—two versions of AVL trees and a demonstration example—have `assume`.

⁷<https://lean-lang.org/documentation/widgets/>

Table 2. Benchmark predicates and synthesis times. The value being generated is \mathbf{v} ; other variables are universally quantified unless specified. External definitions (e.g., `isBST`) are presented in [Appendix B](#). Generators above the line are assume-free; the ones below are not. Benchmarks were run on an M1 MacBook Pro with 8 cores and 16GB of memory using Lean v4.21.0. Times are averaged over 30 runs; means are presented with standard deviations in parentheses. All times are in seconds.

Predicate	Type	Time (s)
$\mathbf{v} = 2$	Nat	0.04 (0.01)
$2 = \mathbf{v}$	Nat	0.04 (0.00)
$\mathbf{v} = 2 \vee \mathbf{v} = 5$	Nat	0.08 (0.00)
$\mathbf{v} = 2 \vee \mathbf{v} = 5 \wedge \text{True}$	Nat	0.08 (0.00)
$\exists a, a = 3 \wedge \mathbf{v} = a + 1$	Nat	0.04 (0.00)
$5 \leq \mathbf{v} \wedge \mathbf{v} \leq 10$	Nat	0.08 (0.00)
$\mathbf{v} > 5$	Nat	0.07 (0.00)
$\mathbf{v} = 0 \vee \text{lo} \leq \mathbf{v} \wedge \mathbf{v} \leq \text{hi}$	Nat	0.14 (0.00)
<code>isAllTwos</code> $\mathbf{v} = \text{true}$	List Nat	0.84 (0.01)
<code>isAllTwosEvenLen</code> $\mathbf{v} = \text{true}$	List Nat	2.76 (0.02)
<code>isEvenLen</code> $\mathbf{v} = \text{true}$	List Nat	2.22 (0.02)
<code>isIncreasingByOne</code> $\mathbf{v} = \text{true}$	List Nat	1.44 (0.01)
<code>List.length</code> $\mathbf{v} = k$	List Nat	1.89 (0.01)
<code>isLengthKAllTwos</code> $k \ \mathbf{v} = \text{true}$	List Nat	2.37 (0.01)
<code>isSortedBetween</code> $\mathbf{v} (\text{lo}, \text{hi}) = \text{true}$	List Nat	1.72 (0.02)
<code>isTrue</code> $\mathbf{v} = \text{true}$	List Nat	2.21 (0.01)
<code>isAllTwos</code> $\mathbf{v} = \text{true}$	Tree Nat	0.61 (0.01)
<code>isBST</code> $\mathbf{v} (\text{lo}, \text{hi}) = \text{true}$	Tree Nat	1.86 (0.01)
<code>isComplete</code> $\mathbf{v} \ n = \text{true}$	Tree Nat	2.35 (0.02)
<code>isIncreasingByOne</code> $\mathbf{v} = \text{true}$	Tree Nat	1.57 (0.02)
<code>isNonempty</code> $\mathbf{v} = \text{true}$	Tree Nat	1.66 (0.04)
<code>isGoodStack</code> $\mathbf{v} \ n = \text{true}$	Stack	5.39 (0.05)
<code>isWellScoped</code> $\mathbf{v} \ 0 = \text{true}$	Term	2.85 (0.03)
<code>isWellTyped</code> $\Gamma \ \mathbf{v}$	Term	3.66 (0.03)
$\text{lo} \leq \mathbf{v} \wedge \mathbf{v} \leq \text{hi}$	Nat	0.07 (0.00)
<code>isAVL height lo hi</code> $\mathbf{v} = \text{true}$	Tree Nat	34.44 (0.22)

Generating AVL trees without backtracking is tricky. A common failure mode, even for many hand-written AVL tree generators, is running out of valid values for a node before that branch is deep enough to be balanced. Indeed, the AVL tree generator that Prinz and Lampropoulos [44] present in their paper on merging inductive relations also backtracks in this (relatively unlikely) case. However, there are ways to write AVL tree generators that do not backtrack, and we replicate two in [Appendix F](#). The first version more carefully selects values for each node. Rather than choose a value between lo and hi , it leaves $2^{\wedge} (\text{height} - 1)$ values on either side of the range to ensure that the next step of generation will not run out of values. Ideally, PALAMEDES would be able to find a generator like this, but it would require automatically rewriting the predicate far beyond what is currently possible. The other option generates AVL trees by repeatedly inserting values into an empty tree. We would never expect PALAMEDES to find this generator—it would require far too much domain knowledge—but it is likely the approach a PBT user would reach for in practice.

```

1030 def genWellTyped (Γ : List Ty) : Gen Term := by
1031   let τ <- arbTy; Term.unfold
1032     (fun (τ, Γ) => do
1033       let step <- caseTy τ
1034       (fun () =>
1035         pick
1036           (pure TermF.unitStep)
1037           (if (Γ.indexesOf .unit).length > 0 then
1038             pick (do let n <- elements (Γ.indexesOf .unit) (...))
1039               pure (TermF.varStep n))
1040             (do let τ' <- arbTy
1041               pure (TermF.appStep (.arrow τ' .unit) τ'))
1042           else do
1043             let τ' <- arbTy; pure (TermF.appStep (.arrow τ' .unit) τ'))
1044       (fun τ1 τ2 () => -- τ = .arrow τ1 τ2
1045         if (Γ.indexesOf (.arrow τ1 τ2)).length > 0 then
1046           pick (do let n <- elements (Γ.indexesOf (.arrow τ1 τ2)) (...))
1047             pure (TermF.varStep n))
1048           (pick
1049             (pure (TermF.absStep τ1 τ2))
1050             (do let τ' <- arbTy
1051               pure (TermF.appStep (.arrow τ' (.arrow τ1 τ2)) τ'))))
1052       else
1053         pick (pure (TermF.absStep τ1 τ2))
1054         (do let τ' <- arbTy
1055           pure (TermF.appStep (.arrow τ' (.arrow τ1 τ2)) τ'))))
1056   match step with
1057   | TermF.unitStep => pure TermF.unitStep
1058   | TermF.varStep n => pure (TermF.varStep n)
1059   | TermF.absStep τ b => pure (TermF.absStep τ (b, τ :: Γ))
1060   | TermF.appStep b1 b2 => pure (TermF.appStep (b1, Γ) (b2, Γ)) (τ, Γ)

```

Fig. 3. Synthesized generator for well-typed STLC terms. Comments were added manually, and some variables were renamed for clarity.

6.2 Synthesized vs. Handwritten Generators

Now we turn to RQ2, exploring how the generators that PALAMEDES produces compare to ones that expert users might write.

As our main case study, we focus on the generator for well-typed STLC terms that PALAMEDES synthesized, shown in Figure 3. Following the S-BIND rule, `genWellTyped` picks a random type and then generates a random term of that type—this happens to be the method that prior work [37] suggests for a generator like this. The bulk of the generator uses `Term.unfold` to generate terms step by step. Each step, the generator inspects the type and produces a `TermF` that determines the next data constructor that should be generated. In both cases (unit or arrow), the generator might choose to select a variable from the context, generate a function application, or generate a type-specific term—the unit value or a lambda abstraction. Once the next step is chosen, the type of the continuation is packaged up with a modified context and used to seed the next step.

This generator is logically the same as one that an expert user might write to satisfy the same predicate, although the code would be a bit different. An expert user would reduce repetition by waiting to split on τ until it is necessary to do so. They would likely also avoid the match statement at the end, which is an artifact of the synthesis process. Still, the asymptotic performance of the generators would be the same—neither would backtrack, and both would contain the same basic clauses. We provide a hand-written STLC generator that the authors wrote in Appendix D.

We can also compare this generator to ones described in the literature, in particular the STLC generators from the Etna PBT evaluation framework [49]. These generators are available online, but we replicate the Haskell and Rocq versions of the generator in [Appendix E](#) for reference. Focusing on the Haskell version first, it has some code to carefully control the sizes of generated values that the synthesized version does not—we discuss this limitation of PALAMEDES in [Section 7](#). But the high-level control flow of the Haskell version `genExpr` is the same as PALAMEDES’s `genWellTyped`. Interestingly, the QuickCheck version of the STLC generator in the Etna suite, which is automatically derived from an inductive relation specifying well-typedness [27], is *not* logically equivalent to the other versions. In the case where the context does not contain a variable of the appropriate type, the QuickCheck version may try to generate a bound variable and then backtrack. This is not terribly wasteful, but it is sub-optimal.

In [Appendix F](#), we show 5 more generators that PALAMEDES synthesized, side-by-side with ones that we wrote manually. For all but the last one (AVL trees, which we discuss at length above), there were no functional differences between the generators we wrote and the synthesized ones.

To summarize, when PALAMEDES finds an assume-free generator for a predicate, it is faithful to the kinds of generators that expert users write for those same predicates.

7 LIMITATIONS

Our approach has some limitations that it is important to highlight.

First, our synthesis algorithm is necessarily partial; there are many Lean predicates that PALAMEDES cannot handle. For example, PALAMEDES’s tupling and fold normalization pipeline is not currently powerful enough to synthesize a generator for red-black trees [27]. Additionally, since all uses of `assume` need to be explicitly allowed by the user (e.g., in the definitions of `S-CHOOSE` and `S-ELEMENTS`), new generators that need assumptions cannot be automatically synthesized. This is why PALAMEDES cannot currently synthesize lists with unique or duplicated elements [25]. Luckily, PALAMEDES is extremely extensible—users can add new proof automation, library generators, and synthesis rules—so these limitations can be mitigated over time. In [Section 9](#), we discuss plans to extend the set of predicates that PALAMEDES can handle in a variety of ways.

There is also a classic PBT generator feature that PALAMEDES does not yet incorporate: *sizes*. QuickCheck [11] generators have built-in ways to vary the size of generated values over the course of generation, which we do not include in our `Gen` type. Sizes are compatible with our approach—prior work [39] gives a semantics for internal sizing that plays well with our definitions—but we do not yet have a heuristic for how the synthesizer should use it. That said, there are two ways to control the sizes of generated values in PALAMEDES. Users can augment their specifications with explicit size constraints. For example, `isLengthKAllTwos` allows the user to control the length of generated lists by changing the argument `k`; the tupling transformation means this can be done by simply conjoining an extra predicate. Alternatively, users can manually modify their generators after synthesis to add size control.

8 RELATED WORK

In this section we discuss related work on constrained random generation, on generator synthesis, and on deductive synthesis more generally.

8.1 The Constrained Random Generation Problem

As we discuss at length in [Section 2.3](#), many approaches to the constrained random generation problem have been proposed over the years [10, 17, 46, 48, 50, 55]. These all solve a slightly different problem from the one solved by PALAMEDES: they actively guide generation towards valid values *during testing*, rather than searching for an appropriate generator ahead of time. In situations

where a developer writes their generator once and then runs it many times—which is common in PBT [16]—these approaches may not scale as well as ones that synthesize generators ahead of time.

8.2 The Constrained Generator Synthesis Problem

As discussed in the introduction, generating inputs that satisfy some given constraint is a well explored area of research. While most existing work opts to search for inputs during generation, a few papers do tackle the constrained generator synthesis problem, more or less as we present it.

In the Rocq theorem prover [51], QuickChick provides an automated mechanism for deriving generators that are sound and complete with respect to predicates defined as inductive relations [27, 38]. This mechanism is incredibly powerful, efficiently deriving high quality generators, but there are clear trade-offs between QuickChick’s approach and ours. On one hand, QuickChick’s deriver is much faster, and it is also better established with more stable and predictable behavior. On the other hand, PALAMEDES is more flexible; it does not require predicates to be expressed in the rigid structure of inductive relations. Predicates like `List.length xs = k`, which PALAMEDES can synthesize a generator for directly, would need to be re-encoded before QuickChick’s deriver could apply. A more functional approach also means that PALAMEDES applies more naturally outside of theorem provers (see Section 9). Additionally, PALAMEDES works harder to avoid backtracking. In cases like the Etna example discussed in Section 6.2, QuickChick derive backtracking generators without warning the user, whereas PALAMEDES is much more explicit. Ultimately, we hope that Lean’s ecosystem eventually includes both kinds of automation—QuickChick style derivers for inductive relations (maybe better backtracking control) and PALAMEDES for functional predicates.

Similarly, Isabelle’s implementation of QuickCheck provides an automated method for deriving enumerators for preconditions [8, 9]. The approach is a spiritual precursor to QuickChick’s, focusing on a syntactic subset of Isabelle that can be represented as Horn clauses. One important design difference between Isabelle’s approach and PALAMEDES is that the generator is never materialized in Isabelle; the deriving mechanism does not allow the user to inspect, modify, or tune the enumerators.

The most recent addition to this space is Cobb [25], a technique that uses program repair to turn incomplete generators into complete ones. Cobb can handle some predicates that PALAMEDES does not: as discussed in Section 7, PALAMEDES cannot yet synthesize generators for lists with unique or duplicate entries and red-black trees. However, for predicates that both approaches handle PALAMEDES is much faster. Cobb takes over 200 seconds to synthesize a BST generator, compared to PALAMEDES’s less than 2 seconds. Cobb also assumes that users already have at least a sketch (i.e., the basic control flow) of the generator that they want, which requires more up-front effort.

8.3 Deductive Program Synthesis

Deductive program synthesis takes a specification in some logic and a set of inference rules, and searches for a proof that the specification holds; steps in the proof immediately translate to a program proved correct by the proof. The simplest version of deductive synthesis is *type-directed synthesis* [15, 19, 35, 36], where the target proof tree is an application of typing rules. This has been extended to expressive type systems like refinement types [42] and semantic types [18]. Several systems extend this approach to other logics [12, 24, 45], including separation logic [13, 22, 43].

PALAMEDES builds on the extensive body of work in deductive synthesis, particularly on SuSLik [43] and its followup work [13, 22]. However, PALAMEDES targets a domain that has not yet been explored by prior work, and its implementation is different in key ways. Unlike prior work, PALAMEDES relies on Lean for many aspects of proof search; for example, avoiding explicit reasoning about termination and solving auxiliary lemmas with tools like Aesop [29]. Additionally, with the exception of Fiat [12], prior work does not produce mechanized proofs of correctness. PALAMEDES’s deduction rules are proved correct within the Lean theorem prover, and those proofs are combined

to produce a proof that the final generator is correct. This strengthens the guarantee that a resulting program is correct by construction.

9 CONCLUSION AND FUTURE WORK

In this work, we address the constrained generator synthesis problem. We give an algorithm for synthesizing generators that are sound and complete with respect to a predicate, including generators that produce recursive data structures like lists and trees. Our approach combines prior work on deductive synthesis, functional programming, theorem provers, and more into a technique that we think will have significant positive impact in the realm of property-based testing.

We close with some ideas for future work.

Tuning Generator Distributions. The generators produced by PALAMEDES are guaranteed to produce the right set of values, but they may produce those values with a suboptimal *distribution*. For example, the predicate $a = 1 \vee a = 2 \vee a = 3 \vee a = 4$ will yield the generator

```
pick (pure 1) (pick (pure 2) (pick (pure 3) (pure 4)))
```

which will produce 1 with probability 0.5, 2 with probability 0.25, and 3 and 4 with probability 0.125. This bias towards 1 is not something the user expressed that they wanted per se, it is simply an artifact of operator associativity.

Luckily, there are ways to address this problem. As a naïve solution, we could implement an optimization pass that re-associates picks to prefer more balanced trees. For practical cases, this might actually be sufficient. But we need not stop there: recent work has shown that probabilistic programming languages, in particular one called Loaded Dice, can be used to automatically tune generators to user-specified distributions [53]. We plan to implement a translation from our generators into Loaded Dice, giving much more comprehensive control over generator distributions. Along the way, we hope to improve on the way Loaded Dice deals with recursion—currently it unrolls loops before training, but we suspect that the extra structure provided by unfolds will enable a more robust approach.

Correct Generators for Everyday Developers. Implementing PALAMEDES as a Lean library has myriad benefits, but it has one major downside: theorem provers are still inaccessible to everyday software developers, so the tool in its current form is unlikely to see broad adoption. However, we see a clear path towards impact in the software engineering industry, by using PALAMEDES as the backend of more user-focused tools. As a first step, we will target Python. We plan to (1) embed a subset of Python’s semantics in Lean, (2) compile Python predicates to that sub-language, (3) synthesize appropriate generators, and (4) render the synthesized generators as Hypothesis [30] strategies. If successful, we hope to push this paradigm even further, using PALAMEDES as a backend for synthesizing generators in as many major languages and PBT frameworks as possible.

Better Automation and Algorithmic Improvements. As we demonstrate in Section 6, PALAMEDES is already flexible enough to synthesize a wide range of generators, but the algorithm may be further improved by ongoing improvements to Lean’s proof automation. For example, Lean now has tactics for automating proof search with SMT [34] and e-graph rewriting [47], and large-language-model-based proof automation is an active area of study [14, 23, 28, 52, 57, 58].

These approaches dovetail nicely with our current infrastructure—they could be used to implement more powerful versions of the `match_*` tactics and some could even to replace Aesop entirely as the engine for the core synthesis loop. We expect PALAMEDES to naturally grow in power over time, working in concert with the Lean community’s commitment to proof automation.

REFERENCES

- [1] 2023. How We Built Cedar with Automated Reasoning and Differential Testing. <https://www.amazon.science/blog/how-we-built-cedar-with-automated-reasoning-and-differential-testing>.
- [2] Thorsten Altenkirch and Conor McBride. 2006. Towards Observational Type Theory. (2006). <https://personal.cis.strath.ac.uk/conor.mcbride/ott.pdf>
- [3] Thomas Arts, John Hughes, Joakim Johansson, and Ulf Wiger. 2006. Testing Telecoms Software with Quviq QuickCheck. In *Proceedings of the 2006 ACM SIGPLAN Workshop on Erlang (ERLANG '06)*. Association for Computing Machinery, New York, NY, USA, 2–10. <https://doi.org/10.1145/1159789.1159792>
- [4] Thomas Arts, John Hughes, Ulf Norell, and Hans Svensson. 2015. Testing AUTOSAR Software with QuickCheck. In *2015 IEEE Eighth International Conference on Software Testing, Verification and Validation Workshops (ICSTW)*. 1–4. <https://doi.org/10.1109/ICSTW.2015.7107466>
- [5] Jeremy Avigad, Mario Carneiro, and Simon Hudon. 2019. Data Types as Quotients of Polynomial Functors. In *10th International Conference on Interactive Theorem Proving (ITP 2019) (Leibniz International Proceedings in Informatics (LIPIcs), Vol. 141)*, John Harrison, John O’Leary, and Andrew Tolmach (Eds.). Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl, Germany, 6:1–6:19. <https://doi.org/10.4230/LIPIcs.ITP.2019.6> ISSN: 1868-8969.
- [6] R. S. Bird. 1984. Using Circular Programs to Eliminate Multiple Traversals of Data. *Acta Informatica* 21, 3 (Oct. 1984), 239–250. <https://doi.org/10.1007/BF00264249>
- [7] James Bornholt, Rajeev Joshi, Vytautas Astrauskas, Brendan Cully, Bernhard Kragl, Seth Markle, Kyle Sauri, Drew Schleit, Grant Slatten, Serdar Tasiran, Jacob Van Geffen, and Andrew Warfield. 2021. Using Lightweight Formal Methods to Validate a Key-Value Storage Node in Amazon S3. In *Proceedings of the ACM SIGOPS 28th Symposium on Operating Systems Principles (SOSP '21)*. Association for Computing Machinery, New York, NY, USA, 836–850. <https://doi.org/10.1145/3477132.3483540>
- [8] Lukas Bulwahn. 2012. The New Quickcheck for Isabelle - Random, Exhaustive and Symbolic Testing under One Roof. In *2nd International Conference on Certified Programs and Proofs (CPP) (Lecture Notes in Computer Science, Vol. 7679)*. Springer, 92–108. <https://www.irisa.fr/celtique/genet/ACF/Bibliolabelle/quickcheckNew.pdf>
- [9] Lukas Bulwahn. 2012. Smart Testing of Functional Programs in Isabelle. In *18th International Conference on Logic for Programming, Artificial Intelligence, and Reasoning (LPAR) (Lecture Notes in Computer Science, Vol. 7180)*. Springer, 153–167. <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.229.1307&rep=rep1&type=pdf>
- [10] Koen Claessen, Jonas Duregård, and Michal H. Palka. 2015. Generating Constrained Random Data with Uniform Distribution. *J. Funct. Program.* 25 (2015). <https://doi.org/10.1017/S0956796815000143>
- [11] Koen Claessen and John Hughes. 2000. QuickCheck: A Lightweight Tool for Random Testing of Haskell Programs. In *Proceedings of the Fifth ACM SIGPLAN International Conference on Functional Programming (ICFP '00), Montreal, Canada, September 18-21, 2000*, Martin Odersky and Philip Wadler (Eds.). ACM, Montreal, Canada, 268–279. <https://doi.org/10.1145/351240.351266>
- [12] Benjamin Delaware, Clément Pit-Claudel, Jason Gross, and Adam Chlipala. 2015. Fiat: Deductive Synthesis of Abstract Data Types in a Proof Assistant. In *Proceedings of the 42nd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (Mumbai, India) (POPL '15)*. Association for Computing Machinery, New York, NY, USA, 689–700. <https://doi.org/10.1145/2676726.2677006>
- [13] Jonáš Fiala, Shachar Itzhaky, Peter Müller, Nadia Polikarpova, and Ilya Sergey. 2023. Leveraging Rust Types for Program Synthesis. *Proc. ACM Program. Lang.* 7, PLDI, Article 164 (June 2023), 24 pages. <https://doi.org/10.1145/3591278>
- [14] Emily First, Markus N. Rabe, Talia Ringer, and Yuriy Brun. 2023. Baldur: Whole-Proof Generation and Repair with Large Language Models. <https://doi.org/10.48550/arXiv.2303.04910> arXiv:2303.04910 [cs].
- [15] Jonathan Frankle, Peter-Michael Osera, David Walker, and Steve Zdancewic. 2016. Example-directed synthesis: a type-theoretic interpretation. In *Proceedings of the 43rd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (St. Petersburg, FL, USA) (POPL '16)*. Association for Computing Machinery, New York, NY, USA, 802–815. <https://doi.org/10.1145/2837614.2837629>
- [16] Harrison Goldstein, Joseph W. Cutler, Daniel Dickstein, Benjamin C. Pierce, and Andrew Head. 2024. Property-Based Testing in Practice. In *Proceedings of the IEEE/ACM 46th International Conference on Software Engineering (ICSE '24, Vol. 187)*. Association for Computing Machinery, Lisbon, Portugal, 1–13. <https://doi.org/10.1145/3597503.3639581>
- [17] Harrison Goldstein and Benjamin C. Pierce. 2022. Parsing Randomness. *Proceedings of the ACM on Programming Languages* 6, OOPSLA2 (Oct. 2022), 128:89–128:113. <https://doi.org/10.1145/3563291>
- [18] Zheng Guo, David Cao, Davin Tjong, Jean Yang, Cole Schlesinger, and Nadia Polikarpova. 2022. Type-directed program synthesis for RESTful APIs. In *Proceedings of the 43rd ACM SIGPLAN International Conference on Programming Language Design and Implementation (San Diego, CA, USA) (PLDI 2022)*. Association for Computing Machinery, New York, NY, USA, 122–136. <https://doi.org/10.1145/3519939.3523450>
- [19] Tihomir Gvero, Viktor Kuncak, Ivan Kuraj, and Ruzica Piskac. 2013. Complete completion using types and weights. In *Proceedings of the 34th ACM SIGPLAN Conference on Programming Language Design and Implementation (Seattle,*

- Washington, USA) (*PLDI '13*). Association for Computing Machinery, New York, NY, USA, 27–38. <https://doi.org/10.1145/2491956.2462192>
- [20] Catalin Hritcu, John Hughes, Benjamin C. Pierce, Antal Spector-Zabusky, Dimitrios Vytiniotis, Arthur Azevedo de Amorim, and Leonidas Lampropoulos. 2013. Testing noninterference, quickly. *SIGPLAN Not.* 48, 9 (Sept. 2013), 455–468. <https://doi.org/10.1145/2544174.2500574>
- [21] Graham Hutton. 1999. A tutorial on the universality and expressiveness of fold. *Journal of Functional Programming* 9, 4 (July 1999), 355–372. <https://doi.org/10.1017/S0956796899003500>
- [22] Shachar Itzhaky, Hila Peleg, Nadia Polikarpova, Reuben N. S. Rowe, and Ilya Sergey. 2021. Cyclic program synthesis. In *Proceedings of the 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation (PLDI 2021)*. Association for Computing Machinery, New York, NY, USA, 944–959. <https://doi.org/10.1145/3453483.3454087>
- [23] Saketh Ram Kasibatla, Arpan Agarwal, Yuriy Brun, Sorin Lerner, Talia Ringer, and Emily First. 2024. Cobblestone: Iterative Automation for Formal Verification. <https://doi.org/10.48550/arXiv.2410.19940> arXiv:2410.19940 [cs].
- [24] Etienne Kneuss, Ivan Kuraj, Viktor Kuncak, and Philippe Suter. 2013. Synthesis modulo recursive functions. In *Proceedings of the 2013 ACM SIGPLAN International Conference on Object Oriented Programming Systems Languages & Applications (Indianapolis, Indiana, USA) (OOPSLA '13)*. Association for Computing Machinery, New York, NY, USA, 407–426. <https://doi.org/10.1145/2509136.2509555>
- [25] Patrick LaFontaine, Zhe Zhou, Ashish Mishra, Suresh Jagannathan, and Benjamin Delaware. 2025. We've Got You Covered: Type-Guided Repair of Incomplete Input Generators. <https://doi.org/10.48550/arXiv.2504.06421> arXiv:2504.06421 [cs].
- [26] Leonidas Lampropoulos, Diane Gallois-Wong, Catalin Hritcu, John Hughes, Benjamin C. Pierce, and Li-yao Xia. 2017. Beginner's Luck: A Language for Property-Based Generators. *Proceedings of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages, POPL 2017, Paris, France, January 18-20, 2017* (2017), 114–129.
- [27] Leonidas Lampropoulos, Zoe Paraskevopoulou, and Benjamin C. Pierce. 2017. Generating Good Generators for Inductive Relations. *Proceedings of the ACM on Programming Languages* 2, POPL (2017), 1–30.
- [28] Zhaoyu Li, Jialiang Sun, Logan Murphy, Qidong Su, Zenan Li, Xian Zhang, Kaiyu Yang, and Xujie Si. 2024. A Survey on Deep Learning for Theorem Proving. <https://doi.org/10.48550/arXiv.2404.09939> arXiv:2404.09939 [cs].
- [29] Jannis Limperg and Asta Halkjær From. 2023. Aesop: White-Box Best-First Proof Search for Lean. In *Proceedings of the 12th ACM SIGPLAN International Conference on Certified Programs and Proofs (CPP 2023)*. Association for Computing Machinery, New York, NY, USA, 253–266. <https://doi.org/10.1145/3573105.3575671>
- [30] David R MacIver, Zac Hatfield-Dodds, et al. 2019. Hypothesis: A New Approach to Property-Based Testing. *Journal of Open Source Software* 4, 43 (2019), 1891.
- [31] Grant Reynold Malcolm. 1990. Algebraic data types and program transformation. (1990).
- [32] Conor McBride and Ross Paterson. 2008. Applicative Programming with Effects. *Journal of functional programming* 18, 1 (2008), 1–13. <https://doi.org/10.1017/S0956796807006326>
- [33] Agustin Mista and Alejandro Russo. 2021. Deriving Compositional Random Generators. In *Proceedings of the 31st Symposium on Implementation and Application of Functional Languages (IFL '19)*. Association for Computing Machinery, New York, NY, USA, 1–12. <https://doi.org/10.1145/3412932.3412943>
- [34] Abdalrhman Mohamed, Tomaz Mascarenhas, Harun Khan, Haniel Barbosa, Andrew Reynolds, Yicheng Qian, Cesare Tinelli, and Clark Barrett. 2025. Lean-SMT: An SMT tactic for discharging proof goals in Lean. <https://doi.org/10.48550/arXiv.2505.15796> arXiv:2505.15796 [cs].
- [35] Peter-Michael Osera. 2019. Constraint-Based Type-Directed Program Synthesis. In *Proceedings of the 4th ACM SIGPLAN International Workshop on Type-Driven Development (TyDe 2019)*. Association for Computing Machinery, New York, NY, USA, 64–76. <https://doi.org/10.1145/3331554.3342608>
- [36] Peter-Michael Osera and Steve Zdancewic. 2015. Type-and-Example-Directed Program Synthesis. *ACM SIGPLAN Notices* 50, 6 (June 2015), 619–630. <https://doi.org/10.1145/2813885.2738007>
- [37] Michal H. Palika, Koen Claessen, Alejandro Russo, and John Hughes. 2011. Testing an Optimising Compiler by Generating Random Lambda Terms. In *Proceedings of the 6th International Workshop on Automation of Software Test (AST '11)*. ACM, New York, NY, USA, 91–97. <https://doi.org/10.1145/1982595.1982615>
- [38] Zoe Paraskevopoulou, Aaron Eline, and Leonidas Lampropoulos. 2022. Computing Correctly with Inductive Relations. In *Proceedings of the 43rd ACM SIGPLAN International Conference on Programming Language Design and Implementation (PLDI 2022)*. Association for Computing Machinery, New York, NY, USA, 966–980. <https://doi.org/10.1145/3519939.3523707>
- [39] Zoe Paraskevopoulou, Cătălin Hritcu, Maxime Dénès, Leonidas Lampropoulos, and Benjamin C. Pierce. 2015. Foundational Property-Based Testing. In *Interactive Theorem Proving (Lecture Notes in Computer Science)*, Christian Urban and Xingyuan Zhang (Eds.). Springer International Publishing, Cham, 325–343. https://doi.org/10.1007/978-3-319-22102-1_22

- [40] Alberto Pardo. 2003. Generic Accumulations. In *Generic Programming: IFIP TC2 / WG2.1 Working Conference Programming July 11–12, 2002, Dagstuhl, Germany*, Jeremy Gibbons and Johan Jeuring (Eds.). Springer US, Boston, MA, 49–78. https://doi.org/10.1007/978-0-387-35672-3_3
- [41] Alberto Pettorossi, Enrico Pietropoli, and Maurizio Proietti. 1993. The Use of the Tupling Strategy in the Development of Parallel Programs. In *Parallel Algorithm Derivation and Program Transformation*, Robert Paige, John Reif, and Ralph W. Thatcher (Eds.). Springer US, Boston, MA, 111–151. https://doi.org/10.1007/978-0-585-27330-3_4
- [42] Nadia Polikarpova, Ivan Kuraj, and Armando Solar-Lezama. 2016. Program synthesis from polymorphic refinement types. *SIGPLAN Not.* 51, 6 (June 2016), 522–538. <https://doi.org/10.1145/2980983.2908093>
- [43] Nadia Polikarpova and Ilya Sergey. 2019. Structuring the Synthesis of Heap-Manipulating Programs. *SuSLik, Tool Implementation for Article: Structuring the Synthesis of Heap-Manipulating Programs* 3, POPL (Jan. 2019), 72:1–72:30. <https://doi.org/10.1145/3290385>
- [44] Jacob Prinz and Leonidas Lampropoulos. 2023. Merging Inductive Relations. *Reproduction Package for "Merging Inductive Relations"* 7, PLDI (June 2023), 178:1759–178:1778. <https://doi.org/10.1145/3591292>
- [45] Xiaokang Qiu and Armando Solar-Lezama. 2017. Natural synthesis of provably-correct data-structure manipulations. *Proc. ACM Program. Lang.* 1, OOPSLA, Article 65 (Oct. 2017), 28 pages. <https://doi.org/10.1145/3133889>
- [46] Sameer Reddy, Caroline Lemieux, Rohan Padhye, and Koushik Sen. 2020. Quickly Generating Diverse Valid Test Inputs with Reinforcement Learning. In *Proceedings of the ACM/IEEE 42nd International Conference on Software Engineering (ICSE '20)*. Association for Computing Machinery, New York, NY, USA, 1410–1421. <https://doi.org/10.1145/3377811.3380399>
- [47] Marcus Rossel and Andrés Goens. 2024. Bridging Syntax and Semantics of Lean Expressions in E-Graphs. <https://doi.org/10.48550/arXiv.2405.10188> arXiv:2405.10188 [cs].
- [48] Eric L. Seidel, Niki Vazou, and Ranjit Jhala. 2015. Type Targeted Testing. In *Programming Languages and Systems (Lecture Notes in Computer Science)*, Jan Vitek (Ed.). Springer, Berlin, Heidelberg, 812–836. https://doi.org/10.1007/978-3-662-46669-8_33
- [49] Jessica Shi, Alperen Keles, Harrison Goldstein, Benjamin C. Pierce, and Leonidas Lampropoulos. 2023. Etna: An Evaluation Platform for Property-Based Testing (Experience Report). *Proc. ACM Program. Lang.* 7 (2023). <https://doi.org/10.1145/3607860>
- [50] Dominic Steinhöfel and Andreas Zeller. 2022. Input Invariants. In *Proceedings of the 30th ACM Joint European Software Engineering Conference and Symposium on the Foundations of Software Engineering (ESEC/FSE 2022)*. Association for Computing Machinery, New York, NY, USA, 583–594. <https://doi.org/10.1145/3540250.3549139>
- [51] Rocq Team. 2025. Welcome to a World of Rocq. <https://rocq-prover.org>.
- [52] Kyle Thompson, Nuno Saavedra, Pedro Carrott, Kevin Fisher, Alex Sanchez-Stern, Yuriy Brun, João F. Ferreira, Sorin Lerner, and Emily First. 2025. Rango: Adaptive Retrieval-Augmented Proving for Automated Software Verification. <https://doi.org/10.48550/arXiv.2412.14063> arXiv:2412.14063 [cs].
- [53] Ryan Tjoa, Poorva Garg, Harrison Goldstein, Todd Millstein, Benjamin C. Pierce, and Guy Van Den Broeck. 2025. Tuning Random Generators. <https://rtjoa.com/papers/draft-tuning.pdf>
- [54] John Wrenn, Tim Nelson, and Shriram and Krishnamurthi. 2021. Using Relational Problems to Teach Property-Based Testing. *The art science and engineering of programming* 5, 2 (Jan. 2021). <https://doi.org/10.22152/programming-journal.org/2021/5/9>
- [55] Chunqiu Steven Xia, Matteo Paltenghi, Jia Le Tian, Michael Pradel, and Lingming Zhang. 2024. Fuzz4All: Universal Fuzzing with Large Language Models. In *Proceedings of the IEEE/ACM 46th International Conference on Software Engineering*. 1–13. <https://doi.org/10.1145/3597503.3639121> arXiv:2308.04748 [cs].
- [56] Li-yao Xia. 2021. generic-random. [//hackage.haskell.org/package/generic-random](https://hackage.haskell.org/package/generic-random)
- [57] Huajian Xin, Z. Z. Ren, Junxiao Song, Zhihong Shao, Wanjia Zhao, Haocheng Wang, Bo Liu, Liyue Zhang, Xuan Lu, Qishi Du, Wenjun Gao, Qihao Zhu, Dejian Yang, Zhibin Gou, Z. F. Wu, Fuli Luo, and Chong Ruan. 2024. DeepSeek-Prover-V1.5: Harnessing Proof Assistant Feedback for Reinforcement Learning and Monte-Carlo Tree Search. <https://doi.org/10.48550/arXiv.2408.08152> arXiv:2408.08152 [cs].
- [58] Kaiyu Yang, Aidan M. Swope, Alex Gu, Rahul Chalamala, Peiyang Song, Shixing Yu, Saad Godil, Ryan Prenger, and Anima Anandkumar. 2023. LeanDojo: Theorem Proving with Retrieval-Augmented Language Models. <https://doi.org/10.48550/arXiv.2306.15626> arXiv:2306.15626 [cs].
- [59] Zhixuan Yang and Nicolas Wu. 2022. Fantastic Morphisms and Where to Find Them: A Guide to Recursion Schemes. <https://doi.org/10.48550/arXiv.2202.13633> arXiv:2202.13633 [cs].

A GREATER THAN AND LESS THAN

def greaterThan ($n : \mathbb{N}$) := arbNat $\gg \lambda lo \Rightarrow$ pure ($lo + 1 + n$)

Lemma A.1 (Greater Than Support). $a \in \llbracket \text{greaterThan } lo \rrbracket \iff a > lo$.

$$\frac{}{\Gamma \vdash \text{greaterThan } lo : \text{Gen}_{\mathbb{N}} (\lambda a \Rightarrow a > lo)} \text{S-GREATER THAN}$$

def lessThan ($hi : \mathbb{N}$) := choose 0 ($hi - 1$)

Lemma A.2 (Less Than Support). If $0 < hi$, then $a \in \llbracket \text{lessThan } hi \rrbracket \iff a < hi$.

$$\frac{\Gamma \vdash 0 < hi}{\Gamma \vdash \text{lessThan } hi : \text{Gen}_{\mathbb{N}} (\lambda a \Rightarrow a < hi)} \text{S-LESS THAN}$$

$$\frac{}{\Gamma \vdash \text{assume } 0 < hi \text{ in } \text{lessThan } hi : \text{Gen}_{\mathbb{N}} (\lambda a \Rightarrow a < hi)} \text{S-LESS THAN PARTIAL}$$

B PREDICATE DEFINITIONS

```

1422 def List.fold { $\alpha$   $\beta$ : Type} (f :  $\alpha \rightarrow \beta \rightarrow \beta$ ) (z :  $\beta$ ) (xs : List  $\alpha$ ) :=
1423   List.foldr f z xs
1424
1425 inductive Tree ( $\alpha$  : Type) where
1426   | leaf : Tree  $\alpha$ 
1427   | node : (l : Tree  $\alpha$ )  $\rightarrow$  (x :  $\alpha$ )  $\rightarrow$  (r : Tree  $\alpha$ )  $\rightarrow$  Tree  $\alpha$ 
1428
1429 def Tree.fold
1430   { $\alpha$   $\beta$  : Type}
1431   (f :  $\beta \rightarrow \alpha \rightarrow \beta \rightarrow \beta$ )
1432   (z :  $\beta$ )
1433   (t : Tree  $\alpha$ ) :
1434      $\beta$  :=
1435     match t with
1436     | .leaf => z
1437     | .node l x r => f (Tree.fold f z l) x (Tree.fold f z r)
1438
1439 inductive Label where
1440   | low
1441   | high
1442
1443 inductive Atom where
1444   | atm (n : Nat) (l : Label)
1445
1446 inductive Stack where
1447   | mty
1448   | cons (a : Atom) (s : Stack)
1449   | ret_cons (pc : Atom) (s : Stack)
1450
1451 def Stack.fold
1452   { $\alpha$  : Type}
1453   (z :  $\alpha$ )
1454   (f_c : Atom  $\rightarrow$   $\alpha \rightarrow \alpha$ )
1455   (f_rc : Atom  $\rightarrow$   $\alpha \rightarrow \alpha$ )
1456   (s : Stack) :  $\alpha$  :=
1457   match s with
1458   | .mty => z
1459   | .cons x s' => f_c x (Stack.fold z f_c f_rc s')
1460   | .ret_cons pc s' => f_rc pc (Stack.fold z f_c f_rc s')
1461
1462 inductive Ty : Type where
1463   | unit
1464   | arrow ( $\tau_1$   $\tau_2$  : Ty)
1465   deriving BEq
1466
1467 inductive Term : Type where
1468   | unit
1469   | var (n : Nat)
1470   | abs ( $\tau$  : Ty) (t : Term)
1471   | app (t1 t2 : Term)
1472
1473 def Term.fold
1474   { $\alpha$  : Type}
1475   (z :  $\alpha$ )
1476   (zn : Nat  $\rightarrow$   $\alpha$ )

```

```

1471     (f_abs : Ty ->  $\alpha$  ->  $\alpha$ )
1472     (f_app:  $\alpha$  ->  $\alpha$  ->  $\alpha$ )
1473     (t : Term) :
1474      $\alpha$  :=
1475     match t with
1476     | .unit => z
1477     | .var n => zn n
1478     | .abs  $\tau$  t' => f_abs  $\tau$  (Term.fold z zn f_abs f_app t')
1479     | .app t1 t2 =>
1480       f_app (Term.fold z zn f_abs f_app t1) (Term.fold z zn f_abs f_app t2)
1481 def isAllTwos : List Nat -> Bool
1482 | [] => true
1483 | x :: xs => x = 2 && isAllTwos xs
1484 def isEvenLen : List  $\alpha$  -> Bool
1485 | [] => true
1486 | _ :: xs => !(isEvenLen xs)
1487 def isAllTwosEvenLen (xs : List Nat) : Bool :=
1488   isAllTwos xs && isEvenLen xs
1489 def isIncreasingByOneAux (xs : List Nat) (prev : Nat) : Bool :=
1490   match xs with
1491   | [] => true
1492   | x :: xs' => x == prev + 1 && isIncreasingByOneAux xs' x
1493 def isIncreasingByOne (xs : List Nat) : Bool :=
1494   isIncreasingByOneAux xs 0
1495 def isLengthKAllTwos (k : Nat) (xs : List Nat) : Bool :=
1496   xs.length == k && isAllTwos xs
1497 def isSortedBetween : List Nat -> Nat  $\times$  Nat -> Bool := fun xs (lo, hi) =>
1498   match xs with
1499   | [] => true
1500   | x :: xs' => (lo <= x && x <= hi) && isSortedBetween xs' (x, hi)
1501 def isTrue : List  $\alpha$  -> Bool
1502 | [] => true
1503 | x :: xs => (fun _ => true) x && isTrue xs
1504 def isAllTwosFold (xs : List Nat) : Bool :=
1505   List.fold (fun x b => x == 2 && b) true xs
1506 def isAllTwosEvenLenFold (xs : List Nat) : Bool :=
1507   List.fold (fun x b => x == 2 && b) true xs = true
1508    $\wedge$  List.fold (fun _ b => !b) true xs
1509 def isEvenLenFold (xs : List  $\alpha$ ) : Bool :=
1510   List.fold (fun _ b => !b) true xs
1511 def isIncreasingByOneFold (xs : List Nat) : Bool :=
1512   List.fold (fun x b prev => x == prev + 1 && b x) (fun _ => true) xs 0
1513 def lengthFold (xs : List  $\alpha$ ) : Nat :=
1514   List.fold (fun _ b => b + 1) 0 xs
1515 def isLengthKAllTwosFold (k : Nat) (xs : List Nat) :=
1516   List.fold (fun _ b => b + 1) 0 xs = k
1517    $\wedge$  List.fold (fun x b => x == 2 && b) true xs
1518 def isSortedBetweenFold (lo hi : Nat) (xs : List Nat) : Prop :=
1519

```

```

1520   List.fold (fun x b s => decide (s <= x)
1521     && decide (x <= hi) && b x) (fun _ => true) xs lo
1522 def isTrueFold (xs : List  $\alpha$ ) : Bool :=
1523   List.fold (fun _ b => b) true xs
1524 def isAllTwos : Tree Nat -> Bool
1525   | .leaf => true
1526   | .node l x r => x == 2 && isAllTwos l && isAllTwos r
1527 def isBST : Tree Nat -> (Nat * Nat) -> Bool := fun t <lo, hi> =>
1528   match t with
1529   | .leaf => true
1530   | .node l x r =>
1531     (lo <= x && x <= hi) &&
1532     isBST l <lo, x - 1> &&
1533     isBST r <x + 1, hi>
1534 def isComplete : Tree  $\alpha$  -> Nat -> Bool := fun t n =>
1535   match t with
1536   | .leaf => n == 0
1537   | .node l _ r =>
1538     n > 0 &&
1539     isComplete l (n - 1) &&
1540     isComplete r (n - 1)
1541 def isIncreasingByOneAux (t : Tree Nat) (prev : Nat) : Bool :=
1542   match t with
1543   | .leaf => true
1544   | .node l x r =>
1545     x == prev + 1 &&
1546     isIncreasingByOneAux l x &&
1547     isIncreasingByOneAux r x
1548 def isIncreasingByOne (t : Tree Nat) : Bool :=
1549   isIncreasingByOneAux t 0
1550 def isNonempty : Tree  $\alpha$  -> Bool
1551   | .leaf => false
1552   | .node l _ r => true && isNonempty l && isNonempty r
1553 def isAllTwosFold (t : Tree Nat) : Bool :=
1554   Tree.fold (fun bl x br => x == 2 && bl && br) true t
1555 def isBSTFold (lo hi : Nat) (t : Tree Nat) : Bool :=
1556   Tree.fold
1557     (fun bl x br s =>
1558       match s with
1559       | (sl, sr) =>
1560         (decide (sl <= x) &&
1561           decide (x <= sr)) &&
1562         bl (sl, x - 1) && br (x + 1, sr))
1563     (fun _ => true) t (lo, hi)
1564 def isCompleteFold (n : Nat) (t : Tree Nat) : Bool :=
1565   Tree.fold (fun bl _ br s => decide (s > 0) &&
1566     bl (s - 1) &&
1567     br (s - 1)) (fun s => s == 0) t n
1568

```



```

1569 def isIncreasingByOneFold (t : Tree Nat) : Bool :=
1570   Tree.fold
1571     (fun bl x br prev => x == prev + 1 && bl x && br x)
1572     (fun _ => true) t 0
1573 def isNonemptyFold (t : Tree  $\alpha$ ) : Bool :=
1574   Tree.fold (fun _ _ _ => true) false t
1575 def isBalanced : Tree Nat -> Nat -> Bool := fun t height =>
1576   match t with
1577   | .leaf => height <= 1
1578   | .node l _ r =>
1579     height > 0 &&
1580     isBalanced l (height - 1) &&
1581     isBalanced r (height - 1)
1582 def isAVL (height lo hi : Nat) (t : Tree Nat) : Bool :=
1583   isBST t (lo, hi) && isBalanced t height
1584 def isAVLFold (height lo hi : Nat) (t : Tree Nat) : Bool :=
1585   Tree.fold
1586     (fun bl x br bounds =>
1587       match bounds with
1588       | (sl, sr) => decide (sl <= x) && decide (x <= sr)
1589         && bl (sl, x - 1) && br (x + 1, sr))
1590     (fun _ => true) t (lo, hi) = true
1591   ^
1592   Tree.fold
1593     (fun bl _ br h => decide (h > 0) && bl (h - 1) && br (h - 1))
1594     (fun h => decide (h <= 1)) t height
1595 def isGoodNat (n : Nat) : Bool :=
1596   n == 0 || n == 1
1597 def isGoodAtom : Atom -> Bool
1598 | .atm n _ => isGoodNat n
1599 def isGoodStackFold (s : Stack) (n : Nat) : Bool :=
1600   Stack.fold (fun s => s == 0)
1601     (fun x acc s => isGoodAtom x && acc (s - 1))
1602     (fun pc acc s => isGoodAtom pc && acc (s - 1)) s n
1603 def isGoodStack (s : Stack) (n : Nat) : Bool :=
1604   match s with
1605   | .mty => n == 0
1606   | .cons x s' => (n > 0 && isGoodAtom x) && isGoodStack s' (n - 1)
1607   | .ret_cons pc s' => (n > 0 && isGoodAtom pc) && isGoodStack s' (n - 1)
1608 def getType (t : Term) ( $\Gamma$  : List Ty) : Option Ty :=
1609   match t with
1610   | .unit => pure .unit
1611   | .var n =>  $\Gamma$ [n]?
1612   | .abs  $\tau$  t => do
1613     let  $\tau'$  ← getType t ( $\tau :: \Gamma$ )
1614     pure (.arrow  $\tau$   $\tau'$ )
1615   | .app t1 t2 => do
1616     let  $\tau_1$  ← getType t1  $\Gamma$ 

```

```

1618     let  $\tau_2 \leftarrow \text{getType } t_2 \ \Gamma$ 
1619     match  $\tau_1$  with
1620     | .arrow  $\tau_{\text{arg}} \ \tau_{\text{res}} \Rightarrow \text{do}$ 
1621         guard ( $\tau_{\text{arg}} == \tau_2$ )
1622         pure  $\tau_{\text{res}}$ 
1623     | .unit  $\Rightarrow$  failure
1624 def isWellTyped ( $\Gamma : \text{List Ty}$ ) ( $t : \text{Term}$ ) : Prop :=
1625      $\exists (\tau : \text{Ty}), \text{getType } t \ \Gamma = \tau$ 
1626 def isWellScoped : Term -> Nat -> Bool := fun t varCap =>
1627     match t with
1628     | .unit  $\Rightarrow$  true
1629     | .var n  $\Rightarrow$  n < varCap
1630     | .abs _ t  $\Rightarrow$  isWellScoped t (varCap + 1)
1631     | .app t1 t2  $\Rightarrow$  isWellScoped t1 varCap && isWellScoped t2 varCap
1632 def getTypeFold : Term -> List Ty -> Option Ty :=
1633     Term.fold
1634         (fun _  $\Rightarrow$  pure .unit)
1635         (fun n  $\Gamma'$   $\Rightarrow$   $\Gamma'[n]?$ )
1636         (fun  $\tau_1$  b  $\Gamma'$   $\Rightarrow$  do
1637             let  $\tau_2 \leftarrow b \ (\tau_1 :: \Gamma')$ 
1638             pure (.arrow  $\tau_1 \ \tau_2$ ))
1639         (fun b1 b2  $\Gamma'$   $\Rightarrow$  do
1640             let  $\tau_1 \leftarrow b1 \ \Gamma'$ 
1641             let  $\tau_2 \leftarrow b2 \ \Gamma'$ 
1642             match  $\tau_1$  with
1643             | .arrow  $\tau_{\text{arg}} \ \tau_{\text{res}} \Rightarrow \text{do}$ 
1644                 guard ( $\tau_{\text{arg}} == \tau_2$ )
1645                 pure  $\tau_{\text{res}}$ 
1646             | Ty.unit  $\Rightarrow$  failure)
1647 def isWellTypedFold ( $\Gamma : \text{List Ty}$ ) ( $t : \text{Term}$ ) : Prop :=
1648      $\exists \tau, \text{getTypeFold } t \ \Gamma = \text{some } \tau$ 
1649 def isWellScopedFold (varCap : Nat) ( $t : \text{Term}$ ) : Bool :=
1650     Term.fold
1651         (fun _  $\Rightarrow$  true)
1652         (fun n s  $\Rightarrow$  s < n)
1653         (fun _ b s  $\Rightarrow$  b (s + 1))
1654         (fun b1 b2 s  $\Rightarrow$  b1 s && b2 s)
1655     t
1656     varCap
1657
1658
1659
1660
1661
1662
1663
1664
1665
1666

```

C EXTENDED TABLE OF BENCHMARKS

Predicate	Type	Time (s)
$v = 2$	Nat	0.04 (0.01)
$2 = v$	Nat	0.04 (0.00)
$v = 2 \vee v = 5$	Nat	0.08 (0.00)
$v = 2 \vee v = 5 \wedge \text{True}$	Nat	0.08 (0.00)
$\exists a, a = 3 \wedge v = a + 1$	Nat	0.04 (0.00)
$5 \leq v \wedge v \leq 10$	Nat	0.08 (0.00)
$v > 5$	Nat	0.07 (0.00)
$v = 0 \vee lo \leq v \wedge v \leq hi$	Nat	0.14 (0.00)
$\text{isAllTwos } v = \text{true}$	List Nat	0.84 (0.01)
$\text{isAllTwosEvenLen } v = \text{true}$	List Nat	2.76 (0.02)
$\text{isEvenLen } v = \text{true}$	List Nat	2.22 (0.02)
$\text{isIncreasingByOne } v = \text{true}$	List Nat	1.44 (0.01)
$\text{List.length } v = k$	List Nat	1.89 (0.01)
$\text{isLengthKAllTwos } k \ v = \text{true}$	List Nat	2.37 (0.01)
$\text{isSortedBetween } v \ (lo, hi) = \text{true}$	List Nat	1.72 (0.02)
$\text{isTrue } v = \text{true}$	List Nat	2.21 (0.01)
$\text{isAllTwosFold } v = \text{true}$	List Nat	0.34 (0.01)
$\text{isAllTwosEvenLenFold } v = \text{true}$	List Nat	2.73 (0.02)
$\text{isEvenLenFold } v = \text{true}$	List Nat	2.19 (0.01)
$\text{isIncreasingByOneFold } v = \text{true}$	List Nat	1.12 (0.01)
$\text{lengthFold } v = k$	List Nat	1.87 (0.01)
$\text{isLengthKAllTwosFold } k \ v = (\text{true} = \text{true})$	List Nat	2.35 (0.08)
$\text{isSortedBetweenFold } lo \ hi \ v = (\text{true} = \text{true})$	List Nat	1.34 (0.01)
$\text{isTrueFold } v = \text{true}$	List Nat	2.18 (0.01)
$\text{isAllTwos } v = \text{true}$	Tree Nat	0.61 (0.01)
$\text{isBST } v \ (lo, hi) = \text{true}$	Tree Nat	1.86 (0.01)
$\text{isComplete } v \ n = \text{true}$	Tree Nat	2.35 (0.02)
$\text{isIncreasingByOne } v = \text{true}$	Tree Nat	1.57 (0.02)
$\text{isNonempty } v = \text{true}$	Tree Nat	1.66 (0.04)
$\text{isAllTwosFold } v = \text{true}$	Tree Nat	0.60 (0.01)
$\text{isBSTFold } lo \ hi \ v = \text{true}$	Tree Nat	1.89 (0.02)
$\text{isCompleteFold } v \ n = \text{true}$	Tree Nat	2.37 (0.02)
$\text{isIncreasingByOneFold } v = \text{true}$	Tree Nat	1.55 (0.02)
$\text{isNonemptyFold } v = \text{true}$	Tree Nat	1.44 (0.02)
$\text{isGoodStack } v \ n = \text{true}$	Stack	5.39 (0.05)
$\text{isGoodStackFold } v \ n = \text{true}$	Stack	7.70 (0.11)
$\text{isWellScoped } v \ o = \text{true}$	Term	2.85 (0.03)
$\text{isWellTyped } \Gamma \ v$	Term	3.66 (0.03)
$\text{isWellTypedFold } \Gamma \ v$	Term	3.71 (0.03)
$\text{isWellScopedFold } v \ o = \text{true}$	Term	2.84 (0.03)
$lo \leq v \wedge v \leq hi$	Nat	0.07 (0.00)
$\text{isAVL height } lo \ hi \ v = \text{true}$	Tree Nat	34.44 (0.22)
$\text{isAVLFold height } lo \ hi \ v = \text{true}$	Tree Nat	34.37 (0.23)

D MANUALLY WRITTEN STLC GENERATOR

```

1716 def genWellTyped ( $\Gamma$  : List Ty) : Gen Term := by
1717   let  $\tau$  <- arbTy
1718   Term.unfold
1719     (fun ( $\tau$ ,  $\Gamma$ ) => do
1720       pick
1721         (caseTy  $\tau$ 
1722           (fun () =>
1723             --  $\tau$  = .unit
1724             pure TermF.unitStep)
1725           (fun  $\tau_1$   $\tau_2$  () =>
1726             --  $\tau$  = .arrow  $\tau_1$   $\tau_2$ 
1727             pure (TermF.absStep  $\tau_1$  ( $\tau_2$ ,  $\tau_1$  ::  $\Gamma$ ))))
1728   (if ( $\Gamma$ .indexesOf  $\tau$ ).length > 0 then
1729     pick
1730       (do
1731         let  $n$  <- elements ( $\Gamma$ .indexesOf .unit) (...)
1732         pure (TermF.varStep  $n$ )
1733       (do
1734         let  $\tau'$  <- arbTy
1735         pure (TermF.appStep (.arrow  $\tau'$   $\tau$ ,  $\Gamma$ ) ( $\tau'$ ,  $\Gamma$ )))
1736   else do
1737     let  $\tau'$  <- arbTy
1738     pure (TermF.appStep (.arrow  $\tau'$   $\tau$ ,  $\Gamma$ ) ( $\tau'$ ,  $\Gamma$ )))
1739   ( $\tau$ ,  $\Gamma$ )

```

E STLC GENERATORS FROM ETNA

```

1765
1766 genTyp :: Gen Ty
1767 genTyp = sized go
1768   where
1769     go 0 = return TBool
1770     go n =
1771       oneof
1772         [ go 0,
1773           TFun <$> go (div n 2) <*> go (div n 2)
1774         ]
1775
1776 genExpr :: Ctx -> Typ -> Gen Expr
1777 genExpr ctx t = sized $ \n -> go n ctx t
1778   where
1779     go 0 ctx t = oneof $ genOne ctx t : genVar ctx t
1780     go n ctx t =
1781       oneof
1782         ( [genOne ctx t]
1783           ++ [genAbs ctx t1 t2 | TFun t1 t2 <- [t]]
1784           ++ [genApp ctx t]
1785           ++ genVar ctx t
1786         )
1787   where
1788     genAbs ctx t1 t2 = Abs t1 <$> go (n - 1) (t1 : ctx) t2
1789
1790     genApp ctx t = do
1791       t' <- genTyp
1792       e1 <- go (div n 2) ctx (TFun t' t)
1793       e2 <- go (div n 2) ctx t'
1794       return (App e1 e2)
1795
1796 genOne :: Ctx -> Typ -> Gen Expr
1797 genOne _ TBool = Bool <$> elements [True, False]
1798 genOne ctx (TFun t1 t2) = Abs t1 <$> genOne (t1 : ctx) t2
1799
1800 genVar :: Ctx -> Typ -> [Gen Expr]
1801 genVar ctx t = [Var <$> elements vars | not (null vars)]
1802   where
1803     vars = filter (\i -> ctx !! i == t) [0 .. (length ctx - 1)]
1804
1805
1806
1807
1808
1809
1810
1811
1812
1813

```

```

1814 Fixpoint genVar' (ctx: Ctx) (t: Typ) (p: nat) (r: list nat) : list nat :=
1815   match ctx with
1816   | nil => r
1817   | t'::ctx' =>
1818     if t = t' then genVar' ctx' t (p + 1) (p :: r)
1819     else genVar' ctx' t (p + 1) r
1820   end.
1821
1822 Fixpoint genZero env tau : G (option Expr) :=
1823   match tau with
1824   | TBool =>
1825     bindGen arbitrary
1826       (fun b : bool => returnGen (Some (Bool b)))
1827   | TFun T1 T2 =>
1828     bindOpt
1829       (genZero (T1 :: env) T2)
1830       (fun e : Expr => returnGen (Some (Abs T1 e)))
1831   end.
1832
1833
1834 Fixpoint genExpr env tau (sz: nat) : G (option Expr) :=
1835   match sz with
1836   | 0 =>
1837     backtrack
1838       [(1, oneOf_ (ret None)
1839         (map (fun x => returnGen (Some (Var x))) (genVar' env tau 0 [])))
1840        ;(1, genZero env tau)]
1841   | S sz' =>
1842     backtrack
1843       [(1, oneOf_ (ret None)
1844         (map (fun x => returnGen (Some (Var x))) (genVar' env tau 0 [])))
1845        ;
1846        (1, bindGen arbitrary (fun T1 : Typ =>
1847          bindOpt (genExpr env (TFun T1 tau) sz') (fun e1 : Expr =>
1848            bindOpt
1849              (genExpr env T1 sz')
1850              (fun e2 : Expr => returnGen (Some (App e1 e2))))))
1851        ;
1852        (1, match tau with
1853          | TBool =>
1854            bindGen arbitrary (fun b : bool => returnGen (Some (Bool b)))
1855          | TFun T1 T2 =>
1856            bindOpt
1857              (genExpr (T1 :: env) T2 sz')
1858              (fun e : Expr =>
1859                returnGen (Some (Abs T1 e)))
1860          end)]
1861   end.
1862

```


F MORE SIDE-BY-SIDE GENERATOR COMPARISONS

F.1 One Or In Range

```
def genOneOrInRange (lo hi : Nat) : Gen Nat :=
  if h : decide (lo <= hi) = true then
    pick (pure 0) (choose lo hi (s_between_partial._proof_1 h))
  else
    pure 0
```

```
/-
Differences:
- Simplify proof for choose.
-/
```

```
def genOneOrInRange_manual (lo hi : Nat) : Gen Nat :=
  if h : lo <= hi then
    pick (pure 0) (choose lo hi (by omega))
  else
    pure 0
```

F.2 Complete Tree

```
def genCompleteTree (n : Nat) : Gen (Tree Nat) :=
  Tree.unfold
    (fun x =>
      if x.snd = 0 then pure TreeF.leaf
      else do
        let a <- arbNat
        pure (TreeF.node ((), x.2 - 1) a ((), x.2 - 1)))
    ((), n)
```

```
/-
Differences:
- Remove extra unit in collector.
-/
```

```
def genComplete_manual (n : Nat) : Gen (Tree Nat) :=
  Tree.unfold
    (fun height =>
      if height = 0 then
        pure TreeF.leaf
      else do
        let a <- arbNat
        pure (TreeF.node (height - 1) a (height - 1)))
    n
```

F.3 Sorted Between

```

1912 def genSortedBetween (lo hi : Nat) : Gen (List Nat) :=
1913   List.unfold
1914     (fun x =>
1915       if h : decide (x.snd.fst <= x.snd.snd) = true then
1916         pick (pure ListF.nil) do
1917           let a <- choose x.2.1 x.2.2 (s_between_partial._proof_1 h)
1918           pure (ListF.cons a (PUnit.unit, a, x.2.2))
1919       else
1920         pure ListF.nil)
1921     (PUnit.unit, lo, hi)
1922
1923
1924 /-
1925 Differences:
1926 - Simplify proof for choose.
1927 - Remove extra unit in collector.
1928 -/
1929 def genSortedBetween_manual (lo hi : Nat) : Gen (List Nat) :=
1930   List.unfold
1931     (fun (lo, hi) =>
1932       if h : lo <= hi then
1933         pick
1934           (pure ListF.nil)
1935           (do
1936             let a <- choose lo hi (by omega)
1937             pure (ListF.cons a (a, hi)))
1938       else
1939         pure ListF.nil)
1940     (lo, hi)
1941
1942
1943
1944
1945
1946
1947
1948
1949
1950
1951
1952
1953
1954
1955
1956
1957
1958
1959
1960

```

F.4 Length K, All Twos

```

1961 def genLengthKAllTwos (k : Nat): Gen (List Nat) :=
1962   List.unfold
1963     (fun x =>
1964       if x.fst.fst = 0 then pure ListF.nil
1965       else pure (ListF.cons 2 ((Nat.pred x.1.1, PUnit.unit), PUnit.unit, PUnit.unit)))
1966     ((k, PUnit.unit), PUnit.unit, PUnit.unit)
1967
1968
1969 /-
1970 Differences:
1971 - Remove two extra units in collector.
1972 -/
1973 def genLengthKAllTwos_manual (k : Nat): Gen (List Nat) :=
1974   List.unfold
1975     (fun len =>
1976       if len = 0 then
1977         pure ListF.nil
1978       else
1979         pure (ListF.cons 2 (len - 1)))
1980     k
1981
1982
1983
1984
1985
1986
1987
1988
1989
1990
1991
1992
1993
1994
1995
1996
1997
1998
1999
2000
2001
2002
2003
2004
2005
2006
2007
2008
2009

```

F.5 AVL

```

2010 def genAVL (height lo hi : Nat) : Gen (Tree Nat) :=
2011   Tree.unfold
2012     (fun x => do
2013       let __do_lift <-
2014         if x.snd.snd = 0 then pure TreeF.leaf
2015         else
2016           if Nat.pred x.snd.snd = 0 then
2017             if h : decide (x.snd.fst.fst <= x.snd.fst.snd) = true then
2018               pick (pure TreeF.leaf) do
2019                 let a <- choose x.2.1.1 x.2.1.2 (s_between_partial._proof_1 h)
2020                 pure (TreeF.node (PUnit.unit, PUnit.unit) a (PUnit.unit, PUnit.unit))
2021             else pure TreeF.leaf
2022           assume (decide (x.snd.fst.fst <= x.snd.fst.snd)) fun h => do
2023             let a <- choose x.2.1.1 x.2.1.2 (s_between_partial._proof_1 h)
2024             pure (TreeF.node (PUnit.unit, PUnit.unit) a (PUnit.unit, PUnit.unit))
2025       match __do_lift with
2026       | TreeF.leaf => pure TreeF.leaf
2027       | TreeF.node bl x_1 br =>
2028         pure
2029           (TreeF.node (bl, (x.2.1.1, x_1 - 1), x.2.2 - 1) x_1
2030             (br, (x_1 + 1, x.2.1.2), x.2.2 - 1)))
2031       ((PUnit.unit, PUnit.unit), (lo, hi), height)
2032
2033
2034
2035
2036
2037
2038
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058

```

```

2059 /-
2060 Differences:
2061 - Remove two extra units in collector.
2062 - Nicer match on height to reduce some duplication.
2063 - Generator is technically total now; this requires insight about the total
2064   number of values that can appear in a tree of height k.
2065 -/
2066 def genAVL_manual' (height lo hi : Nat) : Gen (Tree Nat) :=
2067   -- Guarantee that there are enough values in the range, given the height.
2068   assume (hi - lo > 2 ^ height) fun _ =>
2069     Tree.unfold
2070       (fun (lo, hi, height) => do
2071         match height with
2072         | 0 => pure TreeF.leaf
2073         | 1 =>
2074           pick (pure TreeF.leaf)
2075             (assume (lo <= hi) fun h => do -- Will always succeed.
2076               -- Choose values so we never truncate the range to be too small.
2077               let a <-
2078                 choose
2079                   (lo + 2 ^ (height - 1))
2080                   (hi - 2 ^ (height - 1)) (by ...))
2081               pure (TreeF.node (lo, a - 1, height - 1) a (a + 1, hi, height - 1)))
2082       | height' + 1 => do
2083         assume (lo <= hi) fun h => do -- Will always succeed.
2084           -- Choose values so we never truncate the range to be too small.
2085           let a <-
2086             choose
2087               (lo + 2 ^ (height - 1))
2088               (hi - 2 ^ (height - 1)) (by ...))
2089           pure (TreeF.node (lo, a - 1, height - 1) a (a + 1, hi, height - 1)))
2090       (lo, hi, height)
2091
2092 /-
2093 Differences:
2094 - Entirely different approach.
2095 - Relies on AVL.insert being correct.
2096 -/
2097 def genAVL_manual'' : Gen (Tree Nat) :=
2098   -- Generate list of arbitrary Nats
2099   let values <-
2100     List.unfold (fun () =>
2101       pick
2102         (pure (ListF.nil))
2103         (do let a <- arbNat; pure (ListF.cons a ())))
2104     ()
2105   -- Insert all values into an empty AVL tree.
2106   pure (List.fold (fun x t => AVL.insert t x) AVL.empty)
2107

```