

Parsing Randomness

Unifying and Differentiating Parsers and Random Generators

Anonymous Author(s)

Abstract

“A generator is a parser of randomness.” This perspective on random generation is well established as folklore, but to our knowledge has not been formalized, nor have its consequences been deeply explored.

We present *free generators*, which unify parsing and generation in a common structure that makes the relationship between the two concepts concrete. Free generators naturally lead us to a proof that many generators can be factored into a parser and a distribution over choice sequences.

We also put this new abstraction to work by observing that free generators have a notion of *derivative* that “previews” the effect of a particular generator choice. This derivative is compatible with related ideas in formal languages, and we use it in a novel algorithm for generating data satisfying preconditions.

Keywords: Random generation, Parsing, Property-based testing, Formal languages

1 Introduction

“A generator is a parser of randomness.” It’s one of those observations that’s confusing until it’s obvious.

A random generator processes a series of random choices into a data structure, just as a parser processes a series of characters into a data structure. But few tend to think of generators this way. Indeed, to our knowledge the framing of generators as parsers is folklore that has not been explored formally.¹ What a shame! The relationship between these core computer science concepts deserves more than a passing mention.

How, exactly, is a random generator like a parser? A generator is a program that builds a data structure by making a sequence of random choices—those choices are the key. A traditional generator makes decisions using a stored source of randomness (e.g., a seed) that updates throughout the computation. But there is nothing stopping us from pre-computing choices for the generator to use; we could provide a list of choices ahead of time and tell the generator to make those choices one after the other. This mode of operation is exactly parsing!

¹The HYPOTHESIS [13] library does use a parser internally for generation, but further implications have not been explored.

This intuition needs some fleshing out. For example, is there a representation of generators that makes this connection concrete? Where does the probability distribution “go” when we view generators as parsers? To answer these questions, we present a data structure called a *free generator* that can be interpreted as either a generator or as a parser. We establish a rich theory of free generators and use them to prove that a subset of generator programs can be factored into a parser and a distribution over sequences of choices.

Besides clarifying folklore, free generators admit transformations that make them useful as a generation representation. The most exciting transformation is a notion of *derivative* which modifies a generator by asking the question: “what would this generator look like after it makes choice *c*?” The derivative gives a way of previewing a particular choice to determine how good or bad that choice is.

We use free generator derivatives to implement a novel algorithm for tackling a well-documented problem, which we refer to as the *valid generation problem*. The valid generation problem requires generating a large number of the random values that satisfy some specified validity condition. This comes up often in property-based testing, where the validity condition is the precondition of some functional specification. Since generator derivatives give a way of previewing the effects of a particular choice, we can use *gradients* (derivatives of many choices at once) to preview all possible choices and pick the most promising one.

In the next section, we describe the high-level motivation behind free generators and the operations we define on them (§2). After that, we present our main contributions:

- We formalize a folklore correspondence between parsers and generators using *free generators* and prove that every “*applicative*” generator can be factored into a parser and a probability distribution (§3).
- We exploit our new formalism to transport an idea from formal languages—the *Brzowski derivative*—to the context of generators (§4).
- We show that all of this theory is worth the trouble! We present an algorithm that uses derivatives to turn a naïve generator into one that produces only values satisfying a Boolean precondition (§5). Our algorithm performs quite well on simple benchmarks: in most cases, it produces more than twice as many valid values as the naïve generator in the same amount of time (§6).

We conclude with related work (§7) and ideas for future research (§8).

2 The High-Level Story

Moving straight into the technical details would risk missing the forest for the trees. Accordingly, this section gives a bird's-eye view of our motivation, theory, and exploratory evaluation.

2.1 Generators and Parsers

Before trying to formalize the connection between generation and parsing, we should see some examples. Let us start by looking at a specific generator and understanding what it has in common with a specific parser.

Consider `genTree`, defined in Figure 1. The program `genTree` is a generator that produces random binary trees of Booleans like

```
Node True Leaf Leaf   and
Node True Leaf (Node False Leaf Leaf)
```

guided by a series of random coin flips. The choices lead the program to generate trees up to a given height via a simple recursive procedure. Note that the generator does not always make the same number of choices: sometimes it chooses to return a `Leaf` and terminate, and other times it chooses to build a `Node`, which requires more choices.²

Now, consider `parseTree` (also in Figure 1), which parses a tree from a string containing the characters `n`, `l`, `t`, and `f`. The parser turns

```
ntll into Node True Leaf Leaf   and
ntlnfl into Node True Leaf (Node False Leaf Leaf).
```

The program operates character by character, consuming characters of the input string with `consume` and sometimes changing the way it handles the next character based on the previous character.

At this point the superficial similarities between `genTree` and `parseTree` should be clear, but what is really going on? It all comes down to *choices*. In `genTree`, choices are made randomly during the execution of the program, while in `parseTree` the choices are made ahead of time and manifest as the characters in the input string. The programs have the same structure, and only differ in their expectation of when choices should be made and the “labels” for those choices.

2.2 Free Generators

We can unify random generation with parsing by abstracting both ideas into a single data structure. For this, we introduce **free generators**.³ Free generators are not traditional programs like `genTree` or `parseTree`, they are more like abstract

²Program synthesis experts might wonder why we use programs like this, rather than PCFGs or other related representations. Our work may very well translate to those domains, but we chose to target “applicative” generators because they are slightly more expressive.

³This document uses the knowledge package in \LaTeX to make definitions interactive. Readers viewing the PDF electronically can click on technical terms and symbols to see where they are defined in the document.

```
genTree h =
  if h == 0 then
    return Leaf
  else
    c ← flip()
    if c == Head then return Leaf
    if c == Tail then
      c ← flip()
      if c == Head then x ← True
      if c == Tail then x ← False
      l ← genTree (h - 1)
      r ← genTree (h - 1)
      return Node x l r

parseTree h =
  if h == 0 then
    return Leaf
  else
    c ← consume()
    if c == l then return Leaf
    if c == n then
      c ← consume()
      if c == t then x ← True
      if c == f then x ← False
      else fail
      l ← parseTree (h - 1)
      r ← parseTree (h - 1)
      return Node x l r
    else fail
```

Figure 1. A generator and a parser for Boolean binary trees.

syntax trees. Free generators can be *interpreted* as programs that either generate or parse.

In §3 we give a domain-specific language for constructing free generators that is almost identical to the normal way users would construct generators with `QUICKCHECK`. But we give an example here that is explicitly built from data constructors order to make it clear that these are data structures, not normal programs. When reading Figure 2, focus on the structural similarities between `fgenTree` and the examples in Figure 1.

The free generator `fgenTree h` produces binary trees of Booleans. When reading the data structure, think of `Pure` in almost the same way as `return` from the previous examples, returning a pure value without making any choices or parsing any characters. `MapR` takes two arguments, a free generator and a function that can be applied to the result that is generated/parsed. The `Pair` constructor does a sort of sequencing: it generates/parses using its first argument, then it does the same with its second argument, and finally it pairs the results together. Finally, the real magic is in the way we interpret the `Select` structure. When we want a generator,

```

fgenTree h =
  if h == 0 then
    Pure Leaf
  else
    Select
      [ (1, Pure Leaf),
        (n, MapR
          (Pair (Select
            [ (t, Pure True),
              (f, Pure False) ]))
          (Pair (fgenTree (h - 1))
                (fgenTree (h - 1))))
        (λ (x, (l, r)) → Node x l r) ]

```

Figure 2. A free generator for binary trees of Booleans.

we treat it as making a uniform random choice, and when we want a parser we treat it as consuming a character, c , and checking it against the first elements of the pairs.

The almost line-to-line correspondence between `fgenTree` and `genTree` (or `parseTree`) is no accident! In §3 we give formal definitions of free generators, along with a number of interpretation functions. We use $\mathcal{G}[\![\cdot]\!]$ to mean the **generator interpretation** of a free generator and $\mathcal{P}[\![\cdot]\!]$ to mean the **parser interpretation** of a free generator. In other words,

$$\mathcal{G}[\![\text{fgenTree } 5]\!] \approx \text{genTree } 5 \quad \text{and} \quad \mathcal{P}[\![\text{fgenTree } 5]\!] \approx \text{parseTree } 5.$$

These functions allow us to write one free generator that can be used in different ways.

With the generator and parser interpretations in hand, we can ask about the formal relationship between them. The key is one final interpretation, $\mathcal{C}[\![\cdot]\!]$, the **choice distribution**. Intuitively, the choice distribution produces the sequences of choices that the generator interpretation can make, or equivalently the sequences that the parser interpretation can parse.

The choice distribution is used in Theorem 3.4 to connect the dots between parsing and generation. The theorem says that for any free generator g ,

$$\mathcal{P}[\![g]\!] \langle \$ \rangle \mathcal{C}[\![g]\!] \approx \mathcal{G}[\![g]\!]$$

(where $\langle \$ \rangle$ is defined as a kind of “mapping” operation that applies a function to samples from a distribution). Since many normal QUICKCHECK generators can be written as free generators, another way to read this theorem is that generators can often be factored into two pieces: a distribution over choice sequences (given by $\mathcal{C}[\![\cdot]\!]$), and a parser of those sequences (given by $\mathcal{P}[\![\cdot]\!]$). This formalizes the intuition from §1!

2.3 Derivatives of Free Generators

Free generators do more than formalize the folklore about generators as parsers. Since a free generator defines a parser, it defines a formal language—we write the **language interpretation** of a free generator as $\mathcal{L}[\![\cdot]\!]$. The language of a free generator is the set of choice sequences that it can parse (or make). Viewing free generators this way suggests some interesting ways that free generators might be manipulated.

Formal languages have a notion of *derivative* due to Brzozowski [1]. For a formal language L , the Brzozowski derivative is defined as:

$$\delta_c L = \{s \mid c \cdot s \in L\}$$

In other words, the derivative of L with respect to c is all strings in L that start with c , with the first c removed.

Some models of languages, including regular expressions and context-free grammars, have syntactic transformations that correspond to derivatives. While parser programs like `parseTree` cannot be automatically modified in this way, we can write out what such a transformation would look like by hand. Conceptually, the derivative of a parser with respect to a character c is whatever parser remains assuming c has just been parsed.

For example, we can think of the derivative of `parseTree 5` with respect to `n` as:

```

δn(parseTree 5) ≈
  c ← consume()
  if c == t then x ← True
  if c == f then x ← False
  else fail
  l ← parseTree 4
  r ← parseTree 4
  return Node x l r

```

This parser is “one step” simpler than `parseTree 5`. After parsing the character `n`, the next step is to parse either `t` or `f` and then construct a `Node`, so the derivative does just that.

We can look at another derivative of $\delta_n(\text{parseTree } 5)$, this time with respect to `t`:

```

δtδn(parseTree 5) ≈
  l ← parseTree 4
  r ← parseTree 4
  return Node True l r

```

Now we have fixed the value `True` for x , and we can continue by making the recursive calls and constructing the final tree.

Our free generators also have a closely related notion of **derivative**, that can be computed easily! We can take derivatives of the free generator produced by `fgenTree` that look almost identical to the ones that we saw for `parseTree`:

```

331  $\delta_n(\text{fgenTree } 5) \approx$ 
332
333   MapR
334     (Pair (Select
335           [ (t, Pure True),
336             (f, Pure False) ])
337           (Pair (fgenTree 4)
338                 (fgenTree 4)))
339     ( $\lambda (x, (l, r)) \rightarrow \text{Node } x \mid r$ )
340

```

```

341
342  $\delta_t \delta_n(\text{fgenTree } 5) \approx$ 
343
344   MapR
345     (Pair (fgenTree 4)
346           (fgenTree 4))
347     ( $\lambda (l, r) \rightarrow \text{Node True } l \mid r$ )
348

```

In §4 we define a simple procedure for computing the derivative of a free generator and show that the definition behaves the way we want. Formally, we prove Theorem 4.2 which says that for all free generators g ,

$$\delta_c \mathcal{L}[\![g]\!] = \mathcal{L}[\![\delta_c g]\!].$$

In other words, the derivative of the language of g is equal to the language of the derivative of g .

We can think of the derivative of a free generator as the generator that remains after a particular choice. This means that we can preview the result of making a choice without actually interpreting the free generator as a generator or parser.

2.4 Putting Free Generators to Use

In §5 and §6 we present and evaluate an algorithm called **CHOICE GRADIENT SAMPLING** that uses free generators to tackle the *valid generation problem*. Given a validity predicate on a data structure, the goal is to generate as many unique, valid structures as possible in a given amount of time.

Given a simple free generator, our algorithm uses derivatives to evaluate choices and search for valid values. We evaluate our algorithm on four small benchmarks, all of which are common in the property-based testing literature. We compare our algorithm to rejection sampling—sampling from a naïve generator repeatedly and discarding invalid results—which is a simple but useful baseline for understanding how well our algorithm performs. Our algorithm does remarkably well on all but one benchmark, generating more than twice as many valid values as rejection sampling in the same period of time.

3 Free Generators

In this section, we develop the theory of *free generators*. We start with some background information on applicative abstractions for parsing and random generation and then

present a new framing of generators that makes choices explicit and enables syntactic manipulations.

3.1 Background: Applicative Parsers and Generators

In §2 we represented generators and parsers with pseudo-code. Here we bridge the gap between the code in that section and the code that users of our theory would actually work with.

We represent both generators and parsers using *applicative functors* [14]⁴—fair warning, things are about to get a bit HASKELL-ey. At a high level, an applicative functor is a type constructor f with operations:

```

349 ( $\langle \$ \rangle$ ) :: (a → b) → f a → f b
350 pure  :: a → f a
351 ( $\langle * \rangle$ ) :: f (a → b) → f a → f b

```

These operations are mainly useful as a way to apply functions to values inside the type constructor f 's structure. For example, the idiom “ $g \langle \$ \rangle x \langle * \rangle y \langle * \rangle z$ ” applies a pure function g to three structures x , y , and z .

We can use these operations to define `genTree` like we would in QUICKCHECK [4], since the type constructor `Gen` representing generators is an applicative functor:

```

352 genTree :: Int → Gen Tree
353 genTree 0 = pure Leaf
354 genTree h =
355   oneof [ pure Leaf,
356           Node ($) genInt
357           (*) genTree (h - 1)
358           (*) genTree (h - 1) ]

```

Here, `pure` is the trivial generator that always generates the same value, and `Node ($) g1 (*) g2 (*) g3` means apply the constructor `Node` to three sub-generators to produce a new generator. (Operationally, this means sampling x_1 from g_1 , x_2 from g_2 , and x_3 from g_3 , and then constructing `Node x_1 x_2 x_3` .) Notice that we need one extra function beyond the applicative interface: `oneof` makes a uniform choice between generators, just as we saw in the pseudo-code.

We can do the same thing for `parseTree`, using combinators inspired by libraries like PARSEC [11]:

```

359 parseTree :: Int → Parser Tree
360 parseTree 0 = pure Leaf
361 parseTree h =
362   choice [ (1, pure Leaf),
363            (n, Node ($) parseInt
364              (*) parseTree (h - 1)
365              (*) parseTree (h - 1)) ]

```

⁴For Haskell experts: we choose to focus on applicatives, not monads, to clarify our development and avoid some efficiency issues in §4 and §5. We suspect that much of our approach would work for monadic generators as well.

In this context, pure is a parser that consumes no characters and never fails. It just produces the value passed to it. We can interpret Node $\langle \$ \rangle p_1 \langle * \rangle p_2 \langle * \rangle p_3$ as running each sub-parser in sequence (failing if any of them fail) and then wrapping the results in the Node constructor. Finally, we have replaced oneof with choice, but the idea is the same: choose between sub-parsers.

Parsers of this form have type $\text{String} \rightarrow \text{Maybe}(\text{a}, \text{String})$. They can be applied to a string to obtain either Nothing or Just (a, s), where a is the parse result and s contains any extra characters.

3.2 Representing Generators

With the applicative interface in mind, we can now give the formal definition of a *free generator*.⁵

3.2.1 Type Definition. We represent free generators as an inductive data type, FGen, defined as:

data FGen a where

Void :: FGen a

Pure :: a → FGen a

Pair :: FGen a → FGen b → FGen (a, b)

Map :: (a → b) → FGen a → FGen b

Select :: [(Char, FGen a)] → FGen a

These constructors form an abstract syntax tree with nodes that roughly correspond to the functions in the applicative interface. Clearly Pure represents pure. Pair is a slightly different form of $\langle * \rangle$ —one is definable from the other, but this version makes more sense as a data constructor. Map corresponds to $\langle \$ \rangle$.⁶ Finally, Select subsumes both oneof and choice—it might mean either, depending on the interpretation. We need Void for technical reasons; it represents failure.

Free generators draw inspiration from *free applicative functors* [2]. Accordingly, we can write transformations from $\text{FGen a} \rightarrow \text{f a}$ for any f with similar structure. This fact motivates the rest of this section.

3.2.2 Language of a Free Generator. We say that the *language of a free generator* is the set of choice sequences that it might make or parse. We define a generator’s language recursively, by cases:

$$\begin{aligned} \mathcal{L}[\![\cdot]\!] &:: \text{FGen a} \rightarrow \text{Set String} \\ \mathcal{L}[\![\text{Void}]\!] &= \emptyset \\ \mathcal{L}[\![\text{Pure a}]\!] &= \varepsilon \\ \mathcal{L}[\![\text{Map f x}]\!] &= \mathcal{L}[\![x]\!] \\ \mathcal{L}[\![\text{Pair x y}]\!] &= \{s \cdot t \mid s \in \mathcal{L}[\![x]\!] \wedge t \in \mathcal{L}[\![y]\!]\} \\ \mathcal{L}[\![\text{Select xs}]\!] &= \{c \cdot s \mid (c, x) \in \text{xs} \wedge s \in \mathcal{L}[\![x]\!]\} \end{aligned}$$

⁵For algebraists: free generators are “free,” in the sense that they admit unique structure-preserving maps to other “generator-like” structures. In particular, the $\mathcal{G}[\![\cdot]\!]$ and $\mathcal{P}[\![\cdot]\!]$ maps are canonical. For the sake of space, we do not explore these ideas rigorously.

⁶Note that the arguments to Map flipped relative to MapR from §2.

3.2.3 Smart Constructors and Simplified Forms. Free generators admit a useful *simplified form*. We ensure that generators are simplified by requiring that free generators are built with *smart constructors*.

Instead of Pair, users should pair free generators with \otimes :

$(\otimes) :: \text{FGen a} \rightarrow \text{FGen b} \rightarrow \text{FGen (a, b)}$

Void \otimes _ = Void

_ \otimes Void = Void

Pure a \otimes y = $(\lambda b \rightarrow (a, b)) \langle \$ \rangle y$

x \otimes Pure b = $(\lambda a \rightarrow (a, b)) \langle \$ \rangle x$

x \otimes y = Pair x y

which makes sure that Void and Pure are simplified as much as possible with respect to Pair. For example, Pure a \otimes Pure b will simplify to Pure (a, b).

Next, $\langle \$ \rangle$ is a version of Map that does similar simplifications:

$(\langle \$ \rangle) :: (a \rightarrow b) \rightarrow \text{FGen a} \rightarrow \text{FGen b}$

f $\langle \$ \rangle$ Void = Void

f $\langle \$ \rangle$ Pure a = Pure (f a)

f $\langle \$ \rangle$ x = Map f x

We define pure and $\langle * \rangle$ to make FGen an applicative functor:

pure :: a → FGen a

pure = Pure

$(\langle * \rangle) :: \text{FGen (a} \rightarrow \text{b)} \rightarrow \text{FGen a} \rightarrow \text{FGen b}$

f $\langle * \rangle$ x = $(\lambda (f, x) \rightarrow f x) \langle \$ \rangle (f \otimes x)$

The smart constructor Select looks like:

select :: [(Char, FGen a)] → FGen a

select xs =

case filter $(\lambda (_, p) \rightarrow p \neq \text{Void})$ xs of

xs | xs == [] || duplicates (map fst xs) → ⊥

xs → Select xs

Unlike the other smart constructors, select can return ⊥ and fail. This ensures that the operations on generators defined later in this section will be well-formed.

Finally, we define:

void :: FGen a

void = Void

for consistency.

When a generator constructed using only these smart constructors, we say it is in simplified form.

3.2.4 Examples. We can generalize our definitions from earlier in this section to get a single free generator fgenTree that subsumes genTree and parseTree:

```

fgenTree :: Int → FGen Tree
fgenTree 0 = pure Leaf
fgenTree h =
  select [ (l, pure Leaf),
           (n, Node ⟨$⟩ fgenInt
                ⟨*⟩ fgenTree (h - 1)
                ⟨*⟩ fgenTree (h - 1)) ]

```

Excitingly, even though we are building a data structure we can write it like a program with exactly the same abstractions that are used for constructing generators and parsers.

Remark. One might be concerned that these free generators can be quite large, growing exponentially in h . We are able to avoid most size-related issues in HASKELL due to laziness—the parts of the structure that are not yet needed are left uninterpreted—but we recognize that relying on laziness is a bit unsatisfying. Luckily it is straightforward to share the recursive calls to `fgenExpr (h - 1)` between all of the branches of the `Select` node (and the `Pairs` below that) in languages with pointers or references. This avoids any blowup.

Another example of a free generator produces random terms of a simply-typed lambda-calculus:

```

fgenExpr :: Int → FGen Expr
fgenExpr 0 =
  select [ (i, Lit ⟨$⟩ fgenInt ),
           (v, Var ⟨$⟩ fgenVar ) ]
fgenExpr h =
  select [ (i, Lit ⟨$⟩ fgenInt ),
           (p, Plus ⟨$⟩ fgenExpr (h - 1)
                ⟨*⟩ fgenExpr (h - 1)),
           (l, Lam ⟨$⟩ fgenType
                ⟨*⟩ fgenExpr (h - 1)),
           (a, App ⟨$⟩ fgenExpr (h - 1)
                ⟨*⟩ fgenExpr (h - 1)),
           (v, Var ⟨$⟩ fgenVar ) ]

```

Structurally this is very similar to the previous generator, it just has more cases and more choices. Our lambda calculus is constructed with de Bruijn indices for variables and has integers and functions as values. The “raw” untyped language is exceedingly simple, but generating well-typed terms is still fairly difficult. We use this example as one of our case studies in §6.

3.3 Interpreting Free Generators

A free generator does not *do* anything on its own—it is simply a data structure. In this section, we see the formal definitions of the interpretation functions that we mentioned in §2 and prove a theorem that links those interpretations together.

3.3.1 As a Generator of Values. The most natural way to interpret a free generator is as a QUICKCHECK generator—that is, as a distribution over data structures. We define the

generator interpretation of a free generator to be:

```

G[·] :: FGen a → Gen a
G[Void] = ⊥
G[Pure v] = pure v
G[Map f x] = f ⟨$⟩ G[x]
G[Pair x y] =
  (λx y → (x, y)) ⟨$⟩ G[x] ⟨*⟩ G[y]
G[Select xs] =
  oneof (map (λ (_, x) → G[x]) xs)

```

Notice that the implementation of this interpretation is quite straightforward: in most cases, it simply maps the “AST node” version of an applicative operation to the concrete version implemented by `Gen`.

One detail worth noting is that the interpretation behaves poorly on `Void`, but we can prove a lemma to show that this does not cause problems in practice:

Lemma 3.1. *If a free generator g is simplified, g contains `Void` $\iff g = \text{Void}$.*

Proof. By induction on the structure of g and inspection of the smart constructors. \square

Thus we can conclude that as long as g is in simplified form and not `Void`, $G[g]$ is defined.

Example 3.2. $G[\text{fgenTree } 5]$ is equivalent to `genTree 5`.

3.3.2 As Parser of Choice Sequences. Of course, there would be no point in defining free generators if we were only going to interpret them as QUICKCHECK generators. We can make use of the choice labels using the free generator’s *parser interpretation*—in other words, viewing it as a parser of choices as we originally wanted. The translation looks like:

```

P[·] :: FGen a → Parser a
P[Void] = λs → Nothing
P[Pure a] = pure a
P[Map f x] = f ⟨$⟩ P[x]
P[Pair x y] =
  (λx y → (x, y)) ⟨$⟩ P[x] ⟨*⟩ P[y]
P[Select xs] =
  choice (map (λ (c, x) → (c, P[x])) xs)

```

This definition uses the representation of parsers as functions of type `String → Maybe (a, String)` that we saw earlier, and, just like $G[·]$, mostly just maps between applicative operations.

Example 3.3. $P[\text{fgenTree } 5]$ is equivalent to `parseTree 5`.

3.3.3 As a Generator of Choice Sequences. Our final interpretation of free generators is a sort of dual to the previous one. Instead of throwing away randomness in favor of deterministic parsing, we can extract only the random distribution and ignore everything about how result values are constructed. We define the *choice distribution* of a free

generator to be:

```

C[·] :: FGen a → Gen String
C[Void]    = ⊥
C[Pure a]  = pure ε
C[Map f x] = C[x]
C[Pair x y] =
  (λs t → s · t) ⟨$⟩ C[x] ⟨*⟩ C[y]
C[Select xs] =
  oneof (map (λ (c, x) → (λs → c · s) ⟨$⟩ C[x]) xs)
    
```

This definition is not quite as obvious as the previous two, but it is still natural. We can think of the result of this interpretation as a distribution over $\mathcal{L}[g]$. The language of a free generator is exactly those choice sequences that the generator interpretation can make and the parser interpretation can parse.

3.3.4 Factoring Generators. These different interpretations of free generators are closely related to one another, and in particular we can reconstruct $\mathcal{G}[\cdot]$ from $\mathcal{P}[\cdot]$ and $\mathcal{C}[\cdot]$. In essence, this means that a free generator's generator interpretation can be factored into a distribution over choice sequences and a parser of those sequences.

To make this more precise, we need a notion of equality for generators like the ones produced via $\mathcal{G}[\cdot]$. We say two QUICKCHECK generators are *equivalent*, written $g_1 \equiv g_2$, if and only if the generators represent the same distribution over values. This is coarser notion than program equality, since two generators might produce the same distribution of values in different ways.

With this in mind, we can state and prove the relationship between different interpretations of free generators:

Theorem 3.4 (Factoring). *Every simplified free generator can be factored into a parser and a distribution over choice sequences. In other words, for all simplified free generators $g \neq \text{Void}$,*

$$\mathcal{P}[g] \langle \$ \rangle \mathcal{C}[g] \equiv (\lambda x \rightarrow \text{Just } (x, \varepsilon)) \langle \$ \rangle \mathcal{G}[g].$$

Proof sketch. We proceed by induction on the structure of g .

Case $g = \text{Pure } a$. Straightforward.

Case $g = \text{Map } f \ x$. Straightforward.

Case $g = \text{Pair } x \ y$. This case is the most interesting one. The difficulty is that it is not immediately obvious why $\mathcal{P}[\text{Pair } x \ y] \langle \$ \rangle \mathcal{C}[\text{Pair } x \ y]$ should be a function of $\mathcal{P}[x] \langle \$ \rangle \mathcal{C}[x]$ and $\mathcal{P}[y] \langle \$ \rangle \mathcal{C}[y]$. Showing the correct relationship requires a lemma that says that for any sequence s generated by $\mathcal{C}[x]$ and an arbitrary sequence t , there is some a such that $\mathcal{P}[x] (s \cdot t) = \text{Just } (a, t)$.

Case $g = \text{Select } xs$. The reasoning in this case is a bit subtle, since it requires certain operations to commute with Select, but the details are not particularly instructive.

See Appendix B for the full proof. \square

A natural corollary of Theorem 3.4 is the following:

Corollary 3.5. *Any applicative generator, γ , written in terms of pure functions, $\langle \$ \rangle$, pure, $\langle * \rangle$, and oneof, can be factored into a parser and distribution over choice sequences.*

Proof. Translate γ into a free generator, g , by replacing operations with the equivalent smart constructor. (For oneof, draw unique labels for each choice and use select.) By induction, $\mathcal{G}[g] = \gamma$.

The resulting free generator can be factored into a parser and a choice distribution via Theorem 3.4. Thus,

$$(\lambda x \rightarrow \text{Just } (x, \varepsilon)) \langle \$ \rangle \gamma \equiv \mathcal{P}[g] \langle \$ \rangle \mathcal{C}[g],$$

and γ can be factored as desired. \square

This corollary gives a concrete way to view the connection between applicative generators and parsers.

3.4 Replacing a Generator's Distribution

Since a generator g can be factored using $\mathcal{C}[\cdot]$ and $\mathcal{P}[\cdot]$, we can explore what it would look like to modify a generator's distribution (i.e., change or replace $\mathcal{C}[\cdot]$) without having to modify the entire generator.

Suppose we have some other distribution that we want our choices to follow, represented as a function from a history of choices to a generator of next choices that we call DistF. (If the next choice is Nothing, then generation stops.) We can pair a distribution function with a “current” history, to get a formal definition of a custom choice distribution:

type DistF = String → Gen (Maybe Char)

type Dist = (String, DistF)

A Dist may be arbitrarily complex—in particular, it might contain information obtained from a machine learning model, example-based tuning, or some other automated tuning process. How would we use such a distribution in place of the standard distribution given by $\mathcal{C}[\cdot]$?

The solution is to replace $\mathcal{C}[\cdot]$ with a distribution to yield the definition:

$\overline{\mathcal{G}}[\cdot] :: (\text{Dist}, \text{FGen } a) \rightarrow \text{Gen } (\text{Maybe } a)$

$\overline{\mathcal{G}}[(h, d), g] = \mathcal{P}[g] \langle \$ \rangle \text{genDist } h$

where

genDist h = d h >>= $\lambda x \rightarrow \text{case } x \text{ of}$

Nothing → pure h

Just c → genDist (h · c)

Whereas before we proved an equivalence between $\mathcal{G}[g]$ and $\mathcal{P}[g] \langle \$ \rangle \mathcal{C}[g]$, we can now use that relationship as a definition of what it means to interpret a generator under a new distribution.

Since replacing a free generator's distribution does not actually change the structure of the generator, we can have a different distribution for each use-case of the free generator. In a property-based testing scenario, one could imagine the

tester finely-tuning a distribution for each property that is carefully optimized to find bugs as quickly as possible.

4 Derivatives of Free Generators

In this section, we review the notion of Brzozowski derivative in formal language theory and we show that a similar operation exists for **free generators**. Free generator derivatives highlight the advantages of taking the correspondence between generators and parsers seriously.

4.1 Background: Derivatives of Languages

The *Brzozowski derivative* [1] of a formal language L with respect to some choice c is defined as

$$\delta_c L = \{s \mid c \cdot s \in L\}.$$

In other words, the derivative is the set of strings in L with c removed from the front. For example,

$$\delta_a \{abc, aaa, bba\} = \{bc, aa\}.$$

Many language representations have syntactic transformations that correspond to Brzozowski derivatives. For example, we can take the derivative of a regular expression:

$$\begin{aligned} \delta_c \emptyset &= \emptyset & \nu \emptyset &= \emptyset \\ \delta_c \varepsilon &= \varepsilon & \nu \varepsilon &= \varepsilon \\ \delta_c c &= \varepsilon \quad (c = c) & \nu c &= \emptyset \\ \delta_c d &= \emptyset \quad (c \neq d) & \nu(r_1 + r_2) &= \nu r_1 + \nu r_2 \\ \delta_c(r_1 + r_2) &= \delta_c r_1 + \delta_c r_2 & \nu(r_1 \cdot r_2) &= \nu r_1 \cdot \nu r_2 \\ \delta_c(r_1 \cdot r_2) &= \delta_c r_1 \cdot r_2 + \nu r_1 \cdot \delta_c r_2 & \nu(r^*) &= \varepsilon \\ \delta_c(r^*) &= \delta_c r \cdot r^* \end{aligned}$$

The ν operator, used in the “.” rule and defined on the right, determines the *nullability* of an expression (whether or not it accepts ε). As one would hope, if r has language L , it is always the case that $\delta_c r$ has language $\delta_c L$.

4.2 The Free Generator Derivative

Since free generators define a language (given by $\mathcal{L}[\![\cdot]\!]$), it makes sense to ask: can we take their derivatives? Yes! We define the *derivative* of a free generator to be:

$$\begin{aligned} \delta &:: \text{Char} \rightarrow \text{FGen } a \rightarrow \text{FGen } a \\ \delta \text{Void} &= \text{void} \\ \delta (\text{Pure } v) &= \text{void} \\ \delta (\text{Map } f \ x) &= f \ \langle \$ \rangle \ \delta_c x \\ \delta (\text{Pair } x \ y) &= \delta_c x \otimes y \\ \delta (\text{Select } xs) &= \text{if } (c, x) \in xs \text{ then } x \text{ else void} \end{aligned}$$

These definitions should be mostly intuitive. The derivative of a generator that does not make a choice (i.e., **Void** and **Pure**) is void, since the corresponding language would be

empty. The derivative commutes with **Map** since the transformation affects choices, not the final result. **Select**’s derivative is just the argument generator corresponding to the appropriate choice.

The one potentially confusing case is the one for **Pair**. We have defined the derivative of a pair of generators by taking the derivative of the first generator in the pair and leaving the second unchanged, but this is inconsistent with the case for “.” in the regular expression derivative. What happens when the first generator’s language is nullable? Luckily, our **simplified form** clears up the confusion: if **Pair** $x \ y$ is in simplified form, x is not nullable. This is a simple corollary of Lemma 4.1.

Lemma 4.1. *If a free generator g is in simplified form, then either $g = \text{Pure } a$ or $\nu(\mathcal{L}[\![g]\!]) = \emptyset$.*

Proof sketch. See Appendix A. \square

Remark. The derivative of a simplified generator is simplified. This follows simply from the definition, since we only use smart constructors and parts of the original generator to build the derivative generators. By induction, this also means that repeated derivatives preserve simplification.

Besides clearing up the issue with **Pair**, Lemma 4.1 also says that we can define *nullability* for free generators simply as:

$$\begin{aligned} \nu &:: \text{FGen } a \rightarrow \text{Set } a \\ \nu(\text{Pure } v) &= \{v\} \\ \nu g &= \emptyset \quad (g \neq \text{Pure } v) \end{aligned}$$

Note that we get a bit more information here than we do from regular expression nullability. For a regular expression r , νr is either \emptyset or ε . Here, we allow the null check to return either \emptyset or the singleton set containing the value in the **Pure** node. This means that ν for free generators extracts a value that can be obtained by making no further choices.

With the derivative operation defined, we can prove a concrete theorem that says our definition of derivative acts the way we expect:

Theorem 4.2 (Language Consistency). *For all simplified free generators g and choices c ,*

$$\delta_c \mathcal{L}[\![g]\!] = \mathcal{L}[\![\delta_c g]\!].$$

Proof sketch. The proof proceeds by mostly straightforward induction, but it is interesting enough to be worth writing out. See Appendix C. \square

This theorem says that the derivative of a free generator’s language is the same as the language of its derivative.

Besides consistency with respect to the language interpretation, we want our derivative to preserve generator outputs for a given sequence of choices. If a free generator making choices

ntll yields **Node True Leaf Leaf**,

we would like for the derivative of that free generator with respect to `n` to produce the same value after choosing `ttl`. We can formalize this via the parser interpretation:

Theorem 4.3 (Value Consistency). *For all **simplified free generators** g , choice sequences s , and choices c ,*

$$\mathcal{P}[\delta_c g] s = \mathcal{P}[g] (c \cdot s).$$

Proof sketch. Again, this proof is mostly straightforward. See Appendix D. \square

The upshot of this theorem is that derivatives do not fundamentally change the results of a free generator, they only fix a particular choice.

These two consistency theorems together mean that we can simulate a free generator's choices by taking repeated derivatives. Each derivative fixes a particular choice, so a sequence of derivatives fixes a choice sequence.

5 Generating Valid Results with Gradients

In this section, we put the theory of **free generators** and their **derivatives** into practice. We introduce CHOICE GRADIENT SAMPLING (CGS), an algorithm for generating data that satisfies a validity condition.

5.1 The Algorithm

Given a simple free generator, **CHOICE GRADIENT SAMPLING** previews choices using derivatives. In fact, it previews all possible choices, essentially taking the *gradient* of the free generator. (This is akin to the gradient in calculus, which is a vector of partial derivatives with respect to each variable.) We can write

$$\nabla g = \langle \delta_a g, \delta_b g, \delta_c g \rangle$$

for the gradient of g with respect to alphabet $\{a, b, c\}$. Each derivative in the gradient can be sampled, using $\mathcal{G}[\cdot]$, to get a sense of how good or bad the respective choice was. This provides a metric that guides the algorithm to valid inputs.

With this intuition in mind, we present the CGS algorithm, shown in Figure 4, which searches for valid results using repeated free generator gradients.

The intuition from earlier plays out in lines 7–14, and is shown pictorially in Figure 5. We take the gradient of g by taking the derivative with respect to each possible choice, in this case `a`, `b`, and `c`. Then we evaluate each of the derivatives by interpreting the generator with $\mathcal{G}[\cdot]$, sampling values from the resulting generator, and counting how many of those results are valid with respect to φ . At the end of sampling, we have values f_a , f_b , and f_c , which we can think of as the “fitness” of each choice. We then pick a choice randomly, weighted based on fitness, and continue until our choices produce a valid output.

Critically, we avoid wasting effort by saving the samples (\mathcal{V}) that we use to evaluate the gradients. Many of those samples will be valid results that we can use, so there is no reason to throw them away.

```

1:  $g \leftarrow G$ 
2:  $\mathcal{V} \leftarrow \emptyset$ 
3: while true do
4:   if  $vg \neq \emptyset$  then return  $vg \cup \mathcal{V}$ 
5:   if  $g = \text{Void}$  then  $g \leftarrow G$ 
6:    $\nabla g \leftarrow \langle \delta_c g \mid c \in C \rangle$   $\triangleright \nabla g$  is the gradient of  $g$ 
7:   for  $\delta_c g \in \nabla g$  do
8:     if  $\delta_c g = \text{Void}$  then
9:        $V \leftarrow \emptyset$ 
10:    else
11:       $x_1, \dots, x_N \leftarrow \mathcal{G}[\delta_c g]$   $\triangleright$  Sample  $\mathcal{G}[\delta_c g]$ 
12:       $V \leftarrow \{x_j \mid \varphi(x_j)\}$ 
13:       $f_c \leftarrow |V|$   $\triangleright f_c$  is the fitness of  $c$ 
14:       $\mathcal{V} \leftarrow \mathcal{V} \cup V$ 
15:    if  $\max_{c \in C} f_c = 0$  then
16:      for  $c \in C$  do  $f_c \leftarrow 1$ 
17:     $g \leftarrow \text{weightedChoice} \{ (f_c, \delta_c g) \mid c \in C \}$ 
    (Where  $N \in \mathbb{N}$  is the sample rate constant.)
    
```

Figure 4. CHOICE GRADIENT SAMPLING: Given a free generator G in **simplified form** and a validity predicate φ , this algorithm produces a set of outputs that all satisfy $\varphi(x)$.

5.2 Modified Distributions

Interestingly, this algorithm works equally well for free generators whose distributions have been replaced (as discussed in §3.4). Recall that we modify the distribution of a free generator g by pairing it with a pair of a history h and a distribution function d . We can define the derivative of such a structure to be:

$$\delta_c((h, d), g) = ((h \cdot c, d), \delta_c g)$$

We take the derivative of g and internalize c into the distribution's history.

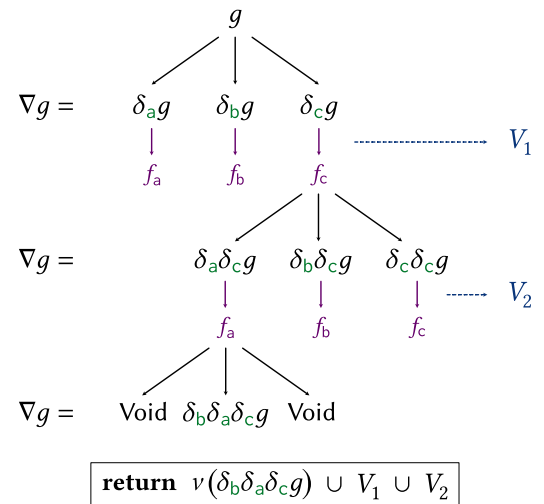


Figure 5. The main loop of CHOICE GRADIENT SAMPLING.

Furthermore, we can say that a modified generator’s nullable set is the same as the nullable set of the underlying generator.

These definitions, along with the ones in §3.4, are enough to replicate CHOICE GRADIENT SAMPLING for a free generator with an external distribution.

6 Exploratory Evaluation

This paper is primarily about the theory of **free generators** and their **derivatives**, but we were curious to see how well CHOICE GRADIENT SAMPLING would perform on a few property-based testing benchmarks. This section describes our experiments in detail and shows that, with a few caveats, CGS presents a promising solution for the valid generation problem.

6.1 Experimental Setup

Our experiments compare CGS to rejection sampling. Rejection sampling takes a naïve generator, samples from it, and simply discards any results that are not valid; it is the default way that QUICKCHECK handles properties with preconditions when no bespoke generator is available. We chose this over more state-of-the-art comparisons [3, 15], because our primary goal was to validate our theory, not to produce a production-ready tool.

We use four simple free generators to test four different benchmarks: **BST**, **SORTED**, **AVL**, and **STLC**. Information about each of these benchmarks is given in Table 1.

Each of these benchmarks requires a free generator as a starting point—in order to avoid any generator cleverness complicating the final results, our generators follow the respective inductive data types as closely as possible. For example, `fgenTree`, shown in §3 and used in the **BST** benchmark, follows the structure of `Tree` exactly. In addition, each benchmark requires an appropriate sample-rate constant N , based on the complexity of the particular task at hand. A higher value for N means that estimates of choice fitness will be better, at the cost of increased time spent sampling. In practice this could be determined by the programmer or via heuristics based on the data type.

6.2 Results

We ran CGS and REJECTION on each benchmark for one minute and recorded the unique valid values produced. We focused on unique values because duplicates are rarely useful (in property-based testing of pure programs, for example, duplicates add nothing). The totals are presented in Table 2.

	BST	SORTED	AVL	STLC
REJECTION	9,785	6,485	156	107,256
CGS	21,763	59,540	221	315,275

Table 2. Unique valid values generated in 60 seconds.

These numbers are quite promising—CGS is always able to generate more unique values than REJECTION in the same amount of time, and it often generates *significantly* more. (Later in this section we explore the **AVL** benchmark to understand why it does not perform as well.)

In addition to the raw numbers, we tracked some other metrics; the charts in Figure 6 give some deeper insights. The first plot (“Unique Terms over Time”) shows that after one minute, CGS has not yet “run out” of unique terms to generate. Additionally “Normalized Size Distribution” chart shows that CGS also generates larger terms on average. This is great from the perspective of property-based testing, where test size is often positively correlated with bug-finding power since larger test inputs tend to exercise more of the implementation code.




Figure 6. Unique values and term sizes for the **STLC** benchmark.

Charts for the remaining benchmarks are in Appendix E.

6.3 Measuring Diversity

In a testing scenario, we care about more than just the number of valid test inputs generated in a period of time—we care about the *diversity* of those inputs. A more diverse test suite will find more bugs more quickly. We wanted to know how CGS impacts the variety of generated values.

To measure this, we compute a *diversity metric* over our values. This metric relies on the fact that each value is roughly isomorphic to the choice sequence that generated it. For

	Free Generator	Validity Condition	N	Depth
BST	Binary trees with values 0–9	Is a valid BST	50	5
SORTED	Lists with values 0–9	Is sorted	50	20
AVL	Binary trees with values and stored heights 0–9	Is a valid AVL tree (balanced)	500	5
STLC	Arbitrary ASTs for λ -terms	Is well-typed	400	5

Table 1. Overview of benchmarks.

example, in the case of **BST**, the sequence `n5l6l1` can be parsed to produce Node 5 Leaf (Node 6 Leaf Leaf) and a simple in-order traversal can recover `n5l6l1` again. This means that it is safe to measure choice sequence diversity as a proxy for value diversity.

We estimated the average Levenshtein distance [12] between pairs of choice sequences in the values generated by each of our algorithms. We could not compute the true mean, so we settled for the mean of a sample of 3000 pairs from each set of valid values. The results are summarized in Table 3.

	BST	SORTED	AVL	STLC
REJ.	7.70(1.71)	4.80(1.15)	4.42(2.01)	12.24(4.55)
CGS	8.89(1.95)	7.28(1.92)	4.35(1.98)	13.62(4.72)

Table 3. Average Levenshtein distance between pairs of choice sequences.

These results are not as dramatic as we had hoped, with diversity only increasing by around 10% for **STLC** and around the same for **BST**. One explanation for this effect rests on the way CGS retains intermediate samples. Late in the algorithm, when a large prefix of choices is already fixed, samples from the derivative generators will tend to be similar. This likely results in some clusters of inputs that are all valid, but that only explore one particular shape of input.

This is disappointing, it is not the end of the world. It is already common practice to test clusters of similar inputs in certain fuzzing contexts [9], so the fact that CGS does this is not unusual. In fact, this method has been shown to be effective at finding bugs in some cases. Additionally, for most of our benchmarks (again, we return to **AVL** in a moment) CGS does increase diversity of tests; combined with the sheer number of valid inputs available, this means that CGS covers a slightly larger space of tests much more thoroughly. This effect should lead to better bug-finding in testing scenarios.

6.4 The Problem with AVL

Of course, the outlier in all of these results seems to be the **AVL** benchmark. CGS only manages to find a modest number of extra valid AVL trees, and their pairwise diversity is actually slightly worse than that of rejection sampling. Why might this be? We suspect that this is effect arises because AVL trees are so difficult to find randomly. Balanced binary search trees are hard to generate on their own, and

AVL trees are even more difficult because the generator must guess the correct height to cache at each node. This is why rejection sampling only finds 156 AVL trees in the time it takes to find 9,762 binary search trees.

This all means that CGS unlikely find *any* valid trees while sampling. In particular, the check in line 15 of Figure 5 will often be true, meaning that choices made uniformly at random rather than guided by the fitness of the appropriate derivatives. We could mitigate this problem by significantly increasing the sample rate constant N , but then sampling time would likely dominate generation time and result in worse performance overall.

Ultimately this failure is disappointing, but not wholly surprising—CGS approximates generator fitness via sampling, and that approximation cannot always be accurate. For now it seems that especially hard-to-satisfy predicates are out of reach for this specific algorithm. Still, gradient-based algorithms might not all fail in this way. In §8 we discuss ideas for alternative algorithms that may perform better in tricky cases.

7 Related Work

We briefly discuss some other theoretical work that is similar to ours, as well as a few other approaches to the valid generation problem.

7.1 Similar Approaches

The PYTHON library HYPOTHESIS [13] implements its generators by parsing a stream of random bits. This is another point in favor of thinking of generators as parsers, but it is more of an implementation detail and is not explored formally.

In, *Generating constrained random data with uniform distribution* [3], Claessen et al. present a structure that is structurally similar to our free generator structure, but which is used in a very different way. They primarily use the syntactic structure of their generators (they call them “spaces”) to control the size distribution of generated outputs; in particular, spaces do not make choice information explicit in the way free generators do. Claessen et al.’s generation approach uses HASKELL’s laziness, rather than derivatives and sampling, to prune unhelpful paths in the generation process. We will consider incorporating some of these ideas into our work in the future.

The CLOTHO [5] library is another interesting point in the generator design space, though it does not rely on parsing.

7.2 The Valid Generation Problem

Many other partial solutions to the valid generation problem exist.

The domain specific generator languages provided by QUICKCHECK [8] make it easier to write manual generators that produce valid inputs by construction. We avoid manual approaches like this in the hopes of making techniques like property-based testing more accessible to those do not have experience writing their own generators.

When validity predicates are expressed as inductive relations, approaches like the one in *Generating Good Generators for Inductive Relations* [10] are extremely powerful. Unfortunately, most programming languages cannot express inductive relations that capture the kinds of preconditions that we care about.

Some approaches have tried to use machine learning to automatically generate valid inputs. LEARN&FUZZ [6] generates valid data using a recurrent neural network. While the results are promising, this solution seems to work best when a large corpus of inputs is already available and the validity condition is more structural than semantic. In the same vein, RLCHECK [15] uses reinforcement learning to guide a generator to valid inputs. We hope to incorporate ideas from RLCHECK into this work in the future, but we felt that getting the theory of free generators right was a critical first step.

8 Future Directions

There are a number of exciting paths forward from this work; some continue our theoretical exploration and others look towards algorithmic improvements.

8.1 Bidirectional Free Generators

We believe that we have only scratched the surface of what is possible with **free generators**. One concrete next step is to merge the theory of free generators with the emerging theory of *ungenerators* [7]. Goldstein expresses generators that can be run both forward (to generate values as usual) and *backward*. In the backward direction, the program takes a value that the generator might have generated and “un-generates” it to give a sequence of choices that the generator might have made when generating that value.

Free generators are quite compatible with these ideas, and turning a free generator into a bidirectional generator that can both generate and ungenerate should be fairly straightforward. From there, we can build on the ideas in the ungenerators work and use the backward direction of the generator to learn a distribution of choices that approximates some user-provided samples of “desirable” values. Used in conjunction with the extended algorithm from §5.2, this would give a better starting point for generation with little extra work from the user.

8.2 Algorithmic Optimizations

In §6.4, we uncover some problems with the **CHOICE GRADIENT SAMPLING** algorithm. Recall that because CGS evaluates derivatives via sampling, it does poorly when validity conditions are particularly difficult to satisfy. This begs the question: might it be possible to evaluate the fitness of a derivative without naïvely sampling?

One potential angle involves staging the sampling process. Given a free generator with a depth parameter, we can first evaluate choices on generators of size 1, then evaluate choices with size 2, etc. These intermediate stages would make gradient sampling more successful at larger sizes, and might significantly improve the results on benchmarks like **AVL**. Unfortunately, this kind of approach might perform poorly on benchmarks like **STLC** where the validity condition is not uniform: size-1 generators would avoid generating variables, leading larger generators to avoid variables as well. In any case, we think this design space is worth exploring.

8.3 Making Choices with Neural Networks

Another algorithmic optimization is a bit farther afield: we think it may be possible to use recurrent neural networks (RNNs) to improve our generation procedure.

As **CHOICE GRADIENT SAMPLING** makes choices, it generates useful data about the frequencies with which choices should be made. Specifically, every iteration of the algorithm produces a pair of a history and a distribution over next choices that looks something like:

$$\text{abcca} \mapsto \{a : 0.3, b : 0.7, c : 0.0\}$$

In the course of CGS, this information is used once (to make the next choice) and then forgotten—what if there was a way to learn from it? Pairs like this could be used to train an RNN to make choices that are similar to the ones made by CGS.

There are still details to work out, including network architecture, hyper-parameters, etc., but in theory we could run CGS for a while, then train the model, and after that point only use the RNN to generate valid data. Setting things up this way would recover some of the time that is currently wasted by the constant sampling of derivative generators.

One could imagine a user writing a definition of a type and a predicate for that type, and then setting the model to train while they work on their algorithm. By the time the algorithm is finished and ready to test, the RNN model would be trained and ready to produce valid test inputs. A workflow like this could significantly increase adoption of property-based testing in industry and give developers more control over the quality of their software.

Free generators and their derivatives are powerful structures that give a unique and flexible perspective on random generation. Our formalism yields a useful algorithm and clarifies the folklore that a generator is a parser of randomness.

References

- [1] Janusz A Brzozowski. 1964. Derivatives of regular expressions. *Journal of the ACM (JACM)* 11, 4 (1964), 481–494.
- [2] Paolo Capriotti and Ambrus Kaposi. 2014. Free applicative functors. *arXiv preprint arXiv:1403.0749* (2014).
- [3] Koen Claessen, Jonas Duregård, and Michal H. Palka. 2015. Generating constrained random data with uniform distribution. *J. Funct. Program.* 25 (2015). <https://doi.org/10.1017/S0956796815000143>
- [4] Koen Claessen and John Hughes. 2000. QuickCheck: a lightweight tool for random testing of Haskell programs. In *Proceedings of the Fifth ACM SIGPLAN International Conference on Functional Programming (ICFP '00), Montreal, Canada, September 18-21, 2000*, Martin Odersky and Philip Wadler (Eds.). ACM, Montreal, Canada, 268–279. <https://doi.org/10.1145/351240.351266>
- [5] Pierce Darragh, William Gallard Hatch, and Eric Eide. 2021. Clotho: A Racket Library for Parametric Randomness. In *Functional Programming Workshop*. 3.
- [6] Patrice Godefroid, Hila Peleg, and Rishabh Singh. 2017. Learn&fuzz: Machine learning for input fuzzing. In *2017 32nd IEEE/ACM International Conference on Automated Software Engineering (ASE)*. IEEE, 50–59.
- [7] Harrison Goldstein. 2021. Ungenerators. In *ICFP Student Research Competition*. <https://harrisongoldste.in/papers/icfpsrc21.pdf>
- [8] John Hughes. 2007. QuickCheck testing for fun and profit. In *International Symposium on Practical Aspects of Declarative Languages*. Springer, 1–32.
- [9] Leonidas Lampropoulos, Michael Hicks, and Benjamin C. Pierce. 2019. Coverage guided, property based testing. *PACMPL* 3, OOPSLA (2019), 181:1–181:29. <https://doi.org/10.1145/3360607>
- [10] Leonidas Lampropoulos, Zoe Paraskevopoulou, and Benjamin C Pierce. 2017. Generating good generators for inductive relations. *Proceedings of the ACM on Programming Languages* 2, POPL (2017), 1–30.
- [11] Daan Leijen and Erik Meijer. 2001. Parsec: Direct style monadic parser combinators for the real world. (2001).
- [12] Vladimir I Levenshtein et al. 1966. Binary codes capable of correcting deletions, insertions, and reversals. In *Soviet physics doklady*, Vol. 10. Soviet Union, 707–710.
- [13] David R MacIver, Zac Hatfield-Dodds, et al. 2019. Hypothesis: A new approach to property-based testing. *Journal of Open Source Software* 4, 43 (2019), 1891.
- [14] Conor McBride and Ross Paterson. 2008. Applicative programming with effects. *Journal of functional programming* 18, 1 (2008), 1–13.
- [15] Sameer Reddy, Caroline Lemieux, Rohan Padhye, and Koushik Sen. 2020. Quickly generating diverse valid test inputs with reinforcement learning. In *ICSE '20: 42nd International Conference on Software Engineering, Seoul, South Korea, 27 June - 19 July, 2020*, Gregg Rothermel and Doo-Hwan Bae (Eds.). ACM, 1410–1421. <https://doi.org/10.1145/3377811.3380399>

Appendix

A Proof of Lemma 4.1

Lemma 4.1. *If a free generator g is in simplified form, then either $g = \text{Pure } a$ or $v(\mathcal{L}[\![g]\!]) = \emptyset$.*

Proof. We proceed by induction on the structure of g .

Case $g = \text{Void}$. Trivial.

Case $g = \text{Pure } a$. Trivial.

Case $g = \text{Pair } x \ y$. By our inductive hypothesis, $x = \text{Pure } a$ or $v\mathcal{L}[\![x]\!] = \emptyset$.

Since the smart constructor \otimes never constructs a Pair with Pure on the left, it must be that $v\mathcal{L}[\![x]\!] = \emptyset$.

Therefore, it must be the case that $v\mathcal{L}[\![\text{Pair } x \ y]\!] = \emptyset$.

Case $g = \text{Map } f \ x$.

Similarly to the previous case, our inductive hypothesis and simplification assumptions imply that $v\mathcal{L}[\![x]\!] = \emptyset$.

Therefore, $v\mathcal{L}[\![\text{Map } f \ y]\!] = \emptyset$.

Case $g = \text{Select } xs$.

It is always the case that $v\mathcal{L}[\![\text{Select } xs]\!] = \emptyset$.

Thus, we have shown that every simplified free generator is either Pure a or has an empty nullable set. \square

B Proof of Theorem 3.4

Lemma B.1. *Pairing two **parser interpretations** and mapping over the concatenation of the associated choice distributions is equal to a function of the two parsers mapped over the distributions individually. Specifically, for all **simplified free generators** x and y ,*

$$\begin{aligned} ((\lambda x y \rightarrow (x, y)) \langle \$ \rangle \mathcal{P} \llbracket x \rrbracket \langle * \rangle \mathcal{P} \llbracket y \rrbracket) \langle \$ \rangle ((\cdot) \langle \$ \rangle \mathcal{C} \llbracket x \rrbracket \langle * \rangle \mathcal{C} \llbracket y \rrbracket) &\equiv (\lambda a_{\perp} b_{\perp} \rightarrow \mathbf{case} (a_{\perp}, b_{\perp}) \mathbf{of} \\ &\quad (\mathbf{Just} (a, _), \mathbf{Just} (b, _)) \rightarrow \mathbf{Just} ((a, b), \varepsilon) \\ &\quad _ \rightarrow \mathbf{Nothing}) \\ &\quad \langle \$ \rangle (\mathcal{P} \llbracket x \rrbracket \langle \$ \rangle \mathcal{C} \llbracket x \rrbracket) \langle * \rangle (\mathcal{P} \llbracket y \rrbracket \langle \$ \rangle \mathcal{C} \llbracket y \rrbracket) \end{aligned}$$

Proof. First, note that for any simplified generator, g , if $\mathcal{C} \llbracket g \rrbracket$ generates a string s , for any other string t $\mathcal{P} \llbracket g \rrbracket (s \cdot t) = \mathbf{Just} (a, t)$ for some value a . This can be shown by induction on the structure of g .

Now, assume $\mathcal{C} \llbracket x \rrbracket$ generates a string s , and $\mathcal{C} \llbracket y \rrbracket$ generates t . This means that $(\cdot) \langle \$ \rangle \mathcal{C} \llbracket x \rrbracket \langle * \rangle \mathcal{C} \llbracket y \rrbracket$ generates $s \cdot t$.

By the above fact, it is simple to show that both sides of the above equation simplify to $\mathbf{Just} ((a, b), \varepsilon)$ for some values a and b that depend on the particular interpretations of x and y .

Since this is true for any s and t that the choice distributions generate, the desired fact holds. \square

Theorem 3.4. *Every **simplified free generator** can be factored into a parser and a distribution over choice sequences. In other words, for all **simplified free generators** $g \neq \mathbf{Void}$,*

$$\mathcal{P} \llbracket g \rrbracket \langle \$ \rangle \mathcal{C} \llbracket g \rrbracket \equiv (\lambda x \rightarrow \mathbf{Just} (x, \varepsilon)) \langle \$ \rangle \mathcal{G} \llbracket g \rrbracket.$$

Proof. We proceed by induction on the structure of g .

Case $g = \mathbf{Pure} \ a$.

$$\begin{aligned} \mathcal{P} \llbracket \mathbf{Pure} \ a \rrbracket \langle \$ \rangle \mathcal{C} \llbracket \mathbf{Pure} \ a \rrbracket &\equiv \mathbf{pure} (\mathbf{Just} (a, \varepsilon)) && \text{(by defn)} \\ &\equiv (\lambda x \rightarrow \mathbf{Just} (x, \varepsilon)) \langle \$ \rangle \mathcal{G} \llbracket \mathbf{Pure} \ a \rrbracket && \text{(by defn)} \end{aligned}$$

Case $g = \mathbf{Pair} \ x \ y$.

$$\begin{aligned} \mathcal{P} \llbracket \mathbf{Pair} \ x \ y \rrbracket \langle \$ \rangle \mathcal{C} \llbracket \mathbf{Pair} \ x \ y \rrbracket &\equiv ((\lambda x y \rightarrow (x, y)) \langle \$ \rangle \mathcal{P} \llbracket x \rrbracket \langle * \rangle \mathcal{P} \llbracket y \rrbracket) \langle \$ \rangle ((\cdot) \langle \$ \rangle \mathcal{C} \llbracket x \rrbracket \langle * \rangle \mathcal{C} \llbracket y \rrbracket) && \text{(by defn)} \\ &\equiv (\lambda a_{\perp} b_{\perp} \rightarrow \mathbf{case} (a_{\perp}, b_{\perp}) \mathbf{of} \\ &\quad (\mathbf{Just} (a, _), \mathbf{Just} (b, _)) \rightarrow \mathbf{Just} ((a, b), \varepsilon) \\ &\quad _ \rightarrow \mathbf{Nothing}) \\ &\quad \langle \$ \rangle (\mathcal{P} \llbracket x \rrbracket \langle \$ \rangle \mathcal{C} \llbracket x \rrbracket) \langle * \rangle (\mathcal{P} \llbracket y \rrbracket \langle \$ \rangle \mathcal{C} \llbracket y \rrbracket) && \text{(by Lemma B.1)} \\ &\equiv (\lambda a_{\perp} b_{\perp} \rightarrow \mathbf{case} (a_{\perp}, b_{\perp}) \mathbf{of} \\ &\quad (\mathbf{Just} (a, _), \mathbf{Just} (b, _)) \rightarrow \mathbf{Just} ((a, b), \varepsilon) \\ &\quad _ \rightarrow \mathbf{Nothing}) \\ &\quad \langle \$ \rangle ((\lambda x \rightarrow \mathbf{Just} (x, \varepsilon)) \langle \$ \rangle \mathcal{G} \llbracket x \rrbracket) \langle * \rangle ((\lambda x \rightarrow \mathbf{Just} (x, \varepsilon)) \langle \$ \rangle \mathcal{G} \llbracket y \rrbracket) && \text{(by IH)} \\ &\equiv (\lambda x \rightarrow \mathbf{Just} (x, \varepsilon)) \langle \$ \rangle ((\lambda x y \rightarrow (x, y)) \langle \$ \rangle \mathcal{G} \llbracket x \rrbracket \langle * \rangle \mathcal{G} \llbracket y \rrbracket) && \text{(by app. properties)} \\ &\equiv (\lambda x \rightarrow \mathbf{Just} (x, \varepsilon)) \langle \$ \rangle \mathcal{G} \llbracket \mathbf{Pair} \ x \ y \rrbracket && \text{(by defn)} \end{aligned}$$

Case $g = \mathbf{Map} \ f \ x$.

$$\begin{aligned} \mathcal{P} \llbracket \mathbf{Map} \ f \ x \rrbracket \langle \$ \rangle \mathcal{C} \llbracket \mathbf{Map} \ f \ x \rrbracket &\equiv (f \langle \$ \rangle \mathcal{P} \llbracket x \rrbracket) \langle \$ \rangle \mathcal{C} \llbracket x \rrbracket && \text{(by defn)} \\ &\equiv f \langle \$ \rangle (\mathcal{P} \llbracket x \rrbracket \langle \$ \rangle \mathcal{C} \llbracket x \rrbracket) && \text{(by functor properties)} \\ &\equiv f \langle \$ \rangle ((\lambda x \rightarrow \mathbf{Just} (x, \varepsilon)) \langle \$ \rangle \mathcal{G} \llbracket x \rrbracket) && \text{(by IH)} \\ &\equiv (\lambda x \rightarrow \mathbf{Just} (x, \varepsilon)) \langle \$ \rangle (f \langle \$ \rangle \mathcal{G} \llbracket x \rrbracket) && \text{(by functor properties)} \\ &\equiv (\lambda x \rightarrow \mathbf{Just} (x, \varepsilon)) \langle \$ \rangle \mathcal{G} \llbracket \mathbf{Map} \ f \ x \rrbracket && \text{(by defn)} \end{aligned}$$

Case $g = \text{Select } xs$.

$$\begin{aligned}
\mathcal{P}[\![\text{Select } xs]\!] \langle \$ \rangle C[\![\text{Select } xs]\!] &\equiv (\text{choice } (\text{map } (\lambda(c, x) \rightarrow (c, \mathcal{P}[\![x]\!])) xs)) && (by \text{ defn}) \\
&\equiv \langle \$ \rangle \text{oneof } (\text{map } (\lambda(c, x) \rightarrow (c \cdot \langle \$ \rangle C[\![x]\!])) xs) && (by \text{ defn}) \\
&\equiv \text{oneof } (\text{map } (\lambda(c, x) \rightarrow && \\
&\quad (\text{choice } (\text{map } (\lambda(c, x) \rightarrow (c, \mathcal{P}[\![x]\!])) xs)) \circ (c \cdot \langle \$ \rangle C[\![x]\!] && \\
&\quad) xs) && (by \text{ generator properties}) \\
&\equiv \text{oneof } (\text{map } (\lambda(_, x) \rightarrow \mathcal{P}[\![x]\!] \langle \$ \rangle C[\![x]\!]) xs) && (by \text{ parser properties}) \\
&\equiv \text{oneof } (\text{map } (\lambda(_, x) \rightarrow (\lambda x \rightarrow \text{Just } (x, \varepsilon)) \langle \$ \rangle \mathcal{G}[\![x]\!]) xs) && (by \text{ IH}) \\
&\equiv (\lambda x \rightarrow \text{Just } (x, \varepsilon)) \langle \$ \rangle \text{oneof } (\text{map } (\lambda(_, x) \rightarrow \mathcal{G}[\![x]\!]) xs) && (by \text{ generator properties}) \\
&\equiv (\lambda x \rightarrow \text{Just } (x, \varepsilon)) \langle \$ \rangle \mathcal{G}[\![\text{Select } xs]\!] && (by \text{ defn})
\end{aligned}$$

Thus, generators can be coherently factored into a parser and a distribution. \square

C Proof of Theorem 4.2

Theorem 4.2. For all *simplified free generators* g and choices c ,

$$\delta_c \mathcal{L} \llbracket g \rrbracket = \mathcal{L} \llbracket \delta_c g \rrbracket.$$

Proof. We again proceed by induction on the structure of g .

Case $g = \text{Void}$. $\emptyset = \emptyset$.

Case $g = \text{Pure } a$. $\emptyset = \emptyset$.

Case $g = \text{Pair } x \ y$.

$$\begin{aligned} \delta_c \mathcal{L} \llbracket \text{Pair } x \ y \rrbracket &= \delta_c (\mathcal{L} \llbracket x \rrbracket \cdot \mathcal{L} \llbracket y \rrbracket) && \text{(by defn)} \\ &= \delta_c (\mathcal{L} \llbracket x \rrbracket) \cdot \mathcal{L} \llbracket y \rrbracket + \nu \mathcal{L} \llbracket x \rrbracket \cdot \delta_c \mathcal{L} \llbracket y \rrbracket && \text{(by defn)} \\ &= \delta_c (\mathcal{L} \llbracket x \rrbracket) \cdot \mathcal{L} \llbracket y \rrbracket && \text{(by Lemma 4.1)} \\ &= \mathcal{L} \llbracket \delta_c x \rrbracket \cdot \mathcal{L} \llbracket y \rrbracket && \text{(by IH)} \\ &= \mathcal{L} \llbracket \text{Pair } (\delta_c x) \ y \rrbracket && \text{(by defn)} \\ &= \mathcal{L} \llbracket \delta_c \text{Pair } x \ y \rrbracket && \text{(by defn)} \end{aligned}$$

Case $g = \text{Map } f \ x$.

$$\begin{aligned} \delta_c \mathcal{L} \llbracket \text{Map } f \ x \rrbracket &= \delta_c \mathcal{L} \llbracket x \rrbracket && \text{(by defn)} \\ &= \mathcal{L} \llbracket \delta_c x \rrbracket && \text{(by IH)} \\ &= \mathcal{L} \llbracket \text{Map } f \ (\delta_c x) \rrbracket && \text{(by defn)} \\ &= \mathcal{L} \llbracket \delta_c (\text{Map } f \ x) \rrbracket && \text{(by defn*)} \end{aligned}$$

*Note that the last step follows because $\text{Map } f \ x$ is assumed to be simplified, so $x \neq \text{Pure } a$. This means that $f \ \langle \$ \rangle \ x$ is equivalent to $\text{Map } f \ x$.

Case $g = \text{Select } xs$. If there is no pair $(c, \ x)$ in xs , then $\emptyset = \emptyset$. Otherwise,

$$\begin{aligned} \delta_c \mathcal{L} \llbracket \text{Select } xs \rrbracket &= \delta_c \{c \cdot s \mid s \in \mathcal{L} \llbracket x \rrbracket\} && \text{(by defn)} \\ &= \mathcal{L} \llbracket x \rrbracket && \text{(by defn)} \\ &= \mathcal{L} \llbracket \delta_c (\text{Select } xs) \rrbracket && \text{(by defn)} \end{aligned}$$

Thus we have shown that the symbolic derivative of free generators is compatible with the derivative of the generator's language. \square

There is another proof of this theorem, suggested by Alexandra Silva, which uses the fact that 2^{Σ^*} is the final coalgebra, along with the observation that FGen has a $2 \times (-)^{\Sigma}$ coalgebraic structure. This approach is certainly more elegant, but it abstracts away some helpful operational intuition.

D Proof of Theorem 4.3

Theorem 4.3. For all *simplified free generators* g , choice sequences s , and choices c ,

$$\mathcal{P}[\![\delta_c g]\!] s = \mathcal{P}[\![g]\!] (c \cdot s).$$

Proof. For simplicity, we prove the point-free version of this claim, i.e.:

$$\mathcal{P}[\![\delta_c g]\!] = \mathcal{P}[\![g]\!] \circ (c \cdot)$$

We proceed by induction on g .

Case $g = \text{Void}$. Nothing = Nothing.

Case $g = \text{Pure } a$. Nothing = Nothing.

Case $g = \text{Pair } x \ y$.

$$\begin{aligned} \mathcal{P}[\![\delta_c(\text{Pair } x \ y)]\!] &= \mathcal{P}[\![\delta_c x \otimes y]\!] && \text{(by defn)} \\ &= (\lambda x \ y \rightarrow (x, y)) \langle \$ \rangle \mathcal{P}[\![\delta_c x]\!] \langle * \rangle \mathcal{P}[\![y]\!] && \text{(by defn)} \\ &= ((\lambda x \ y \rightarrow (x, y)) \langle \$ \rangle \mathcal{P}[\![x]\!] \langle * \rangle \mathcal{P}[\![y]\!]) \circ (c \cdot) && \text{(by IH \& Lemma 4.1*)} \\ &= \mathcal{P}[\![\text{Pair } x \ y]\!] (c \cdot) && \text{(by defn)} \end{aligned}$$

*We can use Lemma 4.1 to show that x must consume at least one character. Thus, we can move the c in the derivative out into the final string, and trust that x will consume it.

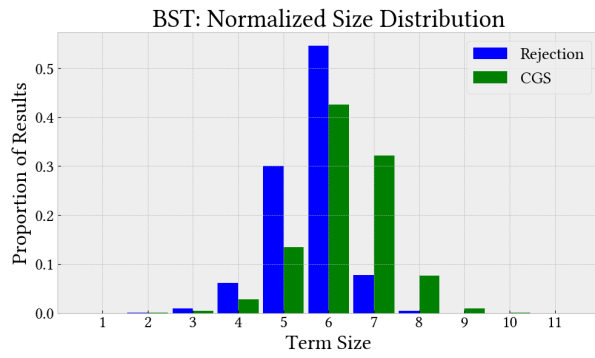
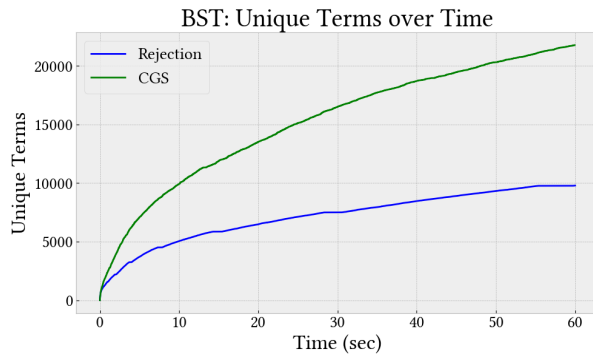
Case $g = \text{Map } f \ x$.

$$\begin{aligned} \mathcal{P}[\![\delta_c(\text{Map } f \ x)]\!] &= \mathcal{P}[\![f \langle \$ \rangle \delta_c x]\!] && \text{(by defn)} \\ &= f \langle \$ \rangle \mathcal{P}[\![\delta_c x]\!] && \text{(by app. properties)} \\ &= f \langle \$ \rangle (\mathcal{P}[\![x]\!] \circ (c \cdot)) && \text{(by IH)} \\ &= (f \langle \$ \rangle \mathcal{P}[\![x]\!]) \circ (c \cdot) && \text{(by app. properties)} \\ &= \mathcal{P}[\![\text{Map } f \ x]\!] \circ (c \cdot) && \text{(by defn)} \end{aligned}$$


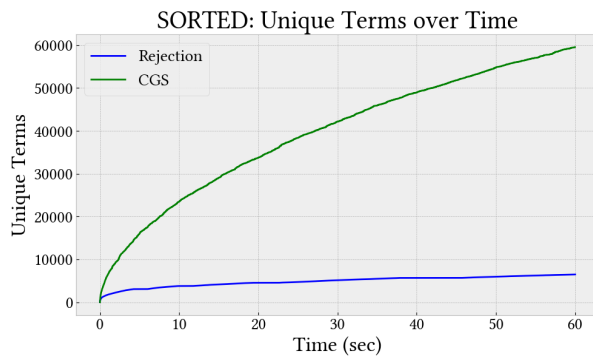
Case $g = \text{Select } xs$. Since both the derivative and the parser simply choose the branch of the Select corresponding to c , this case is trivial.

□

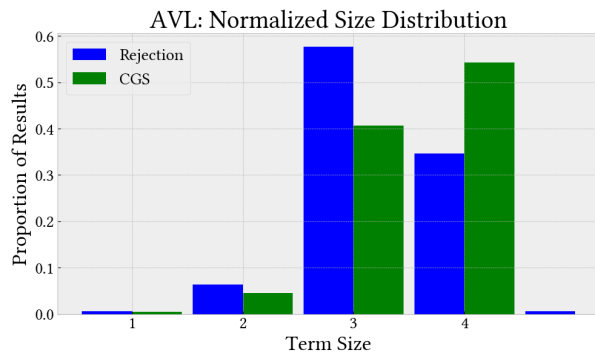

E Full Experimental Results



BST Charts



SORTED Charts



AVL Charts