

Parsing Randomness

HARRISON GOLDSTEIN, University of Pennsylvania, USA

BENJAMIN C. PIERCE, University of Pennsylvania, USA

Random data generators can be thought of as parsers of streams of randomness. This perspective on generators for random data structures is established folklore in the programming languages community, but it has never been formalized, nor have its consequences been deeply explored.

We build on the idea of *freer monads* to develop *free generators*, which unify parsing and generation using a common structure that makes the relationship between the two concepts precise. Free generators lead naturally to a proof that a monadic generator can be factored into a parser plus a distribution over choice sequences. Free generators also support a notion of *derivative*, analogous to the familiar Brzozowski derivatives of formal languages, allowing analysis tools to “preview” the effect of a particular generator choice. This gives rise to a novel algorithm for generating data structures satisfying user-specified preconditions.

1 INTRODUCTION

“A generator is a parser of randomness...” It’s one of those observations that’s totally puzzling right up to the moment it becomes totally obvious: a random generator—such as might be found in a property-based testing tool like QUICKCHECK [Claessen and Hughes 2000]—is a transformer from a series of random choices into a data structure, just as a parser is a transformer from a series of characters into a data structure.

While this connection may be obvious once it is pointed out, few actually think of generators this way. Indeed, to our knowledge the framing of random generators as parsers has never been explored formally. The relationship between these fundamental concepts deserves a deeper look!

We focus on generators written in the *monadic* style popularized by the QUICKCHECK library, which that build random data structures by making a sequence of random choices; those choices are the key. Traditionally, a generator makes decisions using a stored source of randomness (e.g., a seed) that it consults and updates whenever it must make a choice. Equivalently, if we like, we can pre-compute a list of choices and pass it in to the generator, which gradually walks down the list whenever it needs to make random decisions. In this mode of operation, the generator is effectively parsing the sequence of choices into a data structure!

To connect generators and parsers, we introduce *free generators*, syntactic data structures that can be interpreted as *either* generators or parsers. Free generators have a rich theory; in particular,

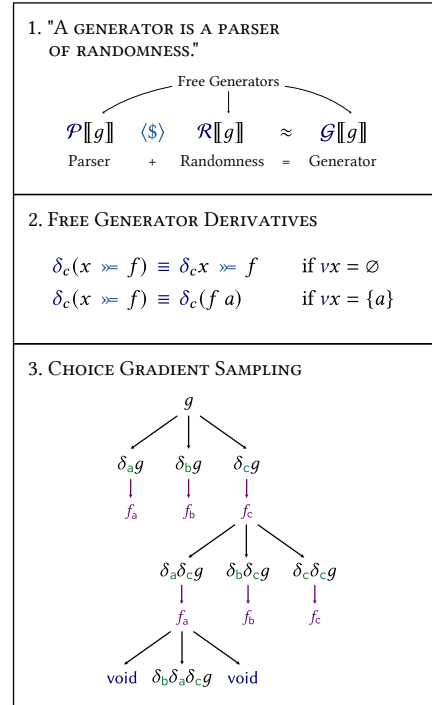


Fig. 1. Our contributions.

we can use them to prove that a large class of random generators can be factored into a parser and a distribution over sequences of choices.

Besides clarifying folklore, free generators admit transformations that do not exist for standard generators and parsers. A particularly exciting one is a notion of *derivative* which modifies a generator by asking the question: “what does this generator look like after it makes choice c ?” The derivative previews a particular choice to determine how likely it is to lead to useful values.

We use derivatives of free generators to tackle a well-known problem—we call it the *valid generation problem*. The challenge is to generate a large number of random values that satisfy some validity condition. This problem comes up often in property-based testing, where the validity condition is the precondition of some functional specification. Since generator derivatives give a way of previewing the effects of a particular choice, we can use *gradients* (derivatives with respect to a vector of choices) to preview all possible choices and pick a promising one. This leads us to an elegant algorithm that takes a free generator and replaces its distribution with one that produces only valid values. Replacing the distribution in this way trades the benefits of the programmer’s tuning effort for a higher chance of finding valid inputs to test with.

In §2 below, we introduce the ideas behind free generators and the operations that can be defined on them. We then present our main contributions:

- We formalize the folklore analogy between parsers and generators using *free generators*, a novel class of structures that make choices explicit and support syntactic transformations (§3). We use free generators to prove that any finitely supported *monadic generator* can be factored into a parser and a distribution over strings.
- We exploit free generators to transport an idea from formal languages—the *Brzowski derivative*—to the context of generators (§4).
- To illustrate the potential applications of these formal results, we present an algorithm that uses derivatives to turn a naïve generator into one with a different distribution, assigning nonzero probability only to values satisfying a Boolean precondition (§5). Our algorithm performs well on a set of simple benchmarks, in most cases producing more than twice as many valid values as a naïve “rejection sampling” generator in the same amount of time (§6).

We conclude with related and future work (§8 and §9).

2 HIGH-LEVEL STORY

To set the stage, let’s clarify the specific formulations of generators and parsers that we plan to discuss. Consider the following programs:

<pre> genTree h = if h = 0 then return Leaf else c ← frequency [(1, False), (3, True)] if c == False then return Leaf if c == True then x ← genInt () l ← genTree (h - 1) r ← genTree (h - 1) return Node l x r </pre>	<pre> parseTree h = if h = 0 then return Leaf else c ← consume () if c == 1 then return Leaf if c == n then x ← parseInt () l ← parseTree (h - 1) r ← parseTree (h - 1) return Node l x r else fail </pre>
--	--

The program on the left, `genTree`, generates random binary trees of integers like

Node Leaf 5 Leaf and Node Leaf 5 (Node Leaf 8 Leaf),

up to a given height h , guided by a series of weighted random Boolean choices made using frequency. Each time the program runs, it produces a random tree—i.e., the program denotes a distribution over trees. Generators like these can describe arbitrary finitely supported distributions of values.

The program on the right, `parseTree`, parses a string into a tree, turning

`n5l1` into Node Leaf 5 Leaf and `n5ln8l1` into Node Leaf 5 (Node Leaf 8 Leaf).

It consumes the input string character by character with `consume` and uses the characters to decide what to do next. This program is deterministic, but its execution (and thus the final tree it produces) is guided by a string of characters it is passed as input. Parsers like these can parse arbitrary computable languages.

These two programs are nearly identical in structure, and both produce the same set of values. The main difference lies in how they make choices: in `genTree` branches are taken at random, whereas in `parseTree` they are controlled by the input string.

This is the key observation that links generators and parsers. To make it more concrete, let us imagine how to recover the distribution of `genTree h` from `parseTree h`. We can do this by choosing a string at random and then parsing it—if we choose strings with the correct distribution, then the result of parsing those strings into values will be the same as if we had run `genTree` in the first place.

Here, we want the distribution over strings given to `parseTree` to satisfy the weighting of the Boolean choices in `genTree`. That is, `n` should appear three times more often than `l`, since `True` is chosen three times more often than `False`.

Free Generators. With these intuitions in hand, let's connect parsing and generation formally. First, we unify random generation with parsing by abstracting both into a single data structure; then we show that a structure of this form can be viewed equivalently as a generator or as a parser and a source of randomness.

Our unifying data structure is called a **free generator**.¹ Free generators are syntactic structures that can be interpreted as programs that either generate or parse. For example:

```
fgenTree h =
  if h = 0 then
    return Leaf
  else
    c ← pick [(1, l, return False), (3, n, return True)]
    if c == False then return Leaf
    if c == True then
      x ← fgenInt ()
      l ← fgenTree (h - 1)
      r ← fgenTree (h - 1)
      return Node l x r
```

The structure of this program is again very similar to that of `genTree` and `parseTree`. The call to `pick` on line 5 combines ideas from both the generator (capturing the relative weights of `False` and `True`) and the parser (capturing the labels `l` and `n` corresponding to different paths in the parser code). However, the meaning of `fgenTree` is very different from that of either `genTree` or `parseTree`.

¹This document uses the knowledge package in \LaTeX to make definitions interactive. Readers viewing the PDF electronically can click on technical terms and symbols to see where they are defined in the document.

The operators in `fgenTree` are entirely syntactic, and the result of running `fgenTree h` is simply an abstract syntax tree (AST).

The syntactic nature of free generators means that they can simultaneously represent generators, parsers, and more. In §3 we give several ways to interpret free generators. We write $\mathcal{G}[\![\cdot]\!]$ for the **random generator interpretation** of a free generator and $\mathcal{P}[\![\cdot]\!]$ for the **parser interpretation**. In other words,

$$\mathcal{G}[\![\text{fgenTree } h]\!] \approx \text{genTree } h \quad \text{and} \quad \mathcal{P}[\![\text{fgenTree } h]\!] \approx \text{parseTree } h.$$

The interpretation functions walk the AST produced by `fgenTree` to recover the behavior of the generator and parser programs.

These two interpretations can be related, formally, with the help of one final interpretation function, $\mathcal{R}[\![\cdot]\!]$, the **randomness interpretation** of the free generator. The randomness interpretation produces the distribution of sequences of choices that the random generator interpretation makes. Now, for any free generator g , we have

$$\mathcal{P}[\![g]\!] \langle \$ \rangle \mathcal{R}[\![g]\!] \approx \mathcal{G}[\![g]\!]$$

where $\langle \$ \rangle$ is a “mapping” operation that applies a function to samples from a distribution (see Theorem 3.1 below). Since a large class of generators (monadic generators with a finitely supported distribution) can also be written as free generators, another way to read this theorem is that such generators can be factored into two pieces: a distribution over choice sequences (given by $\mathcal{R}[\![\cdot]\!]$), and a parser of those sequences (given by $\mathcal{P}[\![\cdot]\!]$).

This precisely formalizes the intuition that “A generator is a parser of randomness.” But wait, there’s more to come!

Derivatives of Free Generators. Since a free generator defines a parser, it also defines a formal language: we write $\mathcal{L}[\![\cdot]\!]$ for this **language interpretation** of a free generator. The language of a free generator is the set of choice sequences that it can parse.

Viewing free generators this way suggests some interesting ways that free generators might be manipulated. In particular, formal languages come with a notion of *derivative*, due to Brzozowski [Brzozowski 1964]. Given a language L , the Brzozowski derivative of L with respect to a character c is

$$\delta_c^{\mathcal{L}} L = \{s \mid c \cdot s \in L\},$$

that is, the set of all strings in L that start with c , with the first c removed.

We can apply the same intuition to parsers by considering the derivative of a parser with respect to c to be whatever parser remains after c has been parsed. Each consecutive derivative fixes certain choices within the parser, simplifying the program:

<pre> parseTree 10 = c ← consume() if c == 1 then return Leaf if c == n then x ← parseInt() l ← parseTree 9 r ← parseTree 9 return Node l x r else fail </pre>	$\delta_n^{\mathcal{L}}(\text{parseTree } 10) \approx$ <pre> x ← parseInt() l ← parseTree 9 r ← parseTree 9 return Node l x r </pre>	$\delta_5^{\mathcal{L}} \delta_n^{\mathcal{L}}(\text{parseTree } 10) \approx$ <pre> l ← parseTree 9 r ← parseTree 9 return Node l 5 r </pre>
--	--	--

The first derivative fixes the character `n`, ensuring that the parser will produce a `Node`. The next fixes the character `5`, which determines the value `5` in the final `Node`.

Free generators have a closely related notion of **derivative**, illustrated by an almost identical set of transformations:

<pre>fgenTree 10 = c ← pick [...] if c == False then return Leaf if c == True then x ← fgenInt() l ← fgenTree 9 r ← fgenTree 9 return Node l x r else fail</pre>	$\delta_n^{\mathcal{L}}(\text{fgenTree } 10) \approx$ <pre> x ← fgenInt() l ← fgenTree 9 r ← fgenTree 9 return Node l x r</pre>	$\delta_5^{\mathcal{L}} \delta_n^{\mathcal{L}}(\text{fgenTree } 10) \approx$ <pre> l ← fgenTree 9 r ← fgenTree 9 return Node l 5 r</pre>
--	---	--

But there is a critical difference between this series of derivatives and the ones for `parseTree`. Whereas the parser derivatives we saw could be thought of *intuitively* as a program transformation on parsers, the analogous transformation on free generators is readily computable! Just as we can compute the derivative of a regular expression or a context-free grammar, we can compute the derivative of a free generator via a simple and efficient syntactic transformation.

In §4 we define a procedure, $\delta_c^{\mathcal{L}}$, for computing the derivative of a free generator and prove it correct, in the sense that, for all free generators g ,

$$\delta_c^{\mathcal{L}} \mathcal{L}[[g]] = \mathcal{L}[[\delta_c g]].$$

In other words, the derivative of the language of g is equal to the language of the derivative of g . (See Theorem 4.2.)

Putting Free Generators to Work. The derivative of a free generator is *the generator that remains after a particular choice*. This gives us a way of “previewing” the effect of making a choice by looking at the generator after fixing that choice.

In §5 and §6 we present and evaluate an algorithm called **Choice Gradient Sampling** that uses free generators to address the *valid generation problem*. Given a validity predicate on a data structure, the goal is to generate as many unique, valid structures as possible in a given amount of time. Starting from a simple free generator, our algorithm uses derivatives to evaluate choices and search for ones that produce valid values.

We evaluate the choice gradient sampling algorithm on four small benchmarks, all standard in the property-based testing literature. For each, we compare our algorithm to rejection sampling—sampling from a naïve generator and discarding invalid results—as a simple but useful baseline for understanding how well our algorithm performs. Our algorithm does remarkably well on three out of our four benchmarks, generating more than twice as many valid values as rejection sampling in the same period of time.

3 FREE GENERATORS

We now turn to developing the theory of free generators, beginning with some background on monadic abstractions for parsing and random generation.

Background: Monadic Parsers and Generators. In §2 we represented generators and parsers as pseudo-code. Here we flesh out the details, presenting all definitions as `HASKELL` programs, both

for the sake of concreteness and also because HASKELL’s abstraction features (e.g., type-classes) allow us to focus on the key concepts. HASKELL is a lazy functional language, but, as we focus our attention on finite programs, our results should apply directly to eager functional languages. It may also be possible, with appropriate domain knowledge, to translate these ideas to idiomatic constructs in popular imperative languages [Petříček 2009].

We represent both generators and parsers using *monads* [Moggi 1991]. A monad is a type constructor (e.g., `List`, `Maybe`, etc.) M equipped with two operations,

`return` $:: a \rightarrow M\ a$

and

$(\gg=) :: M\ a \rightarrow (a \rightarrow M\ b) \rightarrow M\ b$

(pronounced “bind”). Conceptually, `return` is the simplest way to put some value into the monad, while `bind` gives a way to sequence operations that produce monadic values.

We can use these operations to define `genTree` like we would in QUICKCHECK [Claessen and Hughes 2000] and `parseTree` like we would using combinators inspired by libraries like PARSEC [Leijen and Meijer 2001]:

<pre> genTree :: Int → Gen Tree genTree 0 = return Leaf genTree h = do c ← frequency [(1, False), (3, True)] case c of False → return Leaf True → do x ← genInt l ← genTree (h - 1) r ← genTree (h - 1) return (Node l x r) </pre>	<pre> parseTree :: Int → Parser Tree parseTree 0 = return Leaf parseTree h = do c ← consume case c of l → return Leaf n → do x ← parseInt l ← parseTree (h - 1) r ← parseTree (h - 1) return (Node l x r) _ → fail </pre>
---	---

In the first program, `genTree`, we use the monadic operations (along with `frequency`) to generate a random tree of integers. The expression `return Leaf` is a degenerate generator that always produces the value `Leaf`—this is what we mean by the “simplest way to put a value into the `Gen` monad.”

Rather than use $(\gg=)$ explicitly, we use `do`-notation, where

```

do
  a ← x
  f a

```

is syntactic sugar for $x \gg= f$. In the context of the `Gen` type, this operation samples from a generator x to get a value a and then passes it to f for further processing—this is what we mean by “sequencing operations.” Formally, `genTree` denotes a distribution over binary trees (e.g., an arrow in an appropriate category [Giry 1982]), and running the program samples from that distribution.

We can see these same combinators (used with a different monad) in `parseTree`. There, `return a` means “parse nothing and produce a ”, and $x \gg= f$ means “run the parser x to get a value a and then run the parser $f\ a$.” Under the hood, we have:

type `Parser a` = `String → Maybe (a, String)`

A Parser can be applied to a string to obtain either `Nothing` or `Just (a, s)`, where `a` is the parse result and `s` contains any extra characters. The `consume` function pulls the first character off of the string for inspection.

Expressiveness Relative to Other Abstractions. Monadic parsers and generators are maximally expressive in their respective domains. Monadic parsers can parse arbitrary computable languages, subsuming more restricted parser descriptions like context-free grammars and regular expressions. Likewise, monadic generators can generate values satisfying arbitrary computable constraints (e.g., it is possible to write a monadic generator for well-typed System F terms), subsuming less powerful representations like probabilistic context-free grammars.

For example, the following monadic generator generates (only) valid binary search trees:

```
genBST :: (Int, Int) → Gen Tree
genBST (lo, hi) | lo > hi = return Leaf
genBST (lo, hi) = do
  c ← frequency [(1, False), (3, True)]
  case c of
    False → return Leaf
    True → do
      x ← genRange (lo, hi)
      l ← genBST (lo, x - 1)
      r ← genBST (x + 1, hi)
      return (Node l x r)
```

The generator maintains the BST invariant by keeping track of the minimum and maximum values available for a given sub-tree and ensuring that all values to the left of a value are less and that all values to the right of a value are greater. This kind of generator is impossible to express as a stochastic CFG, since there is dependence between the choice of value `x` and the choices of sub-trees. Our examples are mostly focused on simple (non-dependent) generators to streamline the exposition, but our theory applies to the full class of monadic generators with finitely supported distributions.

Representing Free Generators. With the monad interface in mind, we can now give the formal definition of a free generator.²

Type Definition. The actual type of free generators is based on a structure called a *freer monad* [Kise-lyov and Ishii 2015]:

```
data Freer f a where
  Return :: a → Freer f a
  Bind   :: f a → (a → Freer f b) → Freer f b
```

This type looks complicated, but it is essentially just a representation of a monadic syntax tree. The constructors of `Freer` align almost exactly with the monadic operations `return` and `(>=)`, providing syntactic forms that can represent the building blocks of monadic programs.

An eagle-eyed reader might notice that the type of `Bind` here is not quite an instance of the type of `(>=)` above—one would have expected to see

²For algebraists: Free generators are “free” in the sense that they admit unique structure-preserving maps to other “generator-like” structures. In particular, the $\mathcal{G}[\![\cdot]\!]$ and $\mathcal{P}[\![\cdot]\!]$ maps are canonical. For the sake of space, we do not explore these ideas further here.


```
Bind :: Freer f a → (a → Freer f b) → Freer f b
```

with `Freer f a` as the first argument. The version we use is equally powerful, but more convenient. We will see in a moment that syntax trees in a freer monad are normalized by construction.

But what is going on with this `f` that appears throughout `Freer`? The type constructor `f` is a type of *specialized operations* that are specific to a particular monadic program. For example, programs in the `Gen` monad do not just use `return` and `(>=)`, they also use a `Gen`-specific operation, `frequency`. Similarly, representing a `Parser` as a syntax tree requires a way to represent a call to `consume`. In general, `f a` should be a syntactic representation of an operation returning `a`. Thus, we might have a type representing a parser operation that returns a character:

```
data Consume where
  Consume :: Consume Char
```

Since `Freer` is polymorphic over `f`, it can capture any specialized operation necessary to represent the syntax tree of a monad.

For free generators specifically, the specialized operation we need is called `pick`—we saw it in §2. Intuitively, `pick` subsumes both `frequency` and `consume`. We define the `Pick` operation with a data type (since free generators are syntactic objects) simultaneously with our definition of `FGen`, the type of *free generators*:

```
data Pick where
  Pick :: [(Weight, Choice, Freer Pick a)] → Pick a
```

```
type FGen a = Freer Pick a
```

By defining `FGen` as `Freer Pick`, we are really saying that “`FGen` is a monad with operation `Pick`.”

The `Pick` operation takes a list of triples. The first element of type `Weight` represents the weight given to a particular choice; weights are represented by signed integers for efficiency, but for theoretical purposes we treat them as strictly positive. The type `Choice` can theoretically be any type that admits equality, but for the purposes of this paper we take choices to be single characters. This makes the analogy with parsing clearer. Finally, `Freer Pick a` is actually just the type `FGen a`! Thus we should view the third element in the triple as a *nested* free generator that is run iff a specific choice is made.

Together the elements of these triples represent both kinds of choices that we have seen so far, subsuming both the weighted random choices of generators and the input-directed choices of parsers. Depending on our needs, we can interpret `Pick` as either kind of choice. In the rest of the paper, we sometimes speak of free generators “making” or “parsing” a choice, but remember that this is really just an analogy—a free generator is simply syntax, and the interpretation comes later.

Our First Free Generator. The `FGen` structure achieves our goal of unifying monadic generation and parsing, so let’s try writing a free generator. Following the basic structure of `genTree` and `parseTree`, we can start to define `fgenTree`:

```
fgenTree :: Int → FGen Tree
fgenTree 0 = Return Leaf
fgenTree h = Bind
  (Pick [(1, 'l', Return False), (3, 'n', Return True)])
  (λc → case c of
    False → Return Leaf
    True  → ... )
```


The first few lines are relatively easy to translate. The height checks are all the same as before, but now in the $h = 0$ case we produce the syntactic object `Return Leaf` rather than `return Leaf`, whose behavior depends on a particular implementation of `return`. When $h > 0$, we use `Bind` and `Pick` to specify that the generator has two choices: `False` (with weight 1, marked by character `l`) and `True` (with weight 3, marked by `n`).

But things get a bit more complicated when we get into the anonymous function passed as the second argument to `Bind`. In the `False` case we `Return Leaf` again, but in the `True` case the next step should be a call to `fgenInt`. We *could* look at the definitions of `genInt` and `parseInt` to determine the next choice, and then we *could* create a `Bind` node to make that choice, but that would be fairly tedious to do for every choice that the generator might eventually make. In general, while `FGen` is the right type to capture free generators, its constructors are a bit cumbersome to write down directly.

Recovering Monadic Syntax. Luckily, we can use the same monadic machinery used by `genTree` and `parseTree` to make free generators much easier to write. We can define `return` and `(>=)` for `FGen` as follows, allowing us to use `do`-notation to write free generators:

```
return :: a → FGen a
return = Return
```

```
(>=) :: FGen a → (a → FGen b) → FGen b
Return a >= f = f a
Bind p g >= f = Bind p (λa → g a >= f)
```

The `return` operator maps directly to a `Return` syntax node, but there is a bit more going on in the definition of `(>=)`. Specifically, `(>=)` normalizes the structure of the computation, ensuring that there is always an operation at the “front.” The advantage of this is that it is always $O(1)$ to check if a free generator has a choice to make. There is no need to dig through the syntax tree to determine the next step.

Another convenient way to manipulate free generators is via an operation called “`fmap`,” written `f <$> x`. Like `return` and `(>=)`, `(<$>)` is a syntactic transformation, but intuitively `f <$> x` means “apply the function `f` to the result of generating/parsing with `x`.” We define it as:

```
(<$>) :: (a → b) → FGen a → FGen b
f <$> Return a = Return (f a)
g <$> (Bind p f) = Bind p ((g <$>) . f)
```

(Note that all monads have an analogous operation; this will come in handy later.)

Representing Failure. For reasons that will become clear in §4, it is useful to be able to represent a free generator that can “fail.” We call the always-failing free generator `void`, and define it like this:

```
void :: FGen a
void = Bind (Pick []) Return
```

Any reasonable interpretation of this free generator must fail (by either diverging or returning a signal value); with no choices in the `Pick` list, there is no way to get a value of type `a` to pass to the second argument of `Bind`. Additionally, the use of `Return` as the second argument to `Bind` is irrelevant, since any free generator with no choices available will fail. This suggests that we can check if a free generator is certainly `void` by matching on an empty list of choices! In `HASKELL` this is easy to do with a pattern synonym:

```

pattern Void :: FGen a
pattern Void ← Bind (Pick []) _

```

This declaration means that pattern-matching on `Void` is equivalent to matching a `Bind` with no choices to make and ignoring the second argument. It is simple to define a function that uses this new pattern to check if a particular free generator is `void`:

```

isVoid :: FGen a → Bool
isVoid Void = True
isVoid _    = False

```

While `void` is useful as an error case for algorithms that build free generators, it would be incorrect for a user to use `void` in a hand-written free generator. To enforce this constraint, we define a wrapper around `Pick` (called `pick`) that does a few coherence checks to make sure that the generator is constructed properly:

```

pick :: [(Weight, Choice, FGen a)] → FGen a
pick xs =
  case filter (λ (_, _, x) → not (isVoid x)) xs of
    ys | hasDuplicates (map snd ys) → undefined
    [] → undefined
    ys → Bind (Pick ys) Return

```

This function is partial: it yields `undefined` if the list passed to `pick` is invalid. (This is analogous to raising an exception in a conventional imperative language.) The first line filters out any choices that are equivalent to `void`, since making those choices would lead to failure. The second line checks that the user has not duplicated any of the choice labels; this would introduce a nondeterministic choice that would complicate the interpretation considerably (see §7). Finally, the third line ensures that the generator we construct is not itself `void`. In practice, these checks ensure that the various interpretations of free generators presented in the remainder of this section work as intended.

Examples. Now that we have seen the building blocks of free generators, let's look at a couple of concrete examples. First, we can finally write down an ergonomic version of `fgenTree`:

```

fgenTree :: Int → FGen Tree
fgenTree 0 = return Leaf
fgenTree h = do
  c ← pick [(1, l, return False), (3, n, return True)]
  case c of
    False → return Leaf
    True  → do
      x ← fgenInt
      l ← fgenTree (h - 1)
      r ← fgenTree (h - 1)
      return (Node l x r)

```

Remember, the `do`-notation here is no longer sequencing generators or parsers. Instead, each line of a `do`-block builds a new `Bind` node in a syntax tree. Similarly, `return` has no semantics, it only wraps a value in the inert `Return` constructor. In this way `fgenTree` looks like both `genTree` and `parseTree`, but it does not behave like either (yet).

Trees are nice as a running example, but they are by no means the most complicated thing that free generators can represent. Here is a free generator that produces random (possibly ill-typed) terms of a simply-typed lambda-calculus:

```
fgenExpr :: Int → FGen Expr
fgenExpr 0 = pick [ (1, i, Lit ⟨$⟩ fgenInt ), (1, v, Var ⟨$⟩ fgenVar) ]
fgenExpr h =
  pick [ (1, i, Lit ⟨$⟩ fgenInt ),
        (1, p, do { e1 ← fgenExpr (h - 1); e2 ← fgenExpr (h - 1); return (Plus e1 e2) }),
        (1, l, do { t ← fgenType; e ← fgenExpr (h - 1); return (Lam t e) }),
        (1, a, do { e1 ← fgenExpr (h - 1); e2 ← fgenExpr (h - 1); return (App e1 e2) }),
        (1, v, Var ⟨$⟩ fgenVar) ]
```

Structurally `fgenExpr` is similar to `fgenTree`; it just has more cases and more choices. One stylistic difference between `fgenExpr` and `fgenTree` is that `fgenExpr` does not `pick` a coin and use it to decide what should be generated next; instead, it picks among a list of free generators directly. These styles of writing free generators are equivalent.

This version of the lambda calculus uses de Bruijn indices for variables and has integers and functions as values. This is a useful example because, while syntactically valid terms in this language are easy to generate (as we just did), it is more difficult to generate only well-typed terms. We will return to this problem in §6.

Interpreting Free Generators. A free generator does not do anything on its own—it is just a data structure. To actually use these structures, we next define the interpretation functions that we mentioned in §2 and prove a theorem linking those interpretations together.

Free Generators as Generators of Values. The first and most natural way to interpret a free generator is as a `QUICKCHECK` generator—that is, as a distribution over data structures. Plain `QUICKCHECK` generators ignore failure cases like `void` (they throw an error if there are no valid choices to make), but to make things a bit more explicit for our theory we use a modified generator monad: Gen_\perp .³

We define the *random generator interpretation* of a free generator to be:

```
 $\mathcal{G}[\![\cdot]\!] :: \text{FGen } a \rightarrow \text{Gen}_\perp a$ 
 $\mathcal{G}[\![\text{Void}]\!] = \perp$ 
 $\mathcal{G}[\![\text{Return } v]\!] = \text{return } v$ 
 $\mathcal{G}[\![\text{Bind } (\text{Pick } xs) \text{ } f]\!] = \text{do}$ 
   $x \leftarrow \text{frequency } (\text{map } (\lambda (w, \_, x) \rightarrow (w, \text{return } x)) \text{ } xs)$ 
   $a \leftarrow \mathcal{G}[\![x]\!]$ 
 $\mathcal{G}[\![f \text{ } a]\!]$ 
```

Note that the operations on the right-hand side of this definition do *not* build a free generator; they are Gen_\perp operations. This translation turns the syntactic form `Return v` into the semantic action “always generate the value `v`” and the syntactic form `Bind` into an operation that chooses a random sub-generator (with appropriate weight), samples from it, and then continues with `f`.

Note that $\mathcal{G}[\![\text{fgenTree } h]\!]$ has the same distribution as `genTree h`.

Free Generators as Parsers of Random Sequences. The *parser interpretation* of a free generator views it as a parser of sequences of choices. The translation looks like this:

³Our HASKELL development implements partial generators with `Gen (Maybe a)`. We elide these details to streamline the presentation.

```

 $\mathcal{P}[\![\cdot]\!] :: \text{FGen } a \rightarrow \text{Parser } a$ 
 $\mathcal{P}[\![\text{Void}]\!] = \lambda s \rightarrow \text{Nothing}$ 
 $\mathcal{P}[\![\text{Return } a]\!] = \text{return } a$ 
 $\mathcal{P}[\![\text{Bind (Pick } xs) \text{ } f]\!] = \text{do}$ 
   $c \leftarrow \text{consume}$ 
   $x \leftarrow \text{case find } ((= c) \cdot \text{snd}) \text{ } xs \text{ of}$ 
     $\text{Just } (\_, \_, x) \rightarrow \text{return } x$ 
     $\text{Nothing} \rightarrow \text{fail}$ 
   $a \leftarrow \mathcal{P}[\![x]\!]$ 
   $\mathcal{P}[\![f \text{ } a]\!]$ 

```

This time the **do**-notation on the right hand side is interpreted using the Parser monad (as before, defined as $\text{String} \rightarrow \text{Maybe}(a, \text{String})$). In the case for Bind, the parser consumes a character and attempts to make the corresponding choice from the list provided by Pick. If it succeeds, it runs the corresponding sub-parser and continues with f . If it fails, the whole parser fails.

Note that $\mathcal{P}[\![\text{fgenTree } h]\!]$ has the same parsing behavior as $\text{parseTree } h$.

Free Generators as Generators of Random Sequences. Our final interpretation of free generators represents the distribution with which the generator makes choices, ignoring how those choices are used to produce values. In other words, it captures exactly the parts of the structure that the parser interpretation discards. We define the *randomness interpretation* of a free generator to be:

```

 $\mathcal{R}[\![\cdot]\!] :: \text{FGen } a \rightarrow \text{Gen}_{\perp} \text{String}$ 
 $\mathcal{R}[\![\text{Void}]\!] = \perp$ 
 $\mathcal{R}[\![\text{Return } a]\!] = \text{return } \varepsilon$ 
 $\mathcal{R}[\![\text{Bind (Pick } xs) \text{ } f]\!] = \text{do}$ 
   $(c, x) \leftarrow \text{frequency } (\text{map } (\lambda (w, c, x) \rightarrow (w, \text{return } (c, x))) \text{ } xs)$ 
   $s \leftarrow \mathcal{R}[\![x \gg f]\!]$ 
   $\text{return } (c : s)$ 

```

Again, we use Gen_{\perp} and frequency to capture randomness and potential failure.

Factoring Generators. These different interpretations of free generators are closely related to one another; in particular, we can reconstruct $\mathcal{G}[\![\cdot]\!]$ from $\mathcal{P}[\![\cdot]\!]$ and $\mathcal{R}[\![\cdot]\!]$. That is, a free generator's random generator interpretation can be factored into a distribution over choice sequences plus a parser of those sequences.

To make this more precise, we need a notion of equality for generators like the ones produced via $\mathcal{G}[\![\cdot]\!]$. We say two QUICKCHECK generators are *equivalent*, written $g_1 \equiv g_2$, iff the generators represent the same distribution over values. This is coarser notion than program equality, since two generators might produce the same distribution of values in different ways.

With this in mind, we can state and prove the relationship between different interpretations of free generators:

THEOREM 3.1 (FACTORING). *Every free generator can be factored into a parser and a distribution over choice sequences that are, together, equivalent to its interpretation as a generator. In other words, for all free generators g ,*

$$\mathcal{P}[\![g]\!] \langle \$ \rangle \mathcal{R}[\![g]\!] \equiv (\lambda x \rightarrow (x, \varepsilon)) \langle \$ \rangle \mathcal{G}[\![g]\!].$$

PROOF SKETCH. By induction on the structure of g ; see Appendix A for the full proof. \square

COROLLARY 3.2. *Any monadic generator, γ , written using `return`, `(\gg)`, and frequency, can be factored into a parser plus a distribution over choice sequences.*

PROOF. Translate γ into a free generator, g , by replacing `return` and `(\gg)` with the equivalent free generator constructs, and frequency with `pick`. (This will require choosing labels for each choice, but the specific choice of labels is irrelevant.)

By construction, $\gamma = \mathcal{G}[\![g]\!]$.

Additionally, g can be factored into a parser and a source of randomness via Theorem 3.1. Thus,

$$(\lambda x \rightarrow (x, \varepsilon)) \langle \$ \rangle \gamma = (\lambda x \rightarrow (x, \varepsilon)) \langle \$ \rangle \mathcal{G}[\![g]\!] \equiv \mathcal{P}[\![g]\!] \langle \$ \rangle \mathcal{R}[\![g]\!],$$

and γ can be factored as desired. \square

This corollary is what we wanted to show all along. Monadic generators are parsers of randomness.

Free Generators as Formal Language Syntax. One final interpretation will prove useful. The *language of a free generator* is the set of choice sequences that it can make or parse. It is defined recursively, by cases:

$$\begin{aligned} \mathcal{L}[\![\cdot]\!] &:: \text{FGen } a \rightarrow \text{Set } \text{String} \\ \mathcal{L}[\![\text{Void}]\!] &= \emptyset \\ \mathcal{L}[\![\text{Return } a]\!] &= \varepsilon \\ \mathcal{L}[\![\text{Bind } (\text{Pick } xs) \ f]\!] &= [\ c : s \mid (w, c, x) \leftarrow xs, \ s \leftarrow \mathcal{L}[\![x \gg f]\!] \] \end{aligned}$$

This definition uses HASKELL's list comprehension syntax to iterate through the large space of choices sequences in the language of a free generator. To determine the language of a Bind node, we look at each possible choice and then at each possible string in the language $\mathcal{L}[\![x \gg f]\!]$ obtained by continuing with that choice. (This recursion is well-founded as long as the language of the free generator is finite; by monad identities $\text{Bind } (\text{Pick } xs) \ f = \text{Bind } (\text{Pick } xs) \ \text{Return } \gg f$, and x is strictly smaller than $\text{Bind } (\text{Pick } xs) \ \text{Return}$.) For each of these strings, we attach the appropriate choice label to the front. The end result is a list of all of the sequences of choices that, if made in order, would result in a valid output.

We can think of the result of this interpretation as the support of the distribution given by $\mathcal{R}[\![g]\!]$. The language of a free generator is exactly those choice sequences that the random generator interpretation can make and the parser interpretation can parse.

4 DERIVATIVES OF FREE GENERATORS

Next, we review the notion of Brzowski derivative from formal language theory and show that a similar operation exists for *free generators*. The way these derivatives fall out from the structure of free generators highlights the advantages of taking the correspondence between generators and parsers seriously.

Background: Derivatives of Languages. The *Brzowski derivative* [Brzowski 1964] of a formal language L with respect to some choice c is defined as

$$\delta_c^{\mathcal{L}} L = \{s \mid c \cdot s \in L\}.$$
⁴

In other words, it is the set of strings in L that begin with c , with the initial c removed. For example,

$$\delta_a^{\mathcal{L}} \{\text{abc}, \text{aaa}, \text{bba}\} = \{\text{bc}, \text{aa}\}.$$

⁴The superscript \mathcal{L} highlights that is the *language* derivative, distinguishing it from the generator derivative to be defined momentarily.

Many formalisms for defining languages support syntactic transformations that correspond to Brzozowski derivatives. For example, we can take the derivative of a regular expression like this:

$$\begin{aligned}
 \delta_c^{\mathcal{L}} \emptyset &= \emptyset & \nu^{\mathcal{L}} \emptyset &= \emptyset \\
 \delta_c^{\mathcal{L}} \varepsilon &= \varepsilon & \nu^{\mathcal{L}} \varepsilon &= \varepsilon \\
 \delta_c^{\mathcal{L}} \mathbf{c} &= \varepsilon \quad (c = \mathbf{c}) & \nu^{\mathcal{L}} \varepsilon &= \varepsilon \\
 \delta_c^{\mathcal{L}} \mathbf{d} &= \emptyset \quad (c \neq \mathbf{d}) & \nu^{\mathcal{L}} \mathbf{c} &= \emptyset \\
 \delta_c^{\mathcal{L}} (r_1 + r_2) &= \delta_c^{\mathcal{L}} r_1 + \delta_c^{\mathcal{L}} r_2 & \nu^{\mathcal{L}} (r_1 + r_2) &= \nu^{\mathcal{L}} r_1 + \nu^{\mathcal{L}} r_2 \\
 \delta_c^{\mathcal{L}} (r_1 \cdot r_2) &= \delta_c^{\mathcal{L}} r_1 \cdot r_2 + \nu^{\mathcal{L}} r_1 \cdot \delta_c^{\mathcal{L}} r_2 & \nu^{\mathcal{L}} (r_1 \cdot r_2) &= \nu^{\mathcal{L}} r_1 \cdot \nu^{\mathcal{L}} r_2 \\
 \delta_c^{\mathcal{L}} (r^*) &= \delta_c^{\mathcal{L}} r \cdot r^* & \nu^{\mathcal{L}} (r^*) &= \varepsilon
 \end{aligned}$$

The $\nu^{\mathcal{L}}$ operator, used in the “.” rule and defined on the right, determines the *nullability* of an expression—whether or not it accepts ε . If r accepts ε then $\nu^{\mathcal{L}} r = \varepsilon$, otherwise $\nu^{\mathcal{L}} r = \emptyset$.

As one would hope, if r has language L , it is always the case that $\delta_c^{\mathcal{L}} r$ has language $\delta_c^{\mathcal{L}} L$.

The Free Generator Derivative. To define derivatives of free generators, we first need a definition of *nullability* for free generators:

$$\begin{aligned}
 \nu &:: \text{FGen } a \rightarrow \text{Set } a \\
 \nu(\text{Return } v) &= \{v\} \\
 \nu g &= \emptyset \quad (g \neq \text{Return } v)
 \end{aligned}$$

Note that this behaves a bit differently than the $\nu^{\mathcal{L}}$ operation on regular expressions. For a regular expression r , the expression $\nu^{\mathcal{L}} r$ is either \emptyset or ε . Here, the null check returns either \emptyset or the singleton set containing the value in the Return node. That is, ν for free generators extracts a value that can be obtained by making no further choices. Another difference is that, for free generators, “can accept the empty string” and “accepts only the empty string” are equivalent statements; this greatly simplifies the definition of ν .

To see what the derivative operation might look like, we can write down some equations that it should satisfy, based on the equations satisfied by regular expressions:

$$\begin{aligned}
 \delta_c \text{void} &\equiv \text{void} & (1) \\
 \delta_c(\text{return } v) &\equiv \text{void} & (2) \\
 \delta_c(\text{pick } xs) &\equiv x & \text{if } (c, x) \in xs & (3) \\
 \delta_c(\text{pick } xs) &\equiv \text{void} & \text{if } (c, x) \notin xs \\
 \delta_c(x \gg f) &\equiv \delta_c(f a) & \text{if } vx = \{a\} & (4) \\
 \delta_c(x \gg f) &\equiv \delta_c x \gg f & \text{if } vx = \emptyset
 \end{aligned}$$

The derivative of an empty generator, or of one that immediately returns a value without looking at any input, should be **void**. The derivative of **pick** depends on whether or not c is present in the list of possible choices—if it is, we simply make the choice; if not, the result is **void**. Finally, the equations for (\gg) are based on the equation for concatenation of regular expressions, using ν to check to see if the left hand side of the expression is out of choices to make.

Of course, these equations are not definitions. In fact, the actual definition of the *derivative* for a free generator g is much simpler:

```

 $\delta :: \text{Char} \rightarrow \text{FGen } a \rightarrow \text{FGen } a$ 
 $\delta_c(\text{Return } v) = \text{void}$ 
 $\delta_c(\text{Bind (Pick } xs) f) =$ 
  case find ((= c) . snd) xs of
    Just (_, _, x)  $\rightarrow x \gg f$ 
    Nothing  $\rightarrow \text{void}$ 
    
```

Since freer monads are pre-normalized, there is no need to check nullability explicitly in this definition. It is always apparent from the top-level constructor (Return or Bind) whether or not there is a choice available to be made. The definition is not even recursive!

We can use the earlier equations to give us confidence that this definition is correct.

LEMMA 4.1. δ_c satisfies equations (1), (2), (3), and (4). In other words, the free generator derivative behaves similarly to the regular expression derivative.

PROOF SKETCH. See Appendix B for the proofs. Most are immediate. \square

Another way to ensure that the derivative operation acts as expected is to see how it behaves in relation to the free generator's [language interpretation](#). The following theorem makes this concrete:

THEOREM 4.2. The derivative of a free generator's language is the same as the language of its derivative. That is, for all [free generators](#) g and choices c ,

$$\delta_c^{\mathcal{L}} \llbracket g \rrbracket = \llbracket \delta_c g \rrbracket.$$

PROOF SKETCH. Straightforward induction (see Appendix C). \square

Since derivatives behave as expected, we can use them to simulate the behavior of a free generator. Just as we can check if a regular expression matches a string by taking derivatives with respect to each character in the string, we can simulate a free generator's parser interpretation by taking repeated derivatives. Each derivative fixes a particular choice, so a sequence of derivatives fixes a choice sequence.

5 GENERATING VALID RESULTS WITH GRADIENTS

We now put the theory of [free generators](#) and their [derivatives](#) into practice. We introduce Choice Gradient Sampling (CGS), a novel algorithm for generating data that satisfies some given validity condition, given a simple free generator for data of the appropriate type.

The [Choice Gradient Sampling](#) algorithm starts with a free generator for data of some type and uses derivatives to step the generator through choices one at a time. This process guides the generator towards values that are valid with respect to a given validity condition. At each step, the algorithm looks at all available choices and takes the free generator's derivative with respect to each one. Since this is, in a sense, a vector of all possible derivatives, we call this the *gradient* of the free generator, by analogy with calculus. We write

$$\nabla g = \langle \delta_a g, \delta_b g, \delta_c g \rangle$$

for the gradient of g with respect to the available choices $\{a, b, c\}$.

Since each derivative in the gradient is itself a free generator, the derivatives can be interpreted as value generators and sampled. If the derivative with respect to c produces lots of valid samples, then c is a good choice. If it produces mostly invalid samples, maybe other choices would be better. As we discuss below, this process is not faithful to the distribution of the original generator, but

it provides a metric that guides the algorithm toward a series of “good” choices, leading to more valid inputs in many cases.

```

1:  $g \leftarrow G$ 
2:  $\mathcal{V} \leftarrow \emptyset$ 
3: while true do
4:   if  $vg \neq \emptyset$  then return  $vg \cup \mathcal{V}$ 
5:   if isVoid  $g$  then  $g \leftarrow G$ 
6:    $C \leftarrow \text{choices } g$ 
7:    $\nabla g \leftarrow \langle \delta_c g \mid c \in C \rangle$  ▷  $\nabla g$  is the gradient of  $g$ 
8:   for  $\delta_c g \in \nabla g$  do
9:     if isVoid  $\delta_c g$  then
10:        $v \leftarrow \emptyset$ 
11:     else
12:        $x_1, \dots, x_N \leftarrow \mathcal{G}[\![\delta_c g]\!]$  ▷ Sample  $\mathcal{G}[\![\delta_c g]\!]$ 
13:        $v \leftarrow \{x_j \mid \varphi(x_j)\}$ 
14:        $f_c \leftarrow |v|$  ▷  $f_c$  is the fitness of  $c$ 
15:        $\mathcal{V} \leftarrow \mathcal{V} \cup v$ 
16:   if  $\max_{c \in C} f_c = 0$  then
17:     for  $c \in C$  do  $f_c \leftarrow \text{weightOf } c \text{ } G$ 
18:    $g \leftarrow \text{frequency } [(f_c, \delta_c g) \mid c \in C]$ 

```

Fig. 3. Choice Gradient Sampling: Given a free generator G , a sample rate constant N , and a validity predicate φ , this algorithm produces a set of outputs that all satisfy $\varphi(x)$.

We present the CGS algorithm in detail in Figure 3. Lines 7–14 are the core of the algorithm; their execution is shown pictorially in Figure 4. We take the gradient of g by taking the derivative with respect to each possible choice, in this case a , b , and c . Then we evaluate each of the derivatives by interpreting the free generator with $\mathcal{G}[\![\cdot]\!]$, sampling values from the resulting value generator, and counting how many of those results are valid with respect to φ . The precise number of samples is controlled by N , the sample rate constant; this is up to the user, but in general higher values for N will give better information about each derivative at the expense of time spent sampling. At the end of sampling, we have values f_a , f_b , and f_c , which we can think of as the “fitness” of each choice. We then pick a choice randomly, weighted based on fitness, and continue until our choices produce a valid output.

Critically, we avoid wasting effort by saving the samples (\mathcal{V}) that we use to evaluate the gradients. Many of those samples will be valid results that we can use, so there is no reason to throw them

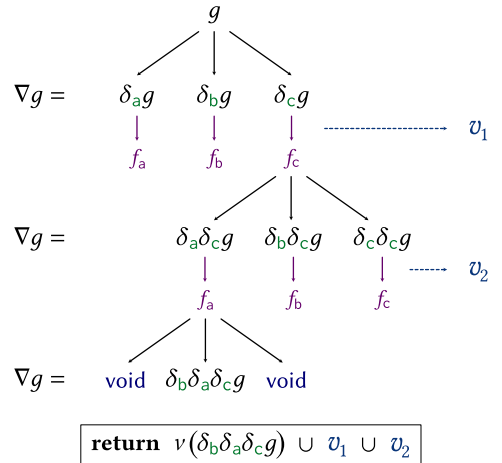


Fig. 4. The main loop of Choice Gradient Sampling.

away. Still, note that the performance of this sampling does depend on $|C|$, the number of choices available at this point. If the generator has many valid choices at a given point, it will need to do a lot of sampling to decide which choice to make.

This sampling procedure would not be possible with a traditional monadic generator: free generators are key. Trying to take a derivative of a traditional monadic generator would be like taking the derivative of a black-box function—there would be no generic way to incrementalize evaluation. Free generators expose more structure, making derivatives (and thus CGS) possible.

Impact on Distribution. As noted above, this algorithm is not faithful to the original distribution of G . In particular, the observable behavior of the algorithm is *not* to sample from the original generator’s distribution, conditioned on validity. While this property would arguably be ideal, it seems quite difficult to obtain. Moreover, its absence need not significantly detract from the value that CGS provides, for two reasons.

First, while the distribution produced by CGS is not faithful to the original distribution, it is certainly informed by it. At any given point in the algorithm, the weight given to a choice is based on how often making future choices, weighted by the original distribution, results in valid values. This means that valid values that are unlikely results from G will be unlikely results from CGS, and likely results from G will also be likely from CGS. Doing better than this would be quite difficult, since the preconditions we care about are black-box functions. This means that the only information they can provide is whether or not a particular value is valid, forcing us into rejection-based approaches. Standard rejection sampling does, in fact, sample from the ideal conditional distribution but it does so very slowly. Rather than sample from that distribution, CGS allows the predicate to guide its generation, reaching valid inputs more quickly.

Second, and more importantly, the primary use case for CGS is to improve the performance of free generators that are either automatically derived or else hand written but not carefully tuned. That is, the algorithm is most effective as a low-effort way to get from a useless generator to a usable one. If a tester has strict requirements for the distribution they are after, CGS will likely not be sufficient; but as a quick way of getting up and running it can be quite helpful.

We have implemented our Choice Gradient Sampling algorithm in HASKELL, along with all of the definitions presented throughout the paper. This code will be submitted for review as an artifact.

6 EXPLORATORY EVALUATION

The [Choice Gradient Sampling](#) algorithm is not a tightly optimized production algorithm: it is a proof of concept. Primarily, CGS exists to illustrate the theory of [free generators](#) and their [derivatives](#). Still, there is much to learn by exploring how well CGS is able to guide realistic generators to valid outputs.

We set out to answer two basic research questions:

RQ1 Does CGS produce more useful test inputs than standard sampling procedures, in the same period of time?

RQ2 Are the test inputs obtained from CGS well distributed in shape and size?

Our experimental results suggest that, with a few (interesting) caveats, these questions can both be answered in the affirmative. We find that CGS generally produces at least twice as many valid values as *rejection sampling* (explained in the next section) in the same period of time, and we also find that CGS’s values tend to be at least as diverse as the ones from rejection sampling. This indicates that using derivatives to guide generation is a promising approach to the valid generation problem.

Experimental Setup. Our experiments explore how well CGS improves on a canonical generation strategy. We compare our algorithm to the standard rejection sampling approach used by default in frameworks like `QUICKCHECK`, which takes a naïve generator, samples from it, and discards any results that are not valid. Rejection sampling is a useful point of comparison because, like our approach, it requires no extra effort from the user.

We use four simple free generators to test four different benchmarks: **BST**, **SORTED**, **AVL**, and **STLC**. Details about each of these benchmarks are given in Table 1.

	Free Generator	Validity Condition	N	Depth
BST	Binary trees with values 0–9	Is a valid BST	50	5
SORTED	Lists with values 0–9	Is sorted	50	20
AVL	AVL trees with values 0–9	Is a balanced AVL tree	500	5
STLC	Arbitrary ASTs for λ -terms	Is well-typed	400	5

Table 1. Overview of benchmarks.

Each of our benchmarks requires a simple free generator to act as a baseline and as a starting point for CGS. For consistency, and to avoid potential biases, our generators follow their respective inductive data types as closely as possible. For example, `fgenTree`, shown in §3 and used in the **BST** benchmark, follows the structure of the definition of the `Tree` type exactly. All generators use uniform choice weights, to avoid potential biases introduced by manual tuning.

The parameter N , used by CGS to decide how many samples to use for each iteration, was chosen via trial and error in order to balance fitness accuracy with sampling time. It is possible that some of our best-case results might improve with a more careful choice of N .

Results. We ran CGS and Rejection on each benchmark for one minute (on a MacBook Pro with an M1 processor and 16GB RAM) and recorded the unique valid values produced. We counted unique values because duplicate tests are generally less useful than fresh ones (if the system under test is pure, duplicate tests add no value). The totals, averaged over 10 trials, are presented in Table 2.

	BST	SORTED	AVL	STLC
Rejection	7,354 (109)	5,768 (88)	129 (6)	70,127 (711)
CGS	22,107 (338)	59,677 (1,634)	219 (2)	280,091 (7,265)

Table 2. Unique valid values generated in 60 seconds ($n = 10$ trials). Standard deviation in parentheses.

These measurements show that CGS is always able to generate more unique values than Rejection in the same amount of time, often significantly more. The exception is the **AVL** benchmark; we discuss this below.

Besides unique values, we measured some other metrics; the charts in Figure 5 show the results for the **STLC** benchmark. The first plot (“Unique Terms over Time”) shows how CGS behaves over time. Not only does CGS find more unique terms than Rejection overall, but its lead continues to grow over time. Additionally, the “Normalized Size Distribution” chart shows the size distributions terms generated by both algorithms. The CGS distribution is skewed farther to the right, showing that it generates larger terms on average; this is good from the perspective of property-based testing, where test size is often positively correlated with bug-finding power, since larger test inputs tend to exercise more of the implementation code. Analogous charts for the remaining benchmarks can be found in Appendix D.

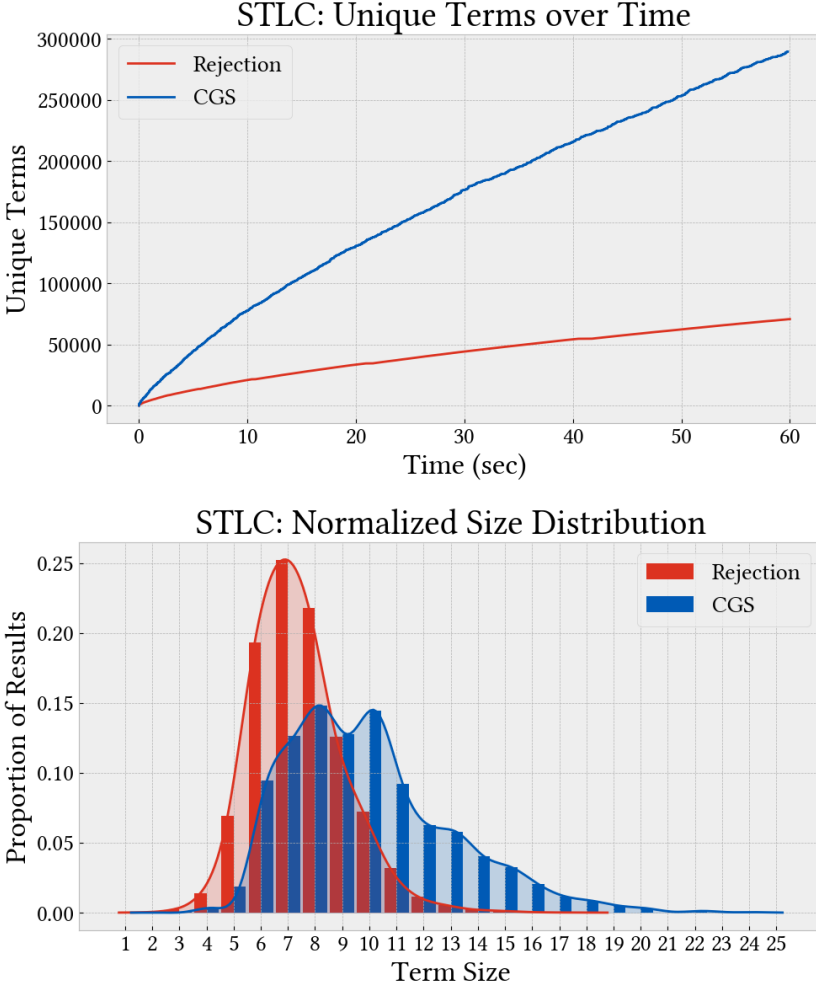


Fig. 5. Unique values and term sizes for the **STLC** benchmark, averaged over values in a single trial.

Measuring Diversity. Nothing in the CGS algorithm guarantees that the values we generate are *diverse*. Test input diversity is critical for effective testing, since a more diverse test suite will find more bugs more quickly, so we present experimental evidence that the values produced by CGS are indeed no less diverse than the valid values produced by rejection sampling.

Our diversity metric relies on the fact that each value is roughly isomorphic to the choice sequence that generated it. For example, in the case of **BST**, the sequence `n51611` can be parsed to produce Node 5 Leaf (Node 6 Leaf Leaf) and a simple in-order traversal can recover `n51611` again. Thus, choice sequence diversity is a reasonable proxy for value diversity.

We estimated the average Levenshtein distance [Levenshtein et al. 1966] (the number of edits needed to turn one string into another) between pairs of choice sequences in the values generated by each of our algorithms. We chose this metric for sequence distance because it is fairly standard and implementations were readily available. Computing an exact mean distance between all pairs in such a large set would be very expensive, so we settled for the mean of a random sample of 3000

pairs from each set of valid values. Figure 6 shows the results of these distance calculations, broken down by value size.




Fig. 6. Levenshtein diversity of generated values, plotted against the size of those values.

Each pair of lines in the chart represents an experiment. For all but **AVL** (the small pair of dash-dotted lines in the lower left), the lines exhibit a clear trend: the per-size diversity of CGS is at least as good as that of rejection sampling. (In fact, the diversity actually gets significantly better at large sizes, but much of this effect can be explained by the fact that CGS simply produces more large values.)

One might hope for even better results than this—why shouldn’t CGS produce much more diverse values at all sizes? A potential explanation lies in the way CGS retains intermediate samples. While the first few samples will be mostly uncorrelated, the samples drawn later on in the generation process (once a number of choices have been fixed) will tend to be similar to one another. This likely results in clusters of inputs that are all valid but that only explore one shape of input.

The Problem with AVL: Very Sparse Validity Conditions. The **AVL** benchmark is an outlier in most of our measurements: CGS only manages to find a modest number of extra valid AVL trees, and their pairwise diversity is actually slightly worse than that of rejection sampling. Understanding this phenomenon provides insight into a critical assumption underlying the CGS algorithm—namely that it is not too difficult to find valid values randomly.

It is clear that AVL trees *are* quite difficult to find randomly: balanced binary search trees are hard to generate on their own, and AVL trees are even more difficult because the generator must guess the correct height to cache at each node. This is why rejection sampling only finds 156 AVL trees in the time it takes to find 9,762 binary search trees.

In domains like this, CGS is unlikely to find *any* valid trees while sampling. In particular, the check in line 15 of Figure 4 will often be true, meaning that choices will be made at random rather than guided by the fitness of the appropriate derivatives. We could reduce this effect by significantly increasing the sample rate constant N , but then sampling time would likely dominate generation time, resulting in worse performance overall.

The lesson here seems to be that the CGS algorithm does not work well with especially hard-to-satisfy validity conditions. In §9, we present an idea that would do some of the hard work ahead of time and help with this issue.

7 LIMITATIONS

Our free generator abstraction is extremely general and demonstrably useful, but a few technical weaknesses are worth discussing.

The biggest limitation has to do with the kinds of distributions our free generators can represent. Our exposition uses weighted choices (frequency) as the randomness primitive, but `QUICKCHECK` is technically built using a primitive like:

```
choose :: Random r => (r, r) -> Gen r
```

Intuitively, `choose (x, y)` uniformly picks a value in the *range* from `x` to `y`, and this range can technically be infinite (e.g., if `r = Rational`). This cannot be replicated with frequency or `pick`. Thus, our results only apply to generators whose distributions are finitely supported.

Another small issue is that we have intentionally neglected one common element of monadic generators in the style of `QUICKCHECK`: `size`. Generators in standard `QUICKCHECK` track size bounds dynamically, allowing the testing framework to externally control the size distribution of the inputs that it generates. This does not impact our theoretical results (sizes can always be passed around manually, as we do in the examples in this paper), and sizes would be relatively easy to add to the free generator language in practice.

Finally, a note on the class of languages that free generators can parse (when interpreted with $\mathcal{P}[\cdot]$). Free generators are limited in their nondeterminism (by the definition of `pick`, and by assumptions made in the definition of $\mathcal{P}[\cdot]$); choices in a free generator are always unambiguous. This means that the parser interpretation of a free generator cannot parse arbitrary languages of choices, even though monadic parsers in general can parse arbitrary languages. Ultimately this is not a practical concern, as free generators parse sequences of choices, not realistic languages, but it is aesthetically disappointing. We believe it would be straightforward to add an operator for explicit nondeterminism and extend the interpretations accordingly.

8 RELATED WORK

We discuss a variety of publications that relate to the present work via connections either to free generators or to our Choice Gradient Sampling algorithm.

Parsing and Generation. The connection between parsers and generators has been employed implicitly in some generator implementations. Two popular property-based testing libraries, `HYPOTHESIS` [MacIver et al. 2019] and `CROWBAR` [Dolan and Preston 2017], implement generators by parsing a stream of random choices. In fact, `HYPOTHESIS` even takes advantage of parsing concepts when *shrinking* test inputs to make failing test-cases more readable for uses. However, neither of these frameworks has formalized the relationship between parsing and generation.

Free Generators. Garnock-Jones et al. present a formalism based on parsing expression grammars (PEGs) with some of the same goals as ours. They give a derivative-based algorithm that somewhat resembles CGS, which constructs sentences that match a particular PEG. Their work does not attempt to solve the valid generation problem for complex validity conditions like the ones we tackle, but it does provide further evidence that connecting parsing and generation is advantageous.

Claessen et al. [2015] present a generator representation that is related to our free generator structure, but used in a very different way. They primarily use the syntactic structure of their

generators (they call them “spaces”) to control the size distribution of generated outputs; in particular, spaces do not make choice information explicit in the way free generators do. [Claessen et al.](#)’s generation approach uses `HASKELL`’s laziness, rather than derivatives and sampling, to prune unhelpful paths in the generation process. This pruning procedure performs well when validity conditions are written to take advantage of laziness, but it is highly dependent on evaluation order and avoids generator choices that immediately lead to exclusively invalid results. In contrast, `CGS` does not require that predicates be written in a specific way and has a much more nuanced notion of “unhelpful” choices.

The Valid Generation Problem. Many other approaches to the valid generation problem have been explored.

The most obvious solution to the valid generation problem is to write a bespoke generator. For example, the `CSMITH` project famously developed a generator for valid C programs that was very successful at finding bugs in C compilers [[Yang et al. 2011](#)]. More generally, the domain-specific language for generators provided by the `QUICKCHECK` library [[Hughes 2007](#)] provides a whole framework for writing manual generators that produce valid inputs by construction. The primary issue with these manual approaches is effort: writing a bespoke generator is labor intensive and difficult. The `CGS` algorithm aims to avoid manual techniques like this in the hopes of making property-based testing more accessible to programmers that do not have the time or expertise to write their own custom generators.

The constraint logic programming (CLP) generators proposed by [Dewey \[2017\]](#) represent a different approach to valid generation, more automated than `QUICKCHECK`. Users of CLP generators have a constraint solver to help them, making it easier to express certain kinds of validity conditions in the generator. Even so, the CLP approach is not truly automatic: testers still need to express validity condition as annotated logic programs. Depending on the testers’ background, this may be ideal or it may be a deal-breaker. In contrast, `CGS` only requires that the validity condition be encoded as a Boolean predicate in the host programming language, which the tester may very well already have written for other reasons.

The `LUCK` language [[Lampropoulos et al. 2017a](#)] provides a similar semi-automatic solution; users are still required to put in some effort, but they are able to define generators and validity predicates at the same time. Again, this solution might be satisfying if users are starting from scratch and willing to learn a domain-specific language, but if validity predicates have already been written or users do not want to learn a new language, a more automated solution such as ours may be preferable.

When validity predicates are expressed as inductive relations, approaches like the one in *Generating Good Generators for Inductive Relations* [[Lampropoulos et al. 2017b](#)] are extremely powerful. In the `QUICKCHICK` framework, users can extract generators from the inductive relations that they likely already have for their proofs. This is incredibly convenient for testing lemmas that will eventually be proved, to establish confidence before attempting the proof. Unfortunately, the kinds of inductive relations that `QUICKCHICK` depends on generally require dependent types to express, so this approach does not work in most mainstream programming languages.

`TARGET` [[Löschner and Sagonas 2017](#)] uses search strategies like hill climbing and simulated annealing to supplement random generation and significantly streamline property-based testing. [Löschner and Sagonas](#)’s approach works well when inputs have a sensible notion of “utility,” but in the case of valid generation the utility is often degenerate—0 if the input is invalid, and 1 if it is valid—with no good way to say if an input is “better” or “worse.” In these cases, derivative-based searches may make more sense.

Some approaches use machine learning to automatically generate valid inputs. `LEARN&FUZZ` [Godefroid et al. 2017] generates valid data using a recurrent neural network. This solution seems to work best when a large corpus of inputs is already available and the validity condition is more structural than semantic. In the same vein, `RLCHECK` [Reddy et al. 2020] uses reinforcement learning to guide a generator to valid inputs. This approach served as early inspiration for our work, and we think that the theoretical advance of generator derivatives may lead improved learning algorithms in the future (see §9).

9 CONCLUSION

Free generators and their derivatives are powerful structures that give a flexible perspective on random generation. This formalism yields a useful algorithm for addressing the valid generation problem, and it clarifies the folklore that a generator is a parser of randomness. Moving forward, there are a number of paths to explore, some continuing our theoretical exploration and others looking towards algorithmic improvements.

Bidirectional Free Generators. We have only scratched the surface of what seems possible with `free generators`. One concrete next step is to merge the theory of free generators with the emerging theory of *ungenerators* [Goldstein 2021]. This work describes generators that can be run both forward (to generate values as usual) and *backward*. In the backward direction, the program takes a value that the generator might have generated and “un-generates” it to give a sequence of choices that the generator might have made when generating that value.

Free generators are quite compatible with these ideas, and turning a free generator into a bidirectional generator that can both generate and ungenerate should be fairly straightforward. From there, we can build on the ideas in the ungenerators work and use the backward direction of the generator to learn a distribution of choices that approximates some user-provided samples of “desirable” values.

Algorithmic Optimizations. In §6, we saw some problems with the `Choice Gradient Sampling` algorithm: because CGS evaluates derivatives via sampling, it does poorly when validity conditions are very difficult to satisfy. This begs the question: might it be possible to evaluate the fitness of a derivative without naïvely sampling?

One potential approach involves staging the sampling process. Given a free generator with a depth parameter, we can first evaluate choices on generators for size 1, then evaluate choices for size 2, etc. These intermediate stages would make gradient sampling more successful at larger sizes, and might significantly improve the results on benchmarks like `AVL`. Unfortunately, this approach might perform poorly on benchmarks like `STLC` where the validity condition is not uniform: size-1 generators would avoid generating variables, leading larger generators to avoid variables as well. Nevertheless, this design space seems well worth exploring.

Making Choices with Neural Networks. Another algorithmic optimization is a bit farther afield: using recurrent neural networks (RNNs) to improve our generation procedure.

As `Choice Gradient Sampling` makes choices, it generates useful data about the frequencies with which choices should be made. Specifically, every iteration of the algorithm produces a pair of a history and a distribution over next choices that looks something like this:

$$\text{abcca} \mapsto \{a : 0.3, b : 0.7, c : 0.0\}$$

In the course of CGS, this information is used once (to make the next choice) and then forgotten—but what if there was a way to learn from it? Pairs like this could be used to train an RNN to make choices that are similar to the ones made by CGS.

There are details to work out, including network architecture, hyper-parameters, etc., but in theory we could run CGS for a while, train an RNN, and after that point only use the RNN to generate valid data. Setting things up this way would recover some of the time that is currently spent sampling of derivative generators.

One could imagine a user writing a definition of a type and a predicate for that type, and then setting the model to train while they work on their algorithm. By the time the algorithm is finished and ready to test, the RNN model would be trained and ready to produce valid test inputs. A workflow like this might help increase adoption of property-based testing in industry.

REFERENCES

- Janusz A Brzozowski. 1964. Derivatives of regular expressions. *Journal of the ACM (JACM)* 11, 4 (1964), 481–494.
- Koen Claessen, Jonas Duregård, and Michal H. Palka. 2015. Generating constrained random data with uniform distribution. *J. Funct. Program.* 25 (2015). <https://doi.org/10.1017/S0956796815000143>
- Koen Claessen and John Hughes. 2000. QuickCheck: a lightweight tool for random testing of Haskell programs. In *Proceedings of the Fifth ACM SIGPLAN International Conference on Functional Programming (ICFP '00)*, Montreal, Canada, September 18–21, 2000, Martin Odersky and Philip Wadler (Eds.). ACM, Montreal, Canada, 268–279. <https://doi.org/10.1145/351240.351266>
- Kyle Thomas Dewey. 2017. *Automated Black Box Generation of Structured Inputs for Use in Software Testing*. University of California, Santa Barbara.
- Stephen Dolan and Mindy Preston. 2017. Testing with crowbar. In *OCaml Workshop*.
- Tony Garnock-Jones, Mahdi Eslamimehr, and Alessandro Warth. 2018. Recognising and generating terms using derivatives of parsing expression grammars. *arXiv preprint arXiv:1801.10490* (2018).
- Michele Giry. 1982. A categorical approach to probability theory. In *Categorical aspects of topology and analysis*. Springer, 68–85.
- Patrice Godefroid, Hila Peleg, and Rishabh Singh. 2017. Learn&fuzz: Machine learning for input fuzzing. In *2017 32nd IEEE/ACM International Conference on Automated Software Engineering (ASE)*. IEEE, 50–59.
- Harrison Goldstein. 2021. Ungenerators. In *ICFP Student Research Competition*. <https://harrisongoldste.in/papers/icfpsrc21.pdf>
- John Hughes. 2007. QuickCheck testing for fun and profit. In *International Symposium on Practical Aspects of Declarative Languages*. Springer, 1–32.
- Oleg Kiselyov and Hiromi Ishii. 2015. Freer monads, more extensible effects. *ACM SIGPLAN Notices* 50, 12 (2015), 94–105.
- Leonidas Lampropoulos, Diane Gallois-Wong, Catalin Hritcu, John Hughes, Benjamin C. Pierce, and Li-yao Xia. 2017a. Beginner’s Luck: a language for property-based generators. In *Proceedings of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages, POPL 2017, Paris, France, January 18–20, 2017*. 114–129. <http://dl.acm.org/citation.cfm?id=3009868>
- Leonidas Lampropoulos, Zoe Paraskevopoulou, and Benjamin C Pierce. 2017b. Generating good generators for inductive relations. *Proceedings of the ACM on Programming Languages* 2, POPL (2017), 1–30.
- Daan Leijen and Erik Meijer. 2001. Parsec: Direct style monadic parser combinators for the real world. (2001).
- Vladimir I Levenshtein et al. 1966. Binary codes capable of correcting deletions, insertions, and reversals. In *Soviet physics doklady*, Vol. 10. Soviet Union, 707–710.
- Andreas Löschner and Konstantinos Sagonas. 2017. Targeted Property-Based Testing. In *Proceedings of the 26th ACM SIGSOFT International Symposium on Software Testing and Analysis (Santa Barbara, CA, USA) (ISSTA 2017)*. Association for Computing Machinery, New York, NY, USA, 46–56. <https://doi.org/10.1145/3092703.3092711>
- David R MacIver, Zac Hatfield-Dodds, et al. 2019. Hypothesis: A new approach to property-based testing. *Journal of Open Source Software* 4, 43 (2019), 1891.
- Eugenio Moggi. 1991. Notions of computation and monads. *Information and computation* 93, 1 (1991), 55–92.
- Tomaš Petříček. 2009. Encoding monadic computations in C# using iterators. *Proceedings of ITAT* (2009).
- Sameer Reddy, Caroline Lemieux, Rohan Padhye, and Koushik Sen. 2020. Quickly generating diverse valid test inputs with reinforcement learning. In *ICSE ’20: 42nd International Conference on Software Engineering, Seoul, South Korea, 27 June - 19 July, 2020*, Gregg Rothermel and Doo-Hwan Bae (Eds.). ACM, 1410–1421. <https://doi.org/10.1145/3377811.3380399>
- Xuejun Yang, Yang Chen, Eric Eide, and John Regehr. 2011. Finding and understanding bugs in C compilers. In *Proceedings of the 32nd ACM SIGPLAN Conference on Programming Language Design and Implementation, PLDI 2011, San Jose, CA, USA, June 4–8, 2011*. 283–294. <https://doi.org/10.1145/1993498.1993532>

Appendix

A PROOF OF THEOREM 3.1

LEMMA A.1.

$$\mathcal{P}[\![x \gg f]\!] \langle \$ \rangle \mathcal{R}[\![x \gg f]\!] \equiv (\mathcal{P}[\![x]\!] \langle \$ \rangle \mathcal{R}[\![x]\!]) \gg \lambda(a, []) \rightarrow (\mathcal{P}[\![fa]\!] \langle \$ \rangle \mathcal{R}[\![fa]\!])$$

PROOF. By induction on the structure of x .

- Case $x = \text{Return } a$

$$\begin{aligned} & \mathcal{P}[\![\text{Return } a \gg f]\!] \langle \$ \rangle \mathcal{R}[\![\text{Return } a \gg f]\!] \\ & \quad \text{-- By definition } (\gg). \\ & \equiv \mathcal{P}[\![f a]\!] \langle \$ \rangle \mathcal{R}[\![f a]\!] \\ & \quad \text{-- By } \eta\text{-expansion and definitions of } \mathcal{P}[\![\cdot]\!] \text{ and } \mathcal{R}[\![\cdot]\!]. \\ & \equiv \mathcal{P}[\![\text{Return } a]\!] \langle \$ \rangle \mathcal{R}[\![\text{Return } a]\!] \gg \lambda(a, []) \rightarrow \mathcal{P}[\![f a]\!] \langle \$ \rangle \mathcal{R}[\![f a]\!] \end{aligned}$$

- Case $x = \text{Bind (Pick } xs) \ k$

$$\begin{aligned} & \mathcal{P}[\![\text{Bind (Pick } xs) \ k \gg f]\!] \langle \$ \rangle \mathcal{R}[\![\text{Bind (Pick } xs) \ k \gg f]\!] \\ & \quad \text{-- By definition } (\gg). \\ & \equiv \mathcal{P}[\![\text{Bind (Pick } xs) (\lambda a \rightarrow k a \gg f)]\!] \langle \$ \rangle \mathcal{R}[\![\text{Bind (Pick } xs) (\lambda a \rightarrow k a \gg f)]\!] \\ & \quad \text{-- By definition } (\mathcal{P}[\![\cdot]\!] \text{ and } \mathcal{R}[\![\cdot]\!]). \\ & \equiv \text{do} \\ & \quad c \leftarrow \text{consume} \\ & \quad x \leftarrow \text{case find } ((= c) . \text{snd}) \ xs \text{ of} \\ & \quad \quad \text{Just } (_, _, x) \rightarrow \text{return } x \\ & \quad \quad \text{Nothing} \rightarrow \text{fail} \\ & \quad \mathcal{P}[\![x \gg \lambda a \rightarrow k a \gg f]\!] \langle \$ \rangle \text{do} \\ & \quad \quad (c, x) \leftarrow \text{frequency } (\text{map } (\lambda (w, c, y) \rightarrow (w, \text{return } (c, y)))) \ xs) \\ & \quad \quad s \leftarrow \mathcal{R}[\![x \gg (\lambda a \rightarrow k a \gg f)]\!] \\ & \quad \quad \text{pure } (c : s) \\ & \quad \text{-- By simplification .} \\ & \equiv \text{do} \\ & \quad (_, x) \leftarrow \text{frequency } (\text{map } (\lambda (w, c, y) \rightarrow (w, \text{return } (c, y)))) \ xs) \\ & \quad \mathcal{P}[\![x \gg \lambda a \rightarrow k a \gg f]\!] \langle \$ \rangle \mathcal{R}[\![x \gg \lambda a \rightarrow k a \gg f]\!] \\ & \quad \text{-- By monad laws.} \\ & \equiv \text{do} \\ & \quad (_, x) \leftarrow \text{frequency } (\text{map } (\lambda (w, c, y) \rightarrow (w, \text{return } (c, y)))) \ xs) \\ & \quad \mathcal{P}[\![(x \gg k) \gg f]\!] \langle \$ \rangle \mathcal{R}[\![(x \gg k) \gg f]\!] \\ & \quad \text{-- By IH.} \\ & \equiv \text{do} \\ & \quad (_, x) \leftarrow \text{frequency } (\text{map } (\lambda (w, c, y) \rightarrow (w, \text{return } (c, y)))) \ xs) \\ & \quad (\mathcal{P}[\![x \gg k]\!] \langle \$ \rangle \mathcal{R}[\![x \gg k]\!]) \gg (\lambda a \rightarrow \mathcal{P}[\![f a]\!] \langle \$ \rangle \mathcal{R}[\![f a]\!]) \\ & \quad \text{-- By expansion and definitions } (\mathcal{P}[\![\cdot]\!] \text{ and } \mathcal{R}[\![\cdot]\!]). \\ & \equiv (\mathcal{P}[\![\text{Bind (Pick } xs) \ k]\!] \langle \$ \rangle \mathcal{R}[\![\text{Bind (Pick } xs) \ k]\!]) \gg (\lambda a \rightarrow \mathcal{P}[\![f a]\!] \langle \$ \rangle \mathcal{R}[\![f a]\!]) \end{aligned}$$

□

THEOREM 3.1 (FACTORING). *Every **free generator** can be factored into a parser and a distribution over choice sequences that are, together, equivalent to its interpretation as a generator. In other words, for all free generators g ,*

$$\mathcal{P}[\![g]\!] \langle \$ \rangle \mathcal{R}[\![g]\!] \equiv (\lambda x \rightarrow (x, \varepsilon)) \langle \$ \rangle \mathcal{G}[\![g]\!].$$

PROOF. By induction on the structure of x .

- Case $x = \text{Return } a$

$$\mathcal{P}[\![\text{Return } a]\!] \langle \$ \rangle \mathcal{R}[\![\text{Return } a]\!] \equiv \text{return } (a, []) \equiv (\lambda a \rightarrow (a, [])) \langle \$ \rangle \mathcal{G}[\![\text{Return } a]\!]$$

(By definition.)

- Case $x = \text{Bind } (\text{Pick } xs) \ k$

$$\mathcal{P}[\![\text{Bind } (\text{Pick } xs) \ k]\!] \langle \$ \rangle \mathcal{R}[\![\text{Bind } (\text{Pick } xs) \ k]\!]$$

-- By definition ($\mathcal{P}[\![\cdot]\!]$ and $\mathcal{R}[\![\cdot]\!]$).

$$\begin{aligned} &\equiv (\text{do } c \leftarrow \text{consume} \\ &\quad x \leftarrow \text{case find } ((= c) . \text{snd}) \ xs \ \text{of} \\ &\quad \quad \text{Just } (_, _, x) \rightarrow \text{return } x \\ &\quad \quad \text{Nothing} \rightarrow \text{fail} \\ &\quad \mathcal{P}[\![x \gg k]\!] \langle \$ \rangle (\text{do} \\ &\quad \quad (c, x) \leftarrow \text{frequency } (\text{map } (\lambda (w, c, x) \rightarrow (w, \text{return } (c, x))) \ xs) \\ &\quad \quad s \leftarrow \mathcal{R}[\![x \gg k]\!] \\ &\quad \quad \text{pure } (c : s)) \end{aligned}$$

-- By simplification .

$$\begin{aligned} &\equiv \text{do } x \leftarrow \text{frequency } (\text{map } (\lambda (w, _, x) \rightarrow (w, \text{return } x)) \ xs) \\ &\quad s \leftarrow \mathcal{R}[\![x \gg k]\!] \\ &\quad \text{return } \mathcal{P}[\![x \gg k]\!] \ s \end{aligned}$$

-- By monad laws.

$$\begin{aligned} &\equiv \text{do } x \leftarrow \text{frequency } (\text{map } (\lambda (w, _, x) \rightarrow (w, \text{return } x)) \ xs) \\ &\quad \mathcal{P}[\![x \gg k]\!] \langle \$ \rangle \mathcal{R}[\![x \gg k]\!] \end{aligned}$$

-- By Lemma A.1.

$$\begin{aligned} &\equiv \text{do } x \leftarrow \text{frequency } (\text{map } (\lambda (w, _, x) \rightarrow (w, \text{return } x)) \ xs) \\ &\quad (\mathcal{P}[\![x]\!] \langle \$ \rangle \mathcal{R}[\![x]\!]) \gg \lambda (a, []) \rightarrow \mathcal{P}[\![k \ a]\!] \langle \$ \rangle \mathcal{R}[\![k \ a]\!] \end{aligned}$$

-- By IH.

$$\begin{aligned} &\equiv \text{do } x \leftarrow \text{frequency } (\text{map } (\lambda (w, _, x) \rightarrow (w, \text{return } x)) \ xs) \\ &\quad ((\lambda a \rightarrow (a, [])) \langle \$ \rangle \mathcal{G}[\![x]\!]) \gg \lambda (a, []) \rightarrow (\lambda a \rightarrow (a, [])) \langle \$ \rangle \mathcal{G}[\![k \ a]\!] \end{aligned}$$

-- By simplification .

$$\begin{aligned} &\equiv (\lambda a \rightarrow (a, [])) \langle \$ \rangle \text{do } x \leftarrow \text{frequency } (\text{map } (\lambda (w, _, x) \rightarrow (w, \text{return } x)) \ xs) \\ &\quad \quad a \leftarrow \mathcal{G}[\![x]\!] \\ &\quad \quad \mathcal{G}[\![k \ a]\!] \end{aligned}$$

-- By definition ($\mathcal{G}[\![\cdot]\!]$)

$$\equiv (\lambda a \rightarrow (a, [])) \langle \$ \rangle \mathcal{G}[\![\text{Bind } (\text{Pick } xs) \ k]\!]$$

Thus the decomposition of a free generator into a parser and a source of randomness is equivalent to interpreting it as a generator. \square

B PROOF OF LEMMA 4.1

LEMMA 4.1. δ_c satisfies equations (1), (2), (3), and (4). In other words, the free generator derivative behaves similarly to the regular expression derivative.

PROOF. We prove each equation individually.

- Equation 1: $\delta_c \text{ void} \equiv \text{void}$
By evaluation.
- Equation 2: $\delta_c(\text{return } v) \equiv \text{void}$
By definition.
- Equation 3:

$$\begin{aligned} \delta_c(\text{pick } xs) &\equiv x && \text{if } (c, x) \in xs \\ \delta_c(\text{pick } xs) &\equiv \text{void} && \text{if } (c, x) \notin xs \end{aligned}$$

Unfold the definition of `pick`, by evaluation.

- Equation 4:

$$\begin{aligned} \delta_c(x \gg f) &\equiv \delta_c(f \ a) && \text{if } vx = \{a\} \\ \delta_c(x \gg f) &\equiv \delta_c x \gg f && \text{if } vx = \emptyset \end{aligned}$$

- Case $x = \text{Return } a$. By definition, $vx = \{a\}$.

$$\begin{aligned} \delta_c(x \gg f) &\equiv \delta_c(\text{Return } a \gg f) \text{ -- By assumption.} \\ &\equiv \delta_c(f \ a) \text{ -- By definition } (\gg). \end{aligned}$$

- Case $x = \text{Bind } (\text{Pick } xs) \ g$. By definition, $vx = \emptyset$.

$$\begin{aligned} \delta_c(x \gg f) &\equiv \delta_c(\text{Bind } (\text{Pick } xs) \ g \gg f) && \text{-- By assumption.} \\ &\equiv \delta_c(\text{Bind } (\text{Pick } xs) \ (\lambda a \rightarrow g \ a \gg f)) && \text{-- By definition } (\gg). \\ &\equiv \text{case find } ((= c) \ . \text{snd}) \ xs \ \text{of} && \text{-- By definition } (\delta). \\ &\quad \text{Just } (_, _, x) \rightarrow x \gg (\lambda a \rightarrow g \ a \gg f) \\ &\quad \text{Nothing} \rightarrow \text{void} \\ &\equiv \text{case find } ((= c) \ . \text{snd}) \ xs \ \text{of} && \text{-- By monad laws.} \\ &\quad \text{Just } (_, _, x) \rightarrow (x \gg g) \gg f \\ &\quad \text{Nothing} \rightarrow \text{void} \\ &\equiv \delta_c x \gg f && \text{-- By definition } (\delta). \end{aligned}$$

Thus all four equations hold. □

C PROOF OF THEOREM 4.2

THEOREM 4.2. *The derivative of a free generator's language is the same as the language of its derivative. That is, for all free generators g and choices c ,*


$$\delta_c^{\mathcal{L}} \mathcal{L}[[g]] = \mathcal{L}[[\delta_c g]].$$

PROOF. $\mathcal{L}[[\delta_c x]] = \mathcal{L}[[\text{case } x \text{ of}$ -- By definition (δ).
 Return $_ \rightarrow \text{void}$
 Bind (Pick xs) $k \rightarrow \text{case find } ((= c) . \text{snd}) \text{ } xs \text{ of}$
 Just $(_, _, y) \rightarrow \mathcal{L}[[y \gg k]]$
 Nothing $\rightarrow []$]]
 = **case** x **of** -- By definition (\mathcal{L}).
 Return $_ \rightarrow []$
 Bind (Pick xs) $k \rightarrow \text{do}$
 $(_, d, y) \leftarrow xs$
 $cs \leftarrow \mathcal{L}[[y \gg k]]$
 guard $(c == d)$
 pure cs
 = **do** -- By Haskell identities .
 $(d : cs) \leftarrow \text{case } x \text{ of}$
 Return $_ \rightarrow [[]]$
 Bind (Pick xs) $k \rightarrow \text{do}$
 $(_, d, y) \leftarrow xs$
 $s \leftarrow \mathcal{L}[[y \gg k]]$
 pure $(d : s)$
 guard $(c == d)$
 pure cs
 = **do** -- By definition (\mathcal{L}).
 $(d : cs) \leftarrow \mathcal{L}[[x]]$
 guard $(c == d)$
 pure cs
 = $\delta_c^{\mathcal{L}}(\mathcal{L}[[x]])$ -- By definition ($\delta^{\mathcal{L}}$)

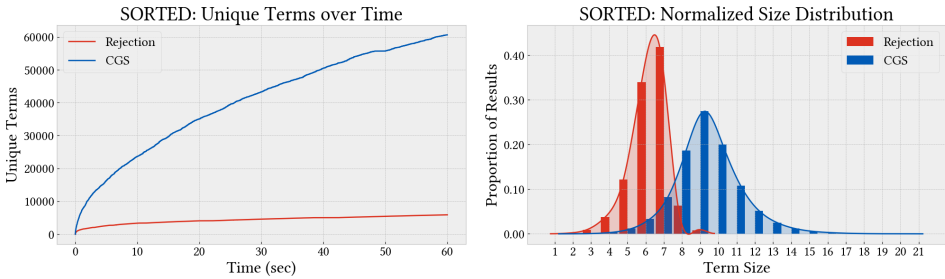
□

There is another proof of this theorem, suggested by Alexandra Silva, which uses the fact that 2^{Σ^*} is the final coalgebra, along with the observation that FGen has a $2 \times (-)^{\Sigma}$ coalgebraic structure. This approach is certainly more elegant, but it abstracts away some helpful operational intuition.


D FULL EXPERIMENTAL RESULTS



BST Charts



SORTED Charts



AVL Charts